

K-means clustering: algorithm

- **Step 1. Initialization:** Select an initial guess for all μ_k^0 for $k=1,2,\dots,K$, set iteration number $n=0$

K-means clustering: algorithm

- **Step 1. Initialization:** Select an initial guess for all μ_k^0 for $k=1,2,\dots,K$, set iteration number $n=0$
- **Step 2. Update configuration:** Set $X^{n+1}=0$, except where $k = \operatorname{argmin}_{k'=1,2,\dots,K} \|x_i - \mu_{k'}\|^2$ for which $X^{n+1}_{ik}=1$

K-means clustering: algorithm

- **Step 1. Initialization:** Select an initial guess for all μ_k^0 for $k=1,2,\dots,K$, set iteration number $n=0$
- **Step 2. Update configuration:** Set $X^{n+1}=0$, except where $k = \operatorname{argmin}_{k'=1,2,\dots,K} \|x_i - \mu_{k'}\|^2$ for which $X_{ik}^{n+1}=1$
- **Step 3. Update centroids:** Compute cluster averages, $\mu_k^{n+1} = 1/N_k \sum_{i=1,2,\dots,N} X_{ik}^{n+1} x_i$ where $N_k = \sum_{i=1,2,\dots,N} X_{ik}^{n+1}$,

K-means clustering: algorithm

- **Step 1. Initialization:** Select an initial guess for all μ_k^0 for $k=1,2,\dots,K$, set iteration number $n=0$
- **Step 2. Update configuration:** Set $X^{n+1}=0$, except where $k = \operatorname{argmin}_{k'=1,2,\dots,K} \|x_i - \mu_{k'}\|^2$ for which $X_{ik}^{n+1}=1$
- **Step 3. Update centroids:** Compute cluster averages, $\mu_k^{n+1} = 1/N_k \sum_{i=1,2,\dots,N} X_{ik}^{n+1} x_i$ where $N_k = \sum_{i=1,2,\dots,N} X_{ik}^{n+1}$,
- **Step 4. Convergence check:** If $n>0$ and $X^{n+1} = X^n$ then exit with solution $X^*=X^{n+1}$ and $\mu^*=\mu^{n+1}$,

K-means clustering: algorithm

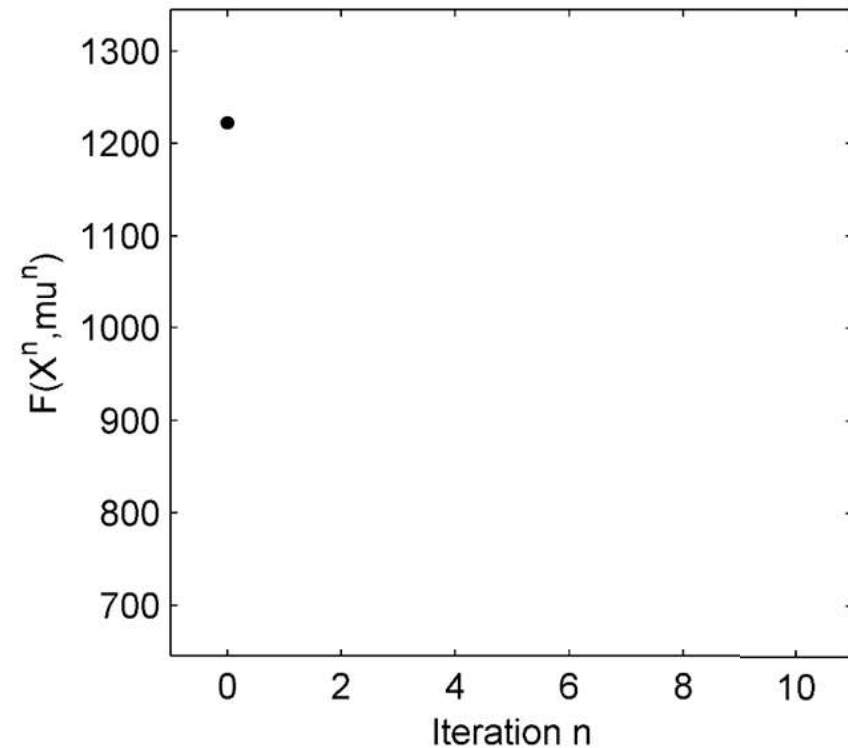
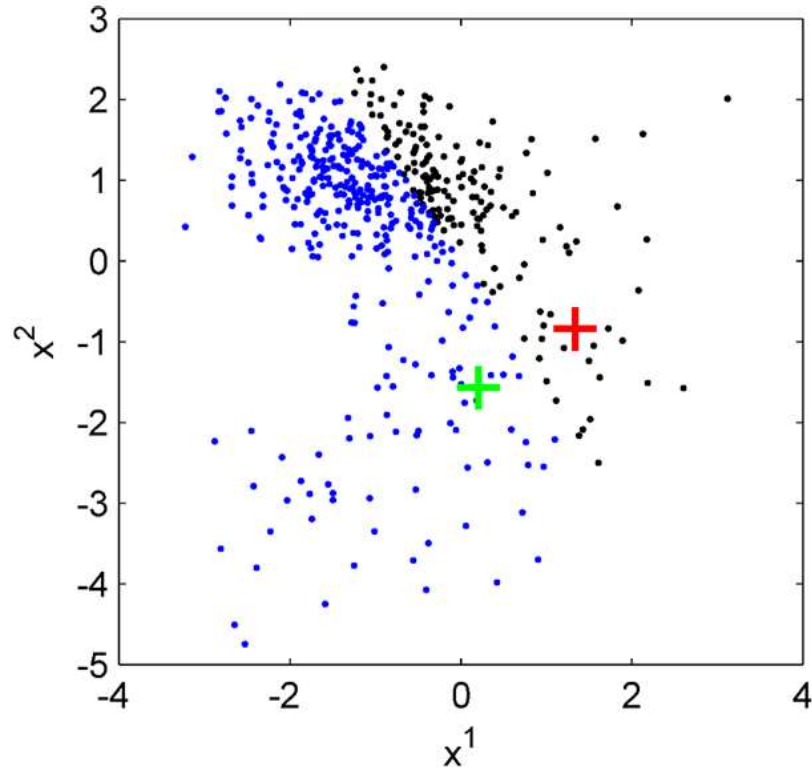
- **Step 1. Initialization:** Select an initial guess for all μ_k^0 for $k=1,2,\dots,K$, set iteration number $n=0$
- **Step 2. Update configuration:** Set $X^{n+1}=0$, except where $k = \operatorname{argmin}_{k'=1,2,\dots,K} \|x_i - \mu_{k'}\|^2$ for which $X^{n+1}_{ik}=1$
- **Step 3. Update centroids:** Compute cluster averages, $\mu_k^{n+1} = 1/N_k \sum_{i=1,2,\dots,N} X^{n+1}_{ik} x_i$ where $N_k = \sum_{i=1,2,\dots,N} X^{n+1}_{ik}$,
- **Step 4. Convergence check:** If $n>0$ and $X^{n+1} = X^n$ then exit with solution $X^*=X^{n+1}$ and $\mu^*=\mu^{n+1}$,
- **Step 5. Iteration:** update $n=n+1$ and go back to step 2.

Attendance Code:
41096890

K-means clustering algorithm in action

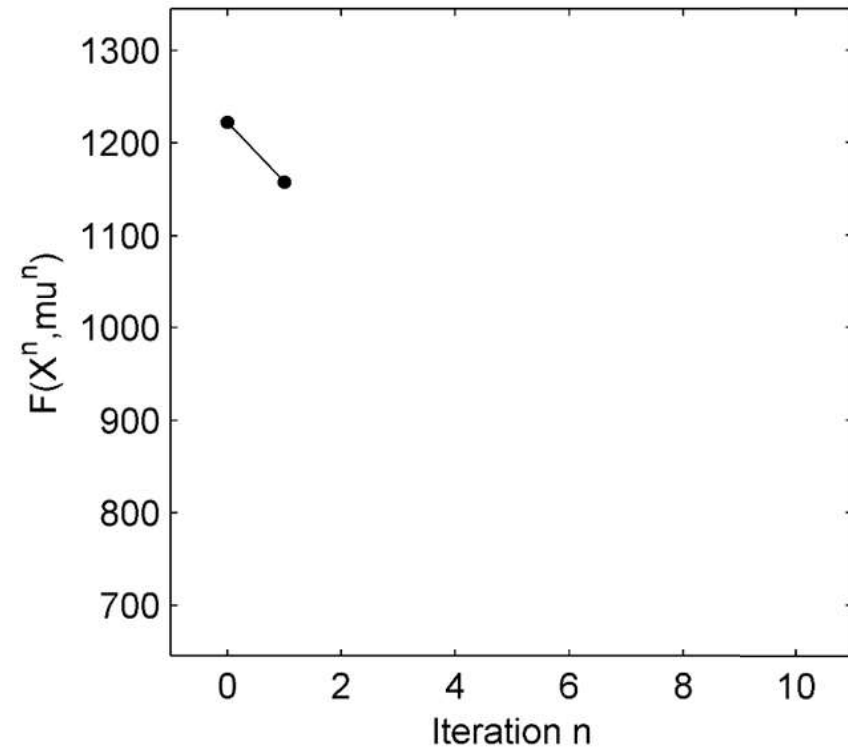
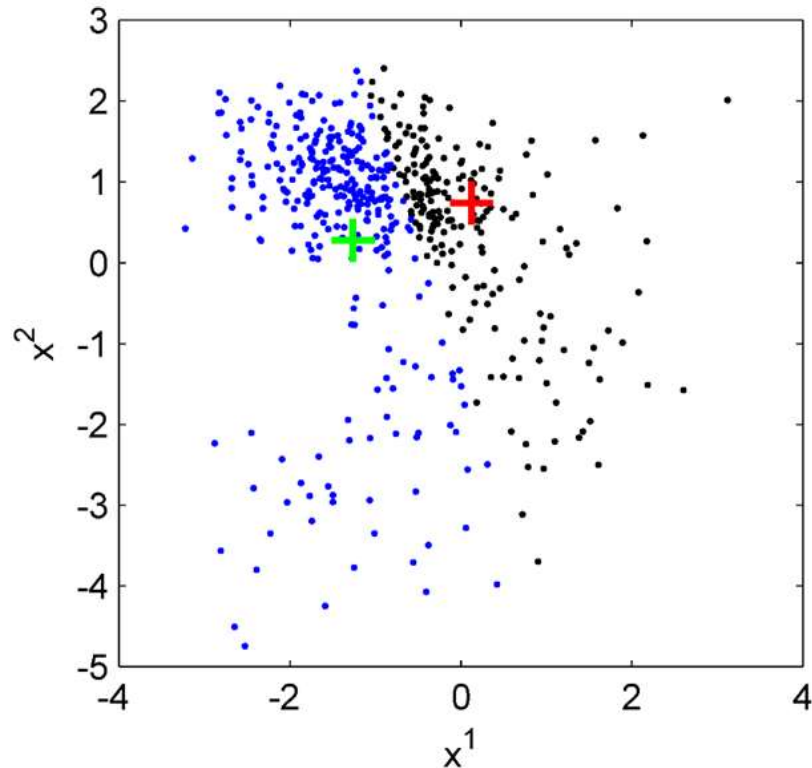
$K = 2$, $n=0$, initial guess for μ_1^0, μ_2^0 (initialize)

Assign X_{ik}^1 to all data points ($i = 1, 2, \dots, N$)



K-means clustering algorithm in action

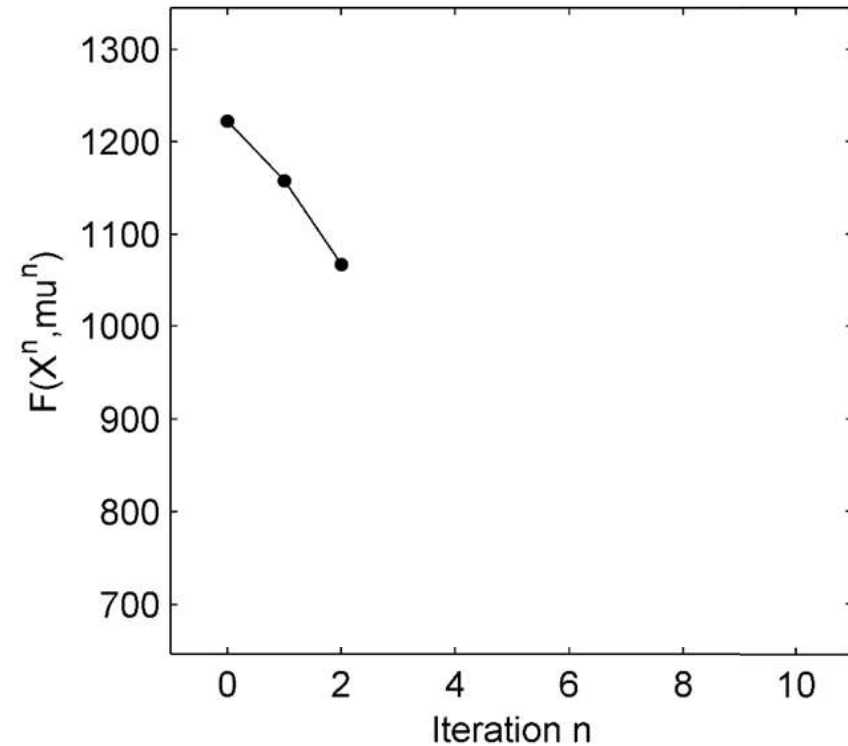
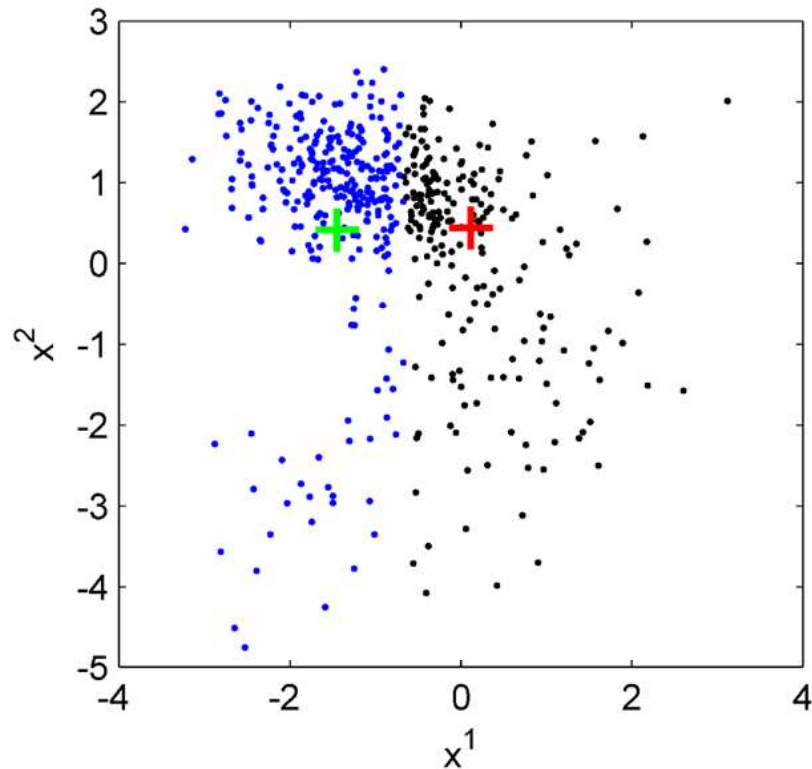
Update centroids μ_1^1, μ_2^1 with current configuration X_{ik}^1
Assign X_{ik}^2 ; new configuration is different from X_{ik}^1



K-means clustering algorithm in action



Update centroids μ_1^2, μ_2^2 with current configuration X_{ik}^2

Assign X_{ik}^3 ; new configuration is different from X_{ik}^2



K-means clustering algorithm in action

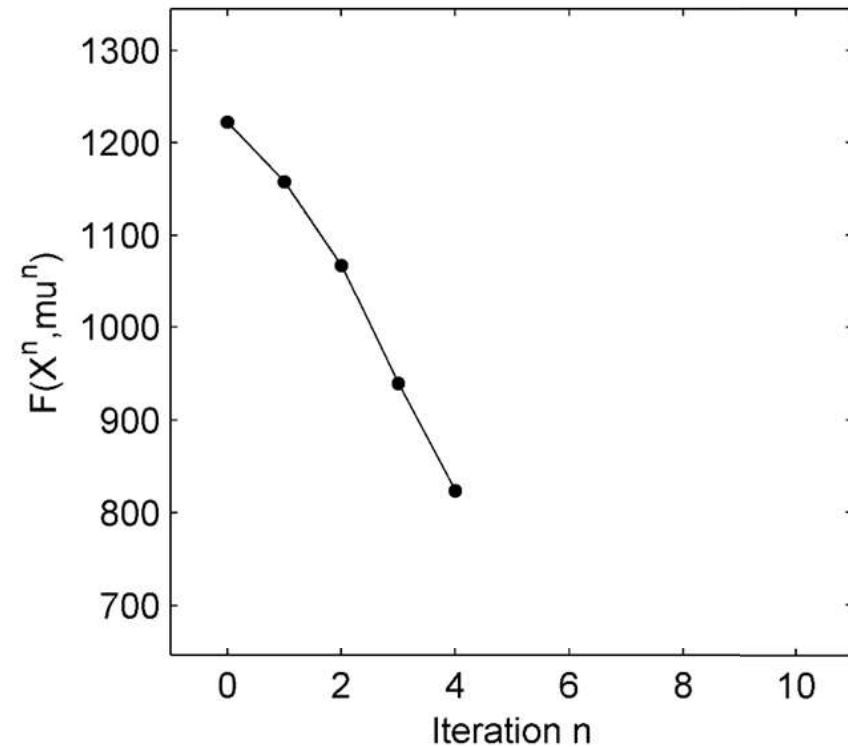

Update centroids μ_1^3, μ_2^3 with current configuration X_{ik}^3
Assign X_{ik}^4 ; new configuration is different from X_{ik}^3



K-means clustering algorithm in action

Update centroids μ_1^4, μ_2^4 with current configuration X_{ik}^4

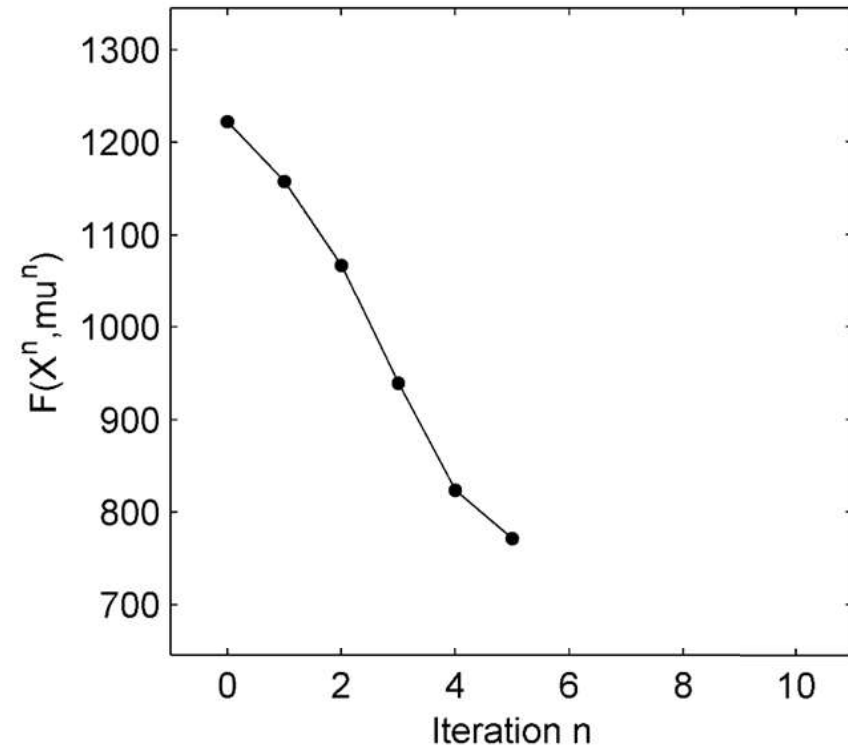
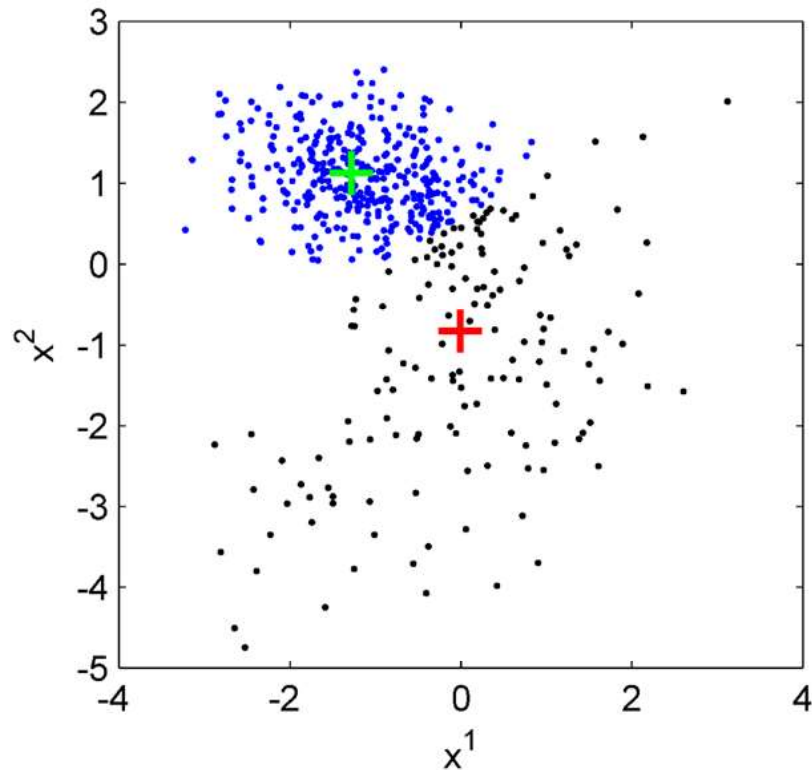
Assign X_{ik}^5 ; new configuration is different from X_{ik}^4



K-means clustering algorithm in action

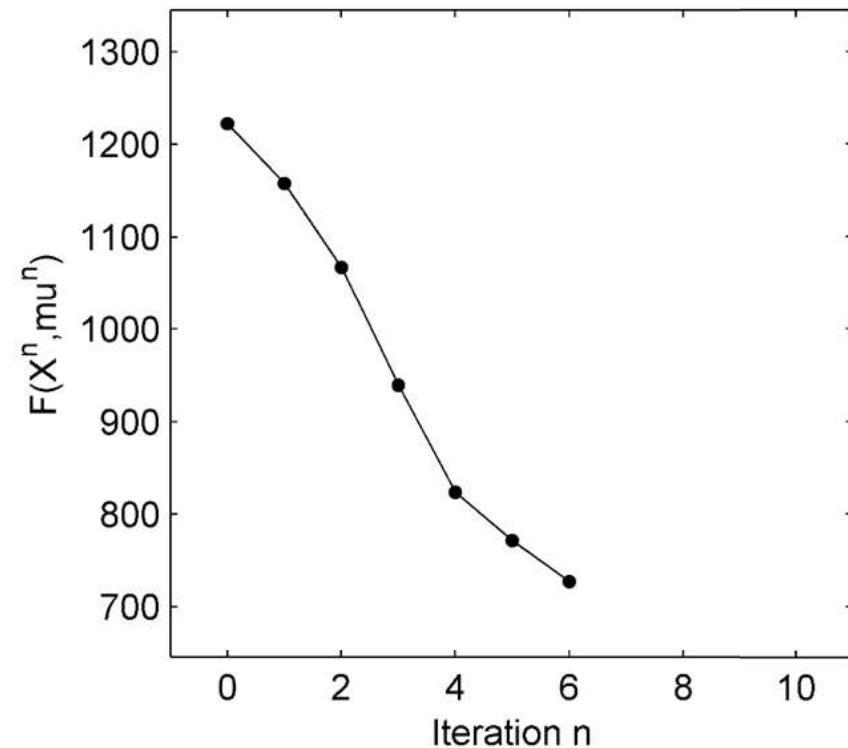
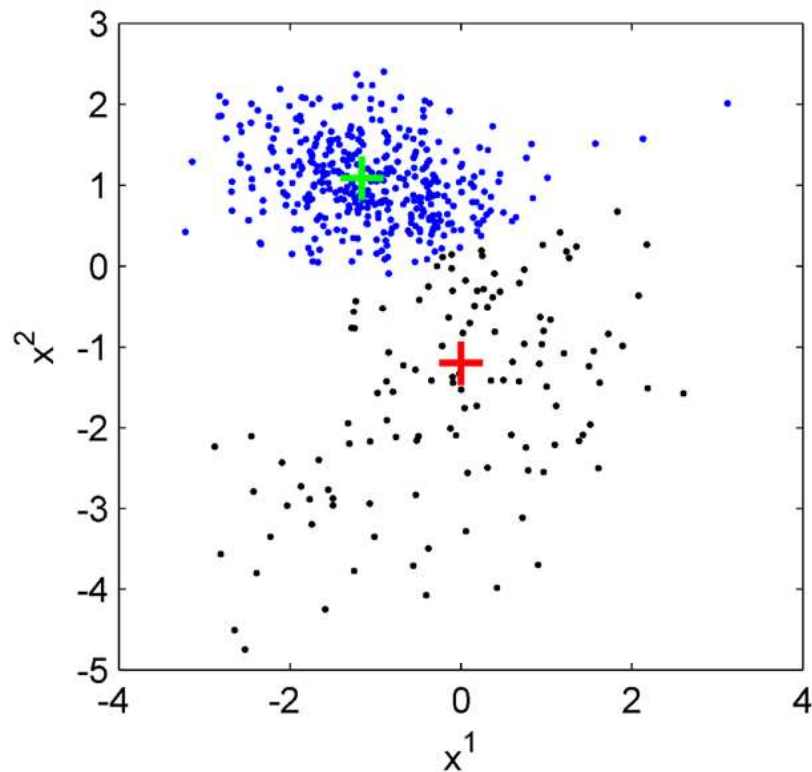
Update centroids μ_1^5, μ_2^5 with current configuration X_{ik}^5

Assign X_{ik}^6 ; new configuration is different from X_{ik}^5





K-means clustering algorithm in action

Update centroids μ_1^6, μ_2^6 with current configuration X_{ik}^6
Assign X_{ik}^7 ; new configuration is different from X_{ik}^6



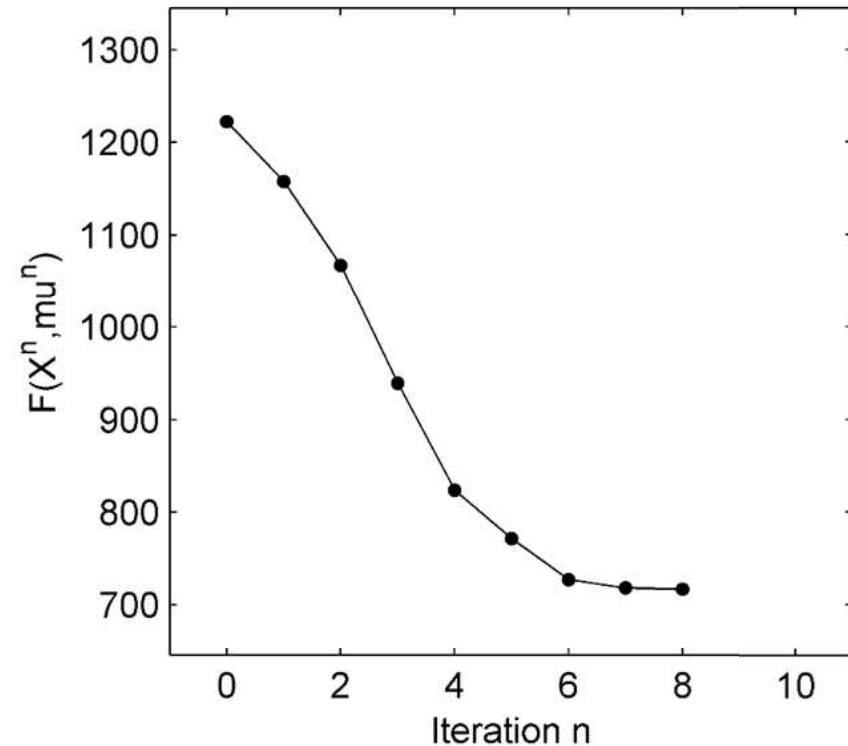

K-means clustering algorithm in action

Update centroids μ_1^7, μ_2^7 with current configuration X_{ik}^7
Assign X_{ik}^8 ; new configuration is different from X_{ik}^7



K-means clustering algorithm in action

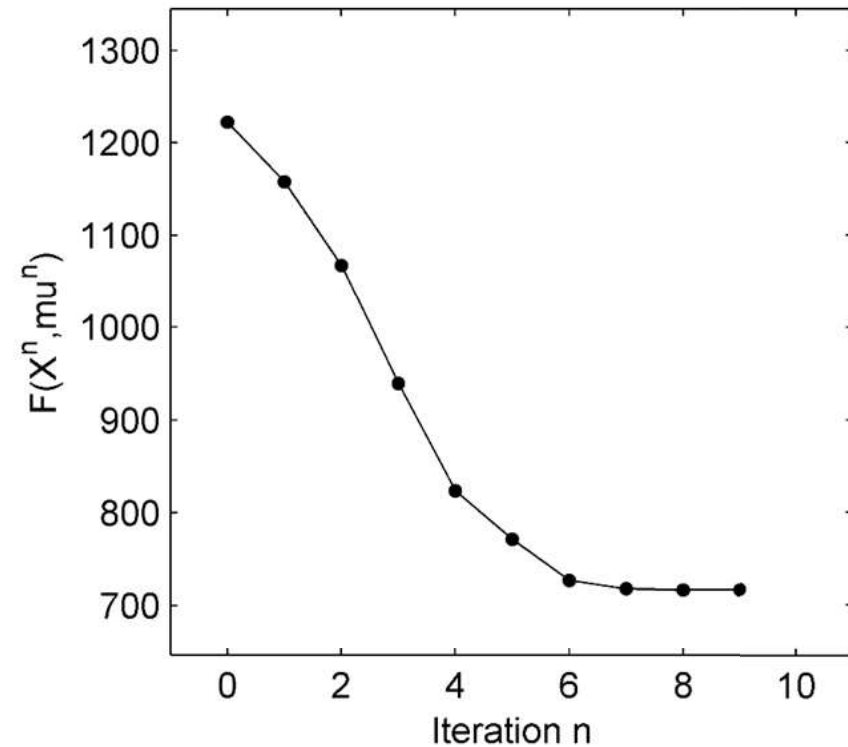

Update centroids μ_1^8, μ_2^8 with current configuration X_{ik}^8
Assign X_{ik}^9 ; new configuration is different from X_{ik}^8



K-means clustering algorithm in action

Update centroids μ_1^9, μ_2^9 with current configuration X_{ik}^9

Assign $X_{ik}^{10} = X_{ik}^9 = X^*$

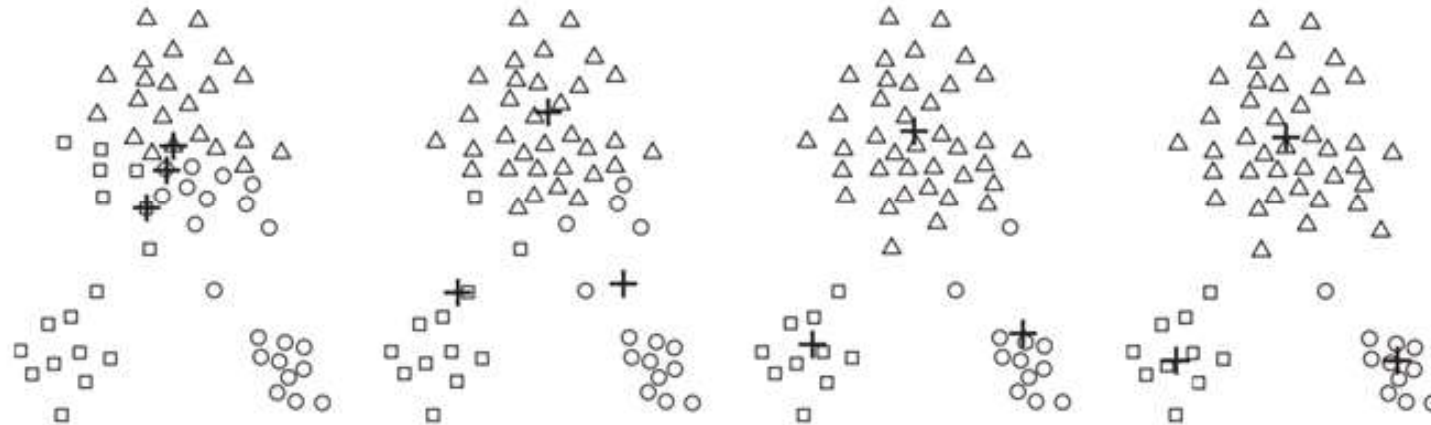


K-means clustering: analysis

- We can show that $F(X^{n+1}, \mu^{n+1}) \leq F(X^n, \mu^n)$ i.e. the K-means objective function is **never increasing** (because given μ we can find the X with globally smallest F , and vice-versa)
- Thus, the K-means algorithm always **converges** on a **fixed point** (that is, it finds a **minima** of F and it would stay there for all subsequent iterations)
- But we cannot know if the local minima is also a global minima (this is an approximate method)
- Do not know how long it takes to converge (typically 7-15 iterations from experience, but no guarantees)
- Potential for pathological results (i.e. empty clusters)

Illustration

#Code



Initialization

Iteration 1

Iteration 2

Iteration 3



UNIVERSITY OF
BIRMINGHAM
DUBAI

Example 1: Clustering of Medicines (K=2)

#Code

	Weight index	PH
Med A	1	1
Med B	2	1
Med C	4	3
Med D	5	4

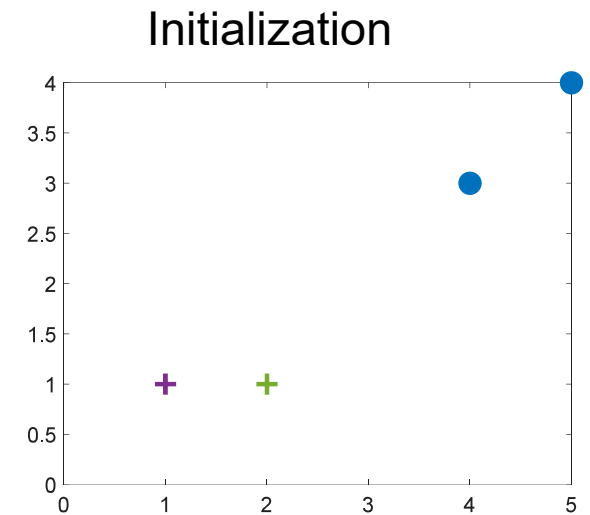
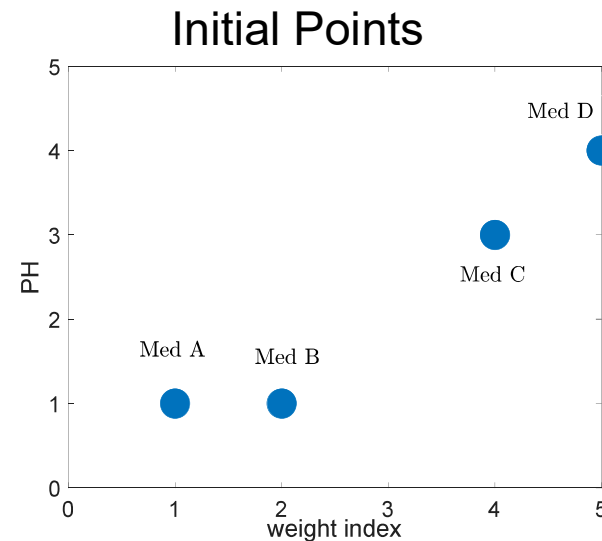


UNIVERSITY OF
BIRMINGHAM
DUBAI

Example 1: Clustering of Medicines (K=2)

#Code

	Weight index	PH
Med A	1	1
Med B	2	1
Med C	4	3
Med D	5	4



Initialization: Initial centroids be Med A and Med B i.e, $c_1 = (1,1)$ and $c_2 = (2,1)$



UNIVERSITY OF
BIRMINGHAM
DUBAI

Iteration 1: Step 1

#Code

1. Calculate (Euclidean) distance of each point to cluster centroids to form an **Object-Centroid Distance Matrix**:

	Med A	Med B	Med C	Med D
c_1	0		13	25
c_2		0	8	18

$$d_{Euc}(Med\ C, C_1)^2 = (4 - 1)^2 + (3 - 1)^2 = 13$$

$$d_{Euc}(Med\ C, C_2)^2 = (4 - 2)^2 + (3 - 1)^2 = 8$$

$$d_{Euc}(Med\ D, C_1)^2 = (5 - 1)^2 + (4 - 1)^2 = 25$$

$$d_{Euc}(Med\ D, C_2)^2 = (5 - 2)^2 + (4 - 1)^2 = 18$$

	Weight index	PH
Med A	1	1
Med B	2	1
Med C	4	3
Med D	5	4



Thus, Medicines B, C and D assigned to Cluster 2.



UNIVERSITY OF
BIRMINGHAM
DUBAI

After step 1 of iteration 1

#Code





UNIVERSITY OF
BIRMINGHAM
DUBAI

#Code

Iteration 1: Step 2

2. Update the centroids of the cluster.

$$c_1 = c_1 \text{ (same)}; c_2 = \frac{\text{Med B} + \text{Med C} + \text{Med D}}{3}$$
$$= \left(\frac{2+4+5}{3}, \frac{1+3+4}{3} \right) = (3.67, 2.67)$$



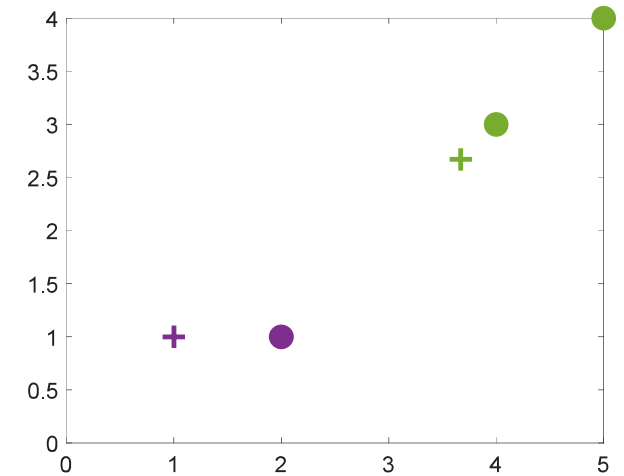
UNIVERSITY OF
BIRMINGHAM
DUBAI

Iteration 2: Step 1

1. Calculate distance of each point to new cluster centroids.

	Med A	Med B	Med C	Med D
c_1	0	1	13	25
c_2	9.92	5.56	0.22	3.53

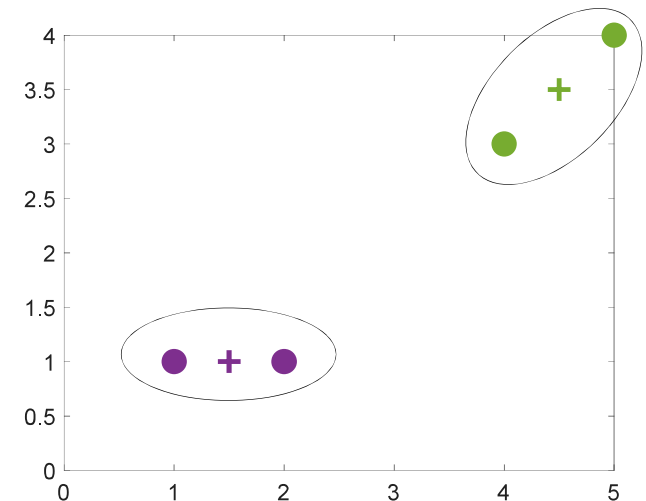
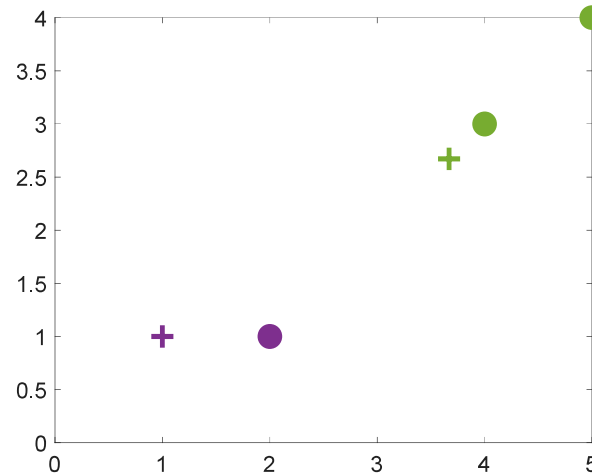
Med B is thus moved to cluster 1.



Iteration 2: Step 2

2. Update the centroids of the cluster.

- $c_1 = \frac{Med A + Med B}{2} = \left(\frac{1+2}{2}, \frac{1+1}{2}\right) = (1.5, 1)$
- $c_2 = \frac{Med C + Med D}{2} = \left(\frac{4+5}{2}, \frac{3+4}{2}\right) = (4.5, 3.5)$



#Code

- Repeat the same steps in iteration 3
- Note that cluster assignments do not change
- Algorithm converged.



UNIVERSITY OF
BIRMINGHAM
DUBAI

K-means clustering: example

- **Histological Analysis:** tissue stained with hemotoxylin and eosin (H&E)

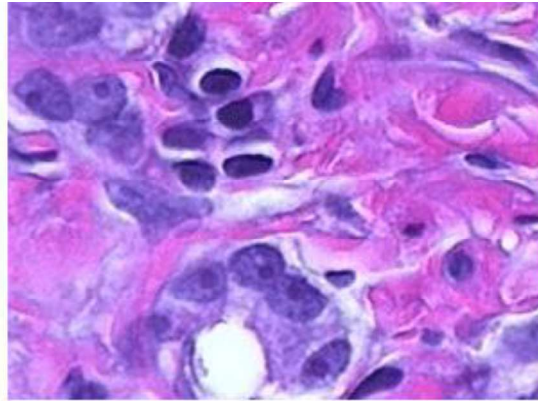




Image courtesy of Alan Partin, Johns Hopkins University


Objects in Cluster 1



Objects in Cluster 2



Objects in Cluster 3



Machine learning (ML): overview

- Given some **training data**, ML training finds optimal **model parameters** such that the **prediction error** is minimized; the smaller the error function the better

$$w^{\star} = \arg \min_{w' \in \mathcal{W}} F(w')$$

- Then use the trained parameters in the model, to make predictions about new unseen **test data**
- ML algorithms usually categorized according to availability of labelled data: **supervised**, **unsupervised**, **self-supervised**, **transfer learning**

Sequential gradient descent (SGD)

- General algorithm for finding a value of the model parameters w such that error function $F(w)$ is minimized, expressed using **multivariable calculus** as $F_w(w) = 0$, where F_w is the **(partial) derivative** of F with respect to vector w
- Idea: starting with a guess for w_n , take a "step" in direction of **steepest descent** of the loss function, $-F_w(w_n)$, use this as a better guess w_{n+1}
- Size of the step, $\alpha > 0$ ("learning rate"), determines how quickly the minimum is reached, but can **overshoot** and also **diverge** if α is too large; not guaranteed to find the minimum, unless the error function is **convex** with respect to w

SGD: algorithm

- **Step 1.** *Initialization:* Select an initial guess for w_0 , a convergence tolerance $\varepsilon > 0$, step size (learning rate) parameter $\alpha > 0$, set iteration number $n=0$
- **Step 2.** *Gradient descent step:* Compute new model parameters,

$$w_{n+1} = w_n - \alpha F_w(w_n)$$

- **Step 3.** *Convergence test:* Compute new loss function value $F(w_{n+1})$, and loss function improvement, $\Delta F = |F(w_{n+1}) - F(w_n)|$ and if $\Delta F < \varepsilon$, exit with solution $w^*=w_{n+1}$
- **Step 4.** *Iteration:* update $n=n+1$ and go to step 2.

SGD Example: Regression

#Code

(Linear Regression)

- Yield prediction of very early potato cultivars before harvest.
- Potato yielding is associated with the amount of nitrogen fertilizer and average temperature in the season. Here is some data collected from several farms.
- Task: predict how many potatoes will yield in a given farm by measuring the independent variables.

Farm	Fertilizer (10*kg)	Avg. Temp. (°C)	Potato (10*t)
A	12.1	12.5	36.5
B	8.7	11.2	26.4
C	14.0	14.7	42.0
D	9.5	11.8	28.8
E	13.2	13.6	39.7

SGD Example: Regression

(Linear Regression)

- Two independent features
 - Fertilizer (x^1)
 - Average Temperature (x^2)

- Regression model

$$\begin{aligned}
 f(w, x) &= w_1 + w_2 x^1 + w_3 x^2 \\
 &= [w_1 \quad w_2 \quad w_3] x \\
 &= w^T x
 \end{aligned}
 \quad x = \begin{bmatrix} 1 \\ x^1 \\ x^2 \end{bmatrix}$$

- Sum-of-squares error function

$$F(w) = \sum_{i=1}^N (w^T x_i - y_i)^2 \quad F_w(w) = 2 \sum_{i=1}^N (w^T x_i - y_i) x_i$$

	x^1	x^2	y
$i = 1$	12.1	12.5	36.5
$i = 2$	8.7	11.2	26.4
$i = 3$	14.0	14.7	42.0
$i = 4$	9.5	11.8	28.8
$i = 5$	13.2	13.6	39.7

SGD Example: Regression

(Linear Regression)

- Initial guess: $w_0^T = [1 \ 0 \ 0]$
- Tolerance: $\epsilon = 1$
- Learning rate: $\alpha = 0.0005$

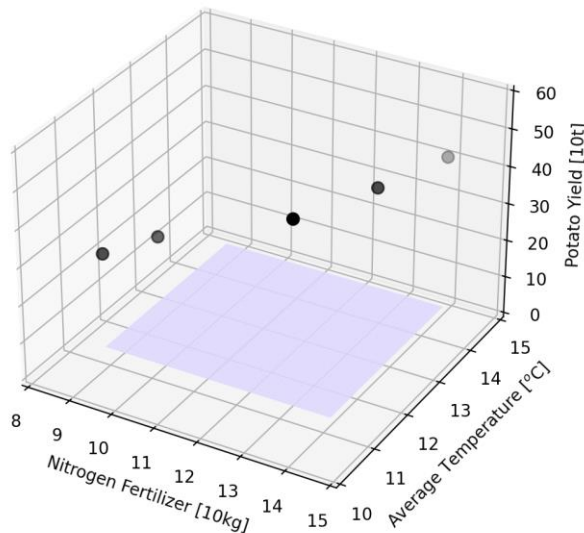
- Error function evaluation:

$$F(w_0) = \sum_{i=1}^5 \left([1 \ 0 \ 0] \begin{bmatrix} 1 \\ x_i^1 \\ x_i^2 \end{bmatrix} - y_i \right)^2 = 5856.9$$

- Gradient evaluation

$$F_w(w_0) = 2 \sum_{i=1}^N (w^T x_i - y_i) \begin{bmatrix} 1 \\ x_i^1 \\ x_i^2 \end{bmatrix} = \begin{bmatrix} -336.8 \\ -3998.9 \\ -4370.58 \end{bmatrix}$$

Iteration: n = 0



- Gradient descent step

$$w_1 = w_0 - \alpha F_w(w_0) = \begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix} - 0.0005 \begin{bmatrix} -336.8 \\ -3998.9 \\ -4370.58 \end{bmatrix} = \begin{bmatrix} 1.168 \\ 1.999 \\ 2.185 \end{bmatrix}$$

- Convergence test

$$\Delta F = |F(w_1) - F(w_0)| = |1513.7 - 5856.9| = 4343.2$$

SGD Example: Regression

(Linear Regression)

- Current weight:

$$w_1^T = [1.168 \quad 1.999 \quad 2.185]$$

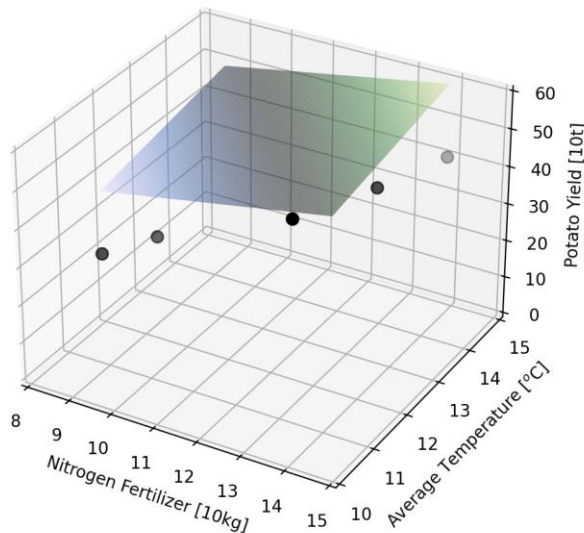
- Error function evaluation:

$$F(w_1) = \sum_{i=1}^5 \left(w_1^T \begin{bmatrix} 1 \\ x_i^1 \\ x_i^2 \end{bmatrix} - y_i \right)^2 = 1513.7$$

- Gradient evaluation

$$F_w(w_1) = 2 \sum_{i=1}^N (w^T x_i - y_i) \begin{bmatrix} 1 \\ x_i^1 \\ x_i^2 \end{bmatrix} = \begin{bmatrix} 173.7 \\ 2011.0 \\ 2227.1 \end{bmatrix}$$

Iteration: n = 1



- Gradient descent step

$$w_2 = w_1 - \alpha F_w(w_1) = \begin{bmatrix} 1.168 \\ 1.999 \\ 2.185 \end{bmatrix} - 0.0005 \begin{bmatrix} 173.7 \\ 2011.0 \\ 2227.1 \end{bmatrix} = \begin{bmatrix} 1.082 \\ 0.994 \\ 1.072 \end{bmatrix}$$

- Convergence test

$$\Delta F = |F(w_2) - F(w_1)| = |399.2 - 1513.7| = 1114.5 > \epsilon$$

SGD Example: Regression

(Linear Regression)

- Current weight:

$$w_2^T = [1.082 \quad 0.994 \quad 1.072]$$

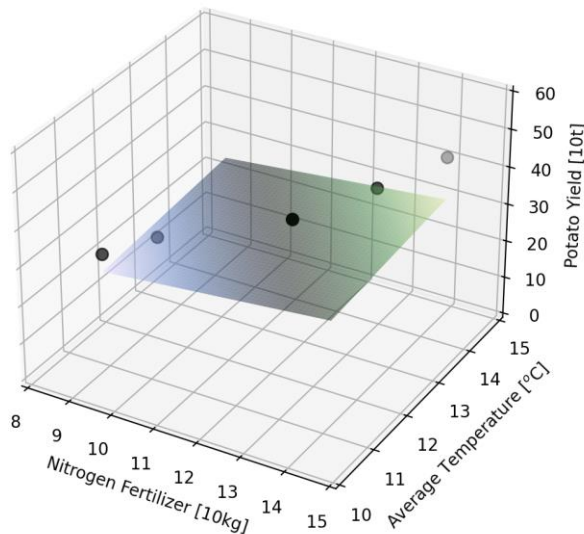
- Error function evaluation:

$$F(w_2) = \sum_{i=1}^5 \left(w_2^T \begin{bmatrix} 1 \\ x_i^1 \\ x_i^2 \end{bmatrix} - y_i \right)^2 = 399.2$$

- Gradient evaluation

$$F_w(w_2) = 2 \sum_{i=1}^N (w^T x_i - y_i) \begin{bmatrix} 1 \\ x_i^1 \\ x_i^2 \end{bmatrix} = \begin{bmatrix} -84.923 \\ -1033.4 \\ -1115.0 \end{bmatrix}$$

Iteration: n = 2



- Gradient descent step

$$w_3 = w_2 - \alpha F_w(w_2) = \begin{bmatrix} 1.082 \\ 0.994 \\ 1.072 \end{bmatrix} - 0.0005 \begin{bmatrix} -84.923 \\ -1033.4 \\ -1115.0 \end{bmatrix} = \begin{bmatrix} 1.124 \\ 1.511 \\ 1.629 \end{bmatrix}$$

- Convergence test

$$\Delta F = |F(w_3) - F(w_2)| = |113.1 - 399.2| = 286.1 > \epsilon$$

SGD Example: Regression

(Linear Regression)

- Current weight:

$$w_3^T = [1.124 \quad 1.511 \quad 1.629]$$

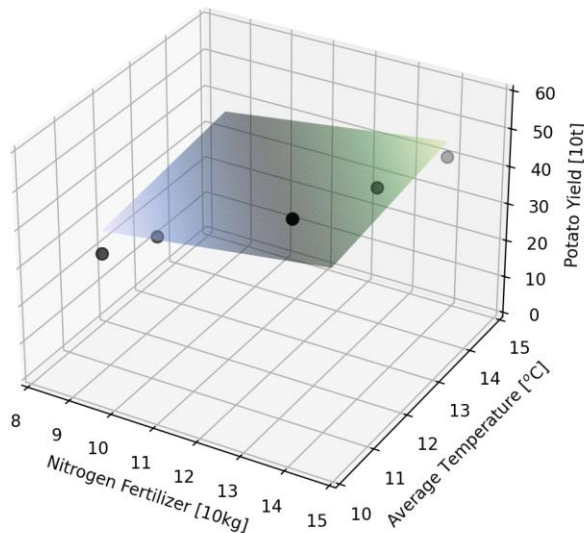
- Error function evaluation:

$$F(w_3) = \sum_{i=1}^5 \left(w_3^T \begin{bmatrix} 1 \\ x_i^1 \\ x_i^2 \end{bmatrix} - y_i \right)^2 = 113.1$$

- Gradient evaluation

$$F_w(w_3) = 2 \sum_{i=1}^N (w^T x_i - y_i) \begin{bmatrix} 1 \\ x_i^1 \\ x_i^2 \end{bmatrix} = \begin{bmatrix} 46.058 \\ 508.84 \\ 577.9 \end{bmatrix}$$

Iteration: n = 3



- Gradient descent step

$$w_4 = w_3 - \alpha F_w(w_3) = \begin{bmatrix} 1.124 \\ 1.511 \\ 1.629 \end{bmatrix} - 0.0005 \begin{bmatrix} 46.058 \\ 508.84 \\ 577.9 \end{bmatrix} = \begin{bmatrix} 1.101 \\ 1.256 \\ 1.340 \end{bmatrix}$$

- Convergence test

$$\Delta F = |F(w_4) - F(w_3)| = |39.661 - 113.1| = 73.5 > \epsilon$$

SGD Example: Regression

(Linear Regression)

- Current weight:

$$w_4^T = [1.101 \quad 1.256 \quad 1.340]$$

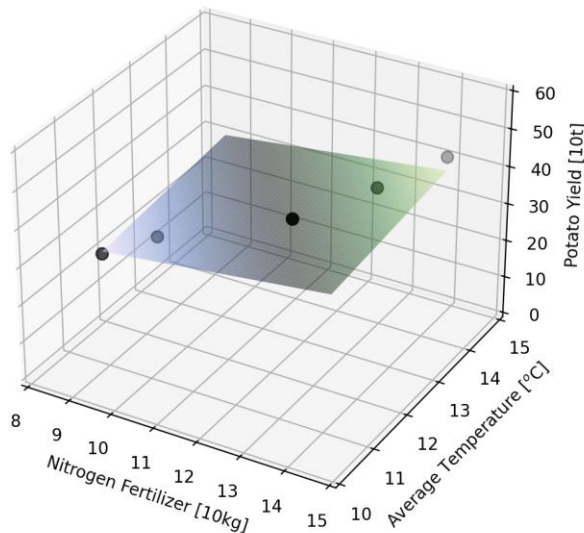
- Error function evaluation:

$$F(w_4) = \sum_{i=1}^5 \left(w_4^T \begin{bmatrix} 1 \\ x_i^1 \\ x_i^2 \end{bmatrix} - y_i \right)^2 = 39.661$$

- Gradient evaluation

$$F_w(w_4) = 2 \sum_{i=1}^N (w^T x_i - y_i) \begin{bmatrix} 1 \\ x_i^1 \\ x_i^2 \end{bmatrix} = \begin{bmatrix} -20.3 \\ -272.3 \\ -279.7 \end{bmatrix}$$

Iteration: n = 4



- Gradient descent step

$$w_5 = w_4 - \alpha F_w(w_4) = \begin{bmatrix} 1.101 \\ 1.256 \\ 1.340 \end{bmatrix} - 0.0005 \begin{bmatrix} -20.3 \\ -272.3 \\ -279.7 \end{bmatrix} = \begin{bmatrix} 1.111 \\ 1.394 \\ 1.480 \end{bmatrix}$$

- Convergence test

$$\Delta F = |F(w_5) - F(w_4)| = |20.747 - 39.661| = 18.9 > \epsilon$$

SGD Example: Regression

(Linear Regression)

- Current weight:

$$w_5^T = [1.111 \quad 1.394 \quad 1.480]$$

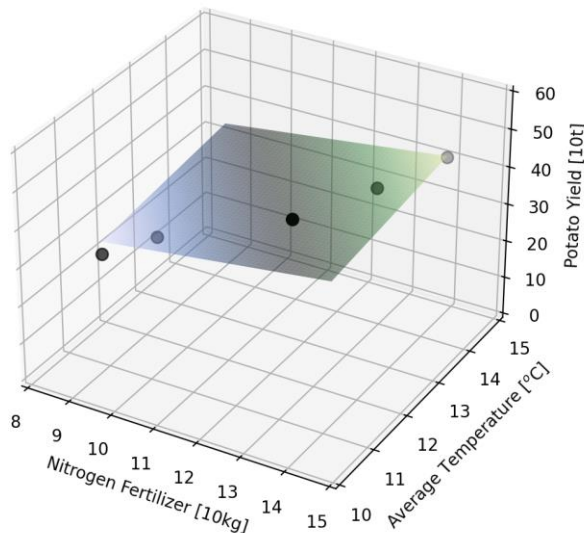
- Error function evaluation:

$$F(w_5) = \sum_{i=1}^5 \left(w_5^T \begin{bmatrix} 1 \\ x_i^1 \\ x_i^2 \end{bmatrix} - y_i \right)^2 = 20.747$$

- Gradient evaluation

$$F_w(w_5) = 2 \sum_{i=1}^N (w^T x_i - y_i) \begin{bmatrix} 1 \\ x_i^1 \\ x_i^2 \end{bmatrix} = \begin{bmatrix} 13.3 \\ 123.4 \\ 154.7 \end{bmatrix}$$

Iteration: n = 5



- Gradient descent step

$$w_6 = w_5 - \alpha F_w(w_5) = \begin{bmatrix} 1.111 \\ 1.394 \\ 1.480 \end{bmatrix} - 0.0005 \begin{bmatrix} 13.3 \\ 123.4 \\ 154.7 \end{bmatrix} = \begin{bmatrix} 1.104 \\ 1.331 \\ 1.403 \end{bmatrix}$$

- Convergence test

$$\Delta F = |F(w_6) - F(w_5)| = |15.830 - 20.747| = 4.92 > \epsilon$$

SGD Example: Regression

(Linear Regression)

- Current weight:

$$w_6^T = [1.104 \quad 1.331 \quad 1.403]$$


- Error function evaluation:

$$F(w_6) = \sum_{i=1}^5 \left(w_6^T \begin{bmatrix} 1 \\ x_i^1 \\ x_i^2 \end{bmatrix} - y_i \right)^2 = 15.830$$

- Gradient evaluation

$$F_w(w_6) = 2 \sum_{i=1}^N (w^T x_i - y_i) \begin{bmatrix} 1 \\ x_i^1 \\ x_i^2 \end{bmatrix} = \begin{bmatrix} -3.727 \\ -76.99 \\ -65.40 \end{bmatrix}$$

Iteration: n = 6



- Gradient descent step

$$w_7 = w_6 - \alpha F_w(w_6) = \begin{bmatrix} 1.104 \\ 1.331 \\ 1.403 \end{bmatrix} - 0.0005 \begin{bmatrix} -3.727 \\ -76.99 \\ -65.40 \end{bmatrix} = \begin{bmatrix} 1.106 \\ 1.369 \\ 1.436 \end{bmatrix}$$

- Convergence test

$$\Delta F = |F(w_7) - F(w_6)| = |14.505 - 15.830| = 1.33 > \epsilon$$

SGD Example: Regression

(Linear Regression)

- Current weight:

$$w_7^T = [1.106 \quad 1.369 \quad 1.436]$$

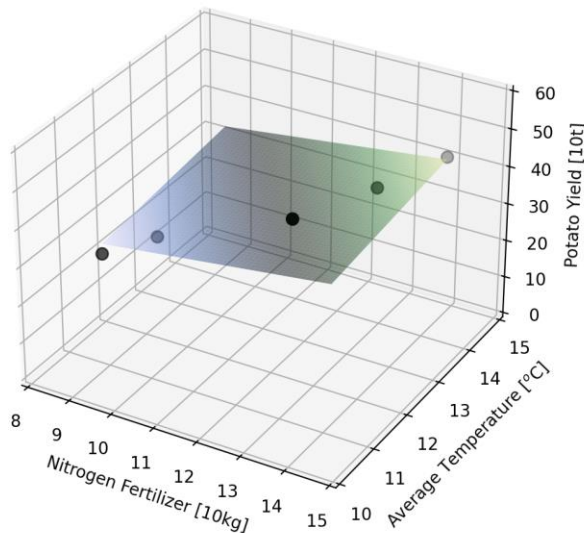
- Error function evaluation:

$$F(w_7) = \sum_{i=1}^5 \left(w_7^T \begin{bmatrix} 1 \\ x_i^1 \\ x_i^2 \end{bmatrix} - y_i \right)^2 = 14.505$$

- Gradient evaluation

$$F_w(w_7) = 2 \sum_{i=1}^N (w^T x_i - y_i) \begin{bmatrix} 1 \\ x_i^1 \\ x_i^2 \end{bmatrix} = \begin{bmatrix} 4.891 \\ 24.57 \\ 46.04 \end{bmatrix}$$

Iteration: n = 7



- Gradient descent step

$$w_8 = w_7 - \alpha F_w(w_7) = \begin{bmatrix} 1.106 \\ 1.369 \\ 1.436 \end{bmatrix} - 0.0005 \begin{bmatrix} 4.891 \\ 24.57 \\ 46.04 \end{bmatrix} = \begin{bmatrix} 1.104 \\ 1.357 \\ 1.413 \end{bmatrix}$$

- Convergence test

$$\Delta F = |F(w_8) - F(w_7)| = |14.103 - 14.505| = 0.40 < \epsilon$$


SGD Example: Regression

(Linear Regression)

- Optimal weight:

$$w^* = w_8 = \begin{bmatrix} 1.104 \\ 1.357 \\ 1.413 \end{bmatrix}$$

- Suppose we have a new farm that used 10 kg of nitrogen fertilizer in a season with average temperature of 15 °C. How many potatoes would you expect to yield on this farm?



- Optimized model

$$f(w, x) = w^T x = \begin{bmatrix} 1.104 & 1.357 & 1.413 \end{bmatrix} \begin{bmatrix} 1 \\ 10 \\ 15 \end{bmatrix} = 35.9 \text{ t}$$

Regression: analysis

- For linear regression, error function is convex so SGD with correct parameters, can converge on the globally optimal solution eventually (N.B. do not need to use SGD, there is a straightforward analytical solution)
- Unlike linear regression, more complex regression can often be made to fit N training data points **exactly** (zero training error), but this model will not **generalize well** to another random sample of data from the same circumstance
- A **trade off** between **model complexity** and **test error**: common feature of most machine learning models: want a model which is **as simple as possible** but **no simpler** (Occam's razor)

To recap

- We reviewed the sequential gradient descent algorithm in the context of regression
- **Next lecture:** Classification in ML