

Artificial Intelligence and Machine Learning (AIML)

2023–24



Attendance Code:
26444542



Attendance Code:
26444542

- **Last lecture:** classification in ML
- **This lecture:** classification using the perceptron algorithm

Example: health insurance company

- Data on whether customers bought the plan

Client	Age (yrs)	Income (k £)	Bought?
1	25	30	No
2	45	60	Yes
3	30	50	Yes
4	22	25	No
5	35	45	Yes
6	55	70	Yes
7	40	55	No
8	60	80	Yes
9	50	40	No
10	28	35	No

- Task: predict whether a new customer is likely to buy or not the plan, given their age and income.
 - **Goal:** split the data into 2 **classes** (bought/didn't buy) that best match **class-labeled training data**.

Example: health insurance company

- Data on whether customers bought the plan

Client	Age (yrs)	Income (k £)	Bought?
1	25	30	No
2	45	60	Yes
3	30	50	Yes
4	22	25	No
5	35	45	Yes
6	55	70	Yes
7	40	55	No
8	60	80	Yes
9	50	40	No
10	28	35	No

- Task: predict whether a new customer is likely to buy or not the plan, given their age and income.
 - **Goal:** split the data into 2 **classes** (bought/didn't buy) that best match **class-labeled training data**.
 - **Classification model:** linear classifier

$$f(w, x) = \text{sign}(w^T x)$$

$$f: \mathbb{R}^D \rightarrow \{-1, +1\}$$

Example: health insurance company

- Data on whether customers bought the plan

Client	Age (yrs)	Income (k £)	Bought?
1	25	30	No
2	45	60	Yes
3	30	50	Yes
4	22	25	No
5	35	45	Yes
6	55	70	Yes
7	40	55	No
8	60	80	Yes
9	50	40	No
10	28	35	No

- Task: predict whether a new customer is likely to buy or not the plan, given their age and income.

- **Goal:** split the data into 2 **classes** (bought/didn't buy) that best match **class-labeled training data**.
- **Classification model:** linear classifier

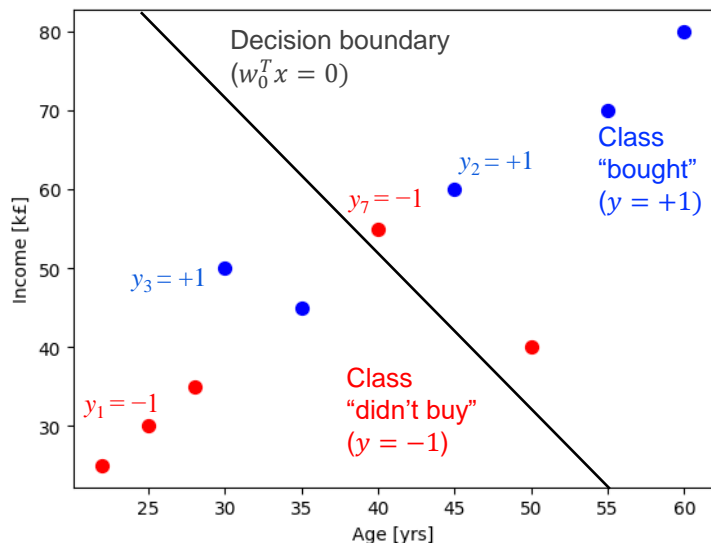
$$f(w, x) = \text{sign}(w^T x)$$

$$f: \mathbb{R}^D \rightarrow \{-1, +1\}$$

- Initial guess: $w_0 = [-130, 2, 1]^T$

Example: health insurance company

- Data on whether customers bought the plan
- Task: predict whether a new customer is likely to buy or not the plan, given their age and income.



- Goal: split the data into 2 **classes** (bought/didn't buy) that best match **class-labeled training data**.
- Classification model: linear classifier


$$f(w, x) = \text{sign}(w^T x)$$

$$f: \mathbb{R}^D \rightarrow \{-1, +1\}$$

- Initial guess: $w_0 = [-130, 2, 1]^T$

Example: health insurance company

- Data on whether customers bought the plan




- **Misclassification error:** number of misclassified data points

$$F(w) = \sum_{i=1}^N \mathbb{I}[f(w, x_i) \neq y_i]$$

- assigns the same penalty to all incorrect decisions, regardless of how 'bad' they are.

Example: health insurance company

Initial guess: $w_0 = [-130, 2, 1]^T$



- Now let's look at y_1

Client	Age (yrs)	Income (k £)	Bought?
1	25	30	No


What should be the loss here?

How about y_2 ?

2	45	60	Yes
---	----	----	-----

Example: health insurance company

Initial guess: $w_0 = [-130, 2, 1]^T$



- Now let's look at y_3

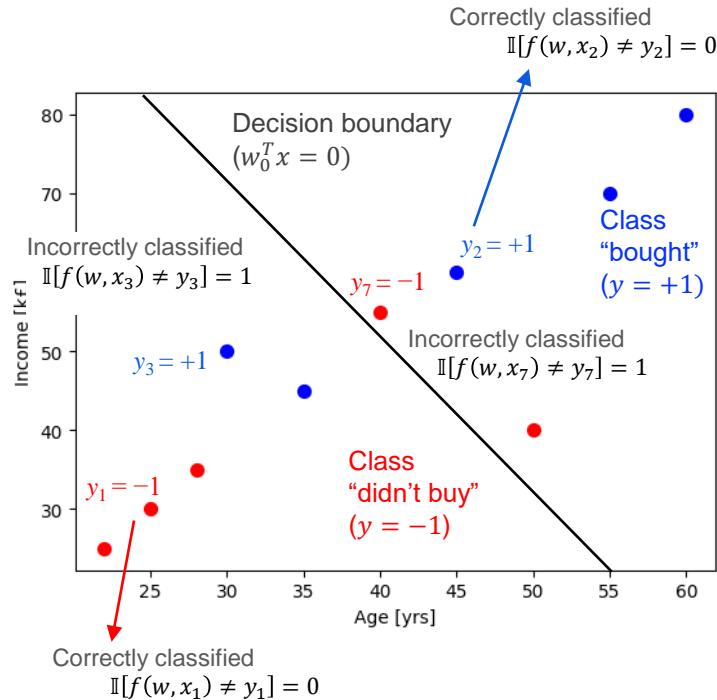
3	30	50	Yes
---	----	----	-----

- How about y_7 ?

7	40	55	No
---	----	----	----

Example: health insurance company

Initial guess: $w_0 = [-130, 2, 1]^T$



○ Now let's look at y_3


3	30	50	Yes
---	----	----	-----

○ How about y_7 ?

7	40	55	No
---	----	----	----

Example: health insurance company

- Data on whether customers bought the plan



- **Misclassification error**: number of misclassified data points

$$F(w) = \sum_{i=1}^N \mathbb{I}[f(w, x_i) \neq y_i]$$

- assigns the same penalty to all incorrect decisions, regardless of how 'bad' they are.

- **Perceptron error**: sum of perpendicular distances of every **misclassified** data point to the decision boundary,

$$F(w) = \sum_{i=1}^N \max (0, -y_i w^T x_i)$$

- 'Penalizes' incorrect decisions by the distance from the decision boundary $w^T x$ in the direction w (perpendicular distance).

SGD: algorithm (Section 9 Lecture Notes)

- **Step 1.** *Initialization:* Select an initial guess for w_0 , a convergence tolerance $\varepsilon > 0$, step size (learning rate) parameter $\alpha > 0$, set iteration number $n=0$
- **Step 2.** *Gradient descent step:* Compute new model parameters,

$$w_{n+1} = w_n - \alpha F_w(w_n)$$

- **Step 3.** *Convergence test:* Compute new loss function value $F(w_{n+1})$, and loss function improvement, $\Delta F = |F(w_{n+1}) - F(w_n)|$ and if $\Delta F < \varepsilon$, exit with solution $w^*=w_{n+1}$
- **Step 4.** *Iteration:* update $n=n+1$ and go to step 2.

Perceptron classification

- Classification model (D -dimensional):
$$f(w, x) = \text{sign}(w_1x^1 + w_2x^2 + \dots + w_Dx^D) = \text{sign}(w^T x)$$

Perceptron classification

- Classification model (D -dimensional):
$$f(w, x) = \text{sign}(w_1x^1 + w_2x^2 + \dots + w_Dx^D) = \text{sign}(w^T x)$$
- Perceptron error function:

$$F(w) = \sum_{i=1}^N \max(0, -y_i w^T x_i)$$

Perceptron classification

- Classification model (D -dimensional):

$$f(w, x) = \text{sign}(w_1 x^1 + w_2 x^2 + \dots + w_D x^D) = \text{sign}(w^T x)$$

- Perceptron error function:

$$F(w) = \sum_{i=1}^N \max(0, -y_i w^T x_i)$$

- Gradient with respect to w :

$$F_w(w) = - \sum_{i=1}^N y_i x_i \mathbb{I}[-y_i w^T x_i \geq 0]$$

- Intuitively, gradient is just sum of $-y_i x_i$ over incorrectly classified points

Perceptron training: algorithm

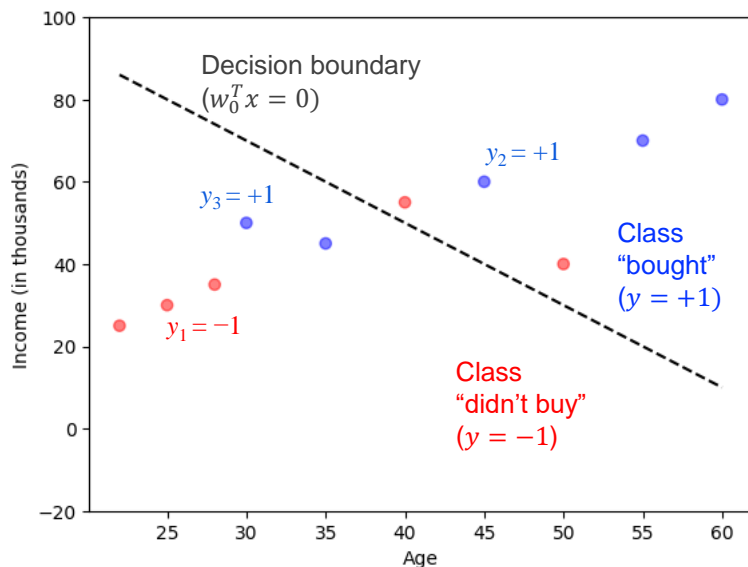
- **Step 1.** *Initialization:* Select a starting candidate classification model w_0 , set iteration number $n = 0$, choose maximum number of iterations R and learning rate $\alpha > 0$
- **Step 2.** *Gradient descent step:* Compute new model parameters: taking each $i = 1, 2, \dots, N$ in turn, if $\text{sign}(w_n^T x_i) \neq y_i$, then
$$w_{n+1} = w_n + \alpha y_i x_i$$
- **Step 3.** *Iteration:* If $n < R$, update $n = n + 1$, go to step 2, otherwise exit with solution $w^* = w_n$.

Perceptron training: algorithm

- **Step 1.** *Initialization:* Select a starting candidate classification model w_0 , set iteration number $n = 0$, choose maximum number of iterations R and learning rate $\alpha > 0$
- **Step 2.** *Gradient descent step:* Compute new model parameters: taking each $i = 1, 2, \dots, N$ in turn, if $\text{sign}(w_n^T x_i) \neq y_i$, then
$$w_{n+1} = w_n + \alpha y_i x_i$$
- **Step 3.** *Iteration:* If $n < R$, update $n = n + 1$, go to step 2, otherwise exit with solution $w^* = w_n$.

Example: health insurance company


- Perceptron algorithm in action



Client	Age (yrs)	Income (k £)	Bought?
1	25	30	No
2	45	60	Yes
3	30	50	Yes
4	22	25	No
5	35	45	Yes
6	55	70	Yes
7	40	55	No
8	60	80	Yes
9	50	40	No
10	28	35	No

Example: health insurance company

- Perceptron algorithm in action




○ $n = 0, w_0 = [-130, 2, 1]^T, R = 10, \alpha = 0.1:$

■ $i = 3, w_1 = w_0 + 0.1 \times (+1) \times \begin{bmatrix} 1 \\ 30 \\ 50 \end{bmatrix} = [-129.9, 5, 6]^T$

Example: health insurance company

- Perceptron algorithm in action




○ $n = 0, w_0 = [-130, 2, 1]^T, R = 10, \alpha = 0.1:$

■ $i = 3, w_1 = w_0 + 0.1 \times (+1) \times \begin{bmatrix} 1 \\ 30 \\ 50 \end{bmatrix} = [-129.9, 5, 6]^T$

■ $i = 4, w_1 = [-129.9, 5, 6]^T + 0.1 \times (-1) \times \begin{bmatrix} 1 \\ 22 \\ 25 \end{bmatrix} = [-130, 2.8, 3.5]^T$

Example: health insurance company

- Perceptron algorithm in action



○ $n = 0, w_0 = [-130, 2, 1]^T, R = 1000, \alpha = 0.1:$

■ $i = 3, w_1 = w_0 + 0.1 \times (+1) \times \begin{bmatrix} 1 \\ 30 \\ 50 \end{bmatrix} = [-129.9, 5, 6]^T$


■ $i = 4, w_1 = [-129.9, 5, 6]^T + 0.1 \times (-1) \times \begin{bmatrix} 1 \\ 22 \\ 25 \end{bmatrix} = [-130, 2.8, 3.5]^T$

■ $i = 7, w_1 = [-130, 2.8, 3.5]^T + 0.1 \times (+1) \times \begin{bmatrix} 1 \\ 40 \\ 55 \end{bmatrix} = [-130.1, -1.2, -2]^T$

....

Example: health insurance company

- Perceptron algorithm in action



- $n = 0, w_0 = [-130, 2, 1]^T, R = 1000, \alpha = 0.1:$

- $i = 3, w_1 = w_0 + 0.1 \times (+1) \times \begin{bmatrix} 1 \\ 30 \\ 50 \end{bmatrix} = [-129.9, 5, 6]^T$

- $i = 4, w_1 = [-129.9, 5, 6]^T + 0.1 \times (-1) \times \begin{bmatrix} 1 \\ 22 \\ 25 \end{bmatrix} = [-130, 2.8, 3.5]^T$

- $i = 7, w_1 = [-130, 2.8, 3.5]^T + 0.1 \times (+1) \times \begin{bmatrix} 1 \\ 40 \\ 55 \end{bmatrix} = [-130.1, -1.2, -2]^T$


....

- $n = 1, w_1 = [-130, -0.2, 2]^T:$

Update weights, and so on... until $n = R$

Example: health insurance company

- Perceptron algorithm in action



- $n = 0, w_0 = [-130, 2, 1]^T, R = 1000, \alpha = 0.1:$

- $i = 3, w_1 = w_0 + 0.1 \times (+1) \times \begin{bmatrix} 1 \\ 30 \\ 50 \end{bmatrix} = [-129.9, 5, 6]^T$

- $i = 4, w_1 = [-129.9, 5, 6]^T + 0.1 \times (-1) \times \begin{bmatrix} 1 \\ 22 \\ 25 \end{bmatrix} = [-130, 2.8, 3.5]^T$

- $i = 7, w_1 = [-130, 2.8, 3.5]^T + 0.1 \times (+1) \times \begin{bmatrix} 1 \\ 40 \\ 55 \end{bmatrix} = [-130.1, -1.2, -2]^T$

....

- $n = 1, w_1 = [-130, -0.2, 2]^T:$



Update weights, and so on... until $n = R$

- $n = 999, w^* = [-128.98, -1.85, 3.55]^T:$



Perceptron algorithm: analysis

- If the data is linearly separable, perceptron algorithm always converges on a decision boundary with zero error but no guarantee on the number of iterations required to reach fixed point
- If data is not linearly separable, no convergence guarantee – can cycle between local optima of the perceptron error function, so we need to stop after some number of iterations R



Perceptron training in action





Perceptron training in action





Perceptron training in action





Perceptron training in action





Perceptron training in action





Perceptron training in action





Perceptron training in action





Perceptron training in action



Perceptron training in action



Perceptron training in action



Some remarks about the perceptron

- Simple linear classifier based on the **perceptron error function** rather than the misclassification error function
- Very important classic algorithm in the history of ML, direct precursor to modern deep learning algorithms
- Extremely simple; there are mathematically better "linear single-layer" algorithms (e.g. **support vector machines**) so the perceptron is rarely used in practice today
- **Understanding the perceptron critical to understanding most of the main principles of modern ML classification**

To recap

- We learned the **perceptron algorithm** for classifying data.
 - It only converges if training data is linearly separable (and solution may not be unique)
- Question: how would you generalize the algorithm to $K > 2$ classes?
- **Next:** Neural networks (we will answer the question above)

Further Reading

- **PRML**, Section 4.1.7
- **R&N**, Section 18.6.3
- **H&T**, Section 4.5.1