

Artificial Intelligence and Machine Learning (AIML)


2023–24




#Code




- **Last lecture:**
 - Neural networks
 - Single-layer perceptron concept (“neuron unit”)



- **Last lecture:**
 - Neural networks
 - Single-layer perceptron concept (“neuron unit”)




- **Last lecture:**
 - Neural networks
 - Single-layer perceptron concept (“neuron unit”)






- **Last lecture:**
 - Neural networks
 - Single-layer perceptron concept (“neuron unit”)
 - Activation functions


ReLU
($\max(0, x)$)




Softplus
($\ln(1 + e^x)$)



Hyperbolic Tangent
($\tanh(x)$)




Logistic
($\frac{1}{1+e^{-x}}$)





- **Neural networks**

- Single-layer perceptron concept (“neuron unit”)
- Activation functions
- Perceptron extensions



$$\begin{aligned} z^1 &= f(w_{1,1}x^1 + w_{2,1}x^2 + \dots + w_{D,1}x^D) \\ &\vdots \\ z^M &= f(w_{1,M}x^1 + w_{2,M}x^2 + \dots + w_{D,M}x^D) \end{aligned}$$


[matrix notation]
 $z = f(W^T x)$

$$\hat{y} = f(v_1 z^1 + v_2 z^2 + \dots + v_M z^M)$$




- **Neural Networks**

- Single-layer perceptron concept (“neuron unit”)
- Activation functions
- Perceptron extensions and deep learning



Neural networks

- Single-layer perceptron concept (“neuron unit”)
- Activation functions
- Perceptron extensions and deep learning



Fully Connected Network:

Input \rightarrow Hidden Layer 1: $D \times M$ weights

HL1 \rightarrow HL 2: $M \times K$ weights


\vdots

Hidden Layer N: $K \times 1$ weights



- **Neural networks**

- Single-layer perceptron concept (“neuron unit”)
- Activation functions
- Perceptron extensions and deep learning
 - Weights problem & weight sharing (CNN)



$$z^1 = \max(0, w^T[x^1 \ x^2]^T)$$

$$z^2 = \max(0, w^T[x^2 \ x^3]^T)$$

...

$$z^M = \max(0, w^T[x^{D-1} \ x^D]^T)$$

$$y = \max(0, v^T z)$$



- **Neural networks**
 - Single-layer perceptron concept (“neuron unit”)
 - Activation functions
 - Perceptron extensions and deep learning
 - Weights problem & weight sharing (CNN)
- **Next:**

How to train deep neural networks

Training ML algorithms

TRAINING:

Iterative procedure for minimization of an error function, with adjustments to the weights being made at each step.


STAGE 1:

Evaluate derivatives of the error function wrt the weights

STAGE 2:

Use derivatives to compute adjustments to be made to the weights
(e.g., gradient descent)

Training a single layer linear perceptron using gradient descent



Error Function

Sum of Squares:


$$F(w) = \sum_{i=1}^N (w_1 x_i^1 - y_i)^2$$

Weight update:

Model:

$$f(w, x) = \hat{y} = w_1 x^1$$

Training a single layer linear perceptron using gradient descent



Error Function

Sum of Squares:

$$F(w) = \sum_{i=1}^N (w_1 x_i^1 - y_i)^2$$

Weight update:


$$w_{n+1} = w_n - \alpha F_w(w_n)$$

Model:

$$f(w, x) = \hat{y} = w_1 x^1$$

Training a single layer linear perceptron using gradient descent

Weight update: $w_{n+1} = w_n - \alpha F_w(w_n)$



Model:

$$f(w, x) = \hat{y} = w_1 x^1$$

Error Function

Sum of Squares:

$$F(w) = \sum_{i=1}^N (w_1 x_i^1 - y_i)^2$$

Derivative


$$F_w(w) = \frac{dF}{dw_1}$$

Sum of Squares:

$$F_w(w) = 2 \sum_{i=1}^N (w_1 x_i^1 - y_i) x_i^1$$

Training a single layer linear perceptron using gradient descent

Weight update: $w_{n+1} = w_n - \alpha F_w(w_n)$



Error Function

Sum of Squares:

$$F(w) = \sum_{i=1}^N ((w_1 x_i^1 + w_2 x_i^2) - y_i)^2$$

Model:

$$f(w, x) = \hat{y} = w_1 x^1 + w_2 x^2$$

Gradient


$$F_w(w) = \begin{bmatrix} \frac{\partial F}{\partial w_1} \\ \frac{\partial F}{\partial w_2} \end{bmatrix}$$

Sum of Squares:

$$F_w(w) = 2 \sum_{i=1}^N ((w_1 x_i^1 + w_2 x_i^2) - y_i) \begin{bmatrix} x_i^1 \\ x_i^2 \end{bmatrix}$$

Training a single layer linear perceptron using gradient descent

$$w_{n+1} = w_n - \alpha F_w(w_n)$$



Error Function

Sum of Squares:

$$F(w) = \sum_{i=1}^N (w^T x_i - y_i)^2$$

Model:

$$f(w, x) = \hat{y} = w_1 x^1 + w_2 x^2 + \dots + w_D x^D$$

(vector notation) $f(w, x) = \hat{y} = w^T x$

Gradient

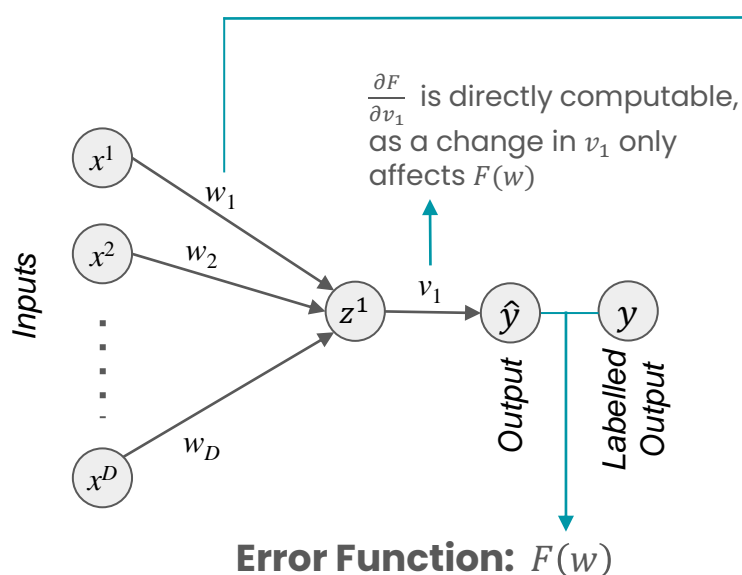
$$F_w(w) = \begin{bmatrix} \frac{\partial F}{\partial w_1} \\ \frac{\partial F}{\partial w_2} \\ \vdots \\ \frac{\partial F}{\partial w_D} \end{bmatrix}$$

$$F_w(w) = 2 \sum_{i=1}^N (w^T x_i - y_i) \begin{bmatrix} x_i^1 \\ x_i^2 \\ \vdots \\ x_i^D \end{bmatrix}$$

Neural networks: training by gradient descent

#Code

$$w_{n+1} = w_n - \alpha F_w(w_n)$$



We need $\frac{\partial F}{\partial w_1}$

A change in w_1 will affect the output of z^1 : $\frac{\partial z^1}{\partial w_1}$

The change in z^1 induced by w_1 will change the predicted output: $\frac{\partial F}{\partial z^1}$

The total change in the output due to a change in w_1 will then be:

$$\frac{\partial F}{\partial w_1} = \frac{\partial F}{\partial z^1} \frac{\partial z^1}{\partial w_1} \quad (\text{chain rule})$$

Gradient

$$F_w(w) = \begin{bmatrix} \frac{\partial F}{\partial w_1} \\ \frac{\partial F}{\partial w_2} \\ \vdots \\ \frac{\partial F}{\partial w_D} \\ \frac{\partial F}{\partial v_1} \end{bmatrix}$$

Chain rule of Calculus

(Gill, Section 5.4)

- Provides a means of differentiating nested functions.
- Given two functions, $f(x)$ and $g(x)$, and the nested form

$$h = f(g(x))$$

we need to account for the actual order of the nesting relationship to correctly differentiate h :

Chain rule of Calculus

(Gill, Section 5.4)

- Provides a means of differentiating nested functions.
- Given two functions, $f(x)$ and $g(x)$, and the nested form

$$h = f(g(x))$$

we need to account for the actual order of the nesting relationship to correctly differentiate h :


$$\frac{dh}{dx} = \frac{d}{dx}f(g(x)) = f'(g(x))g'(x)$$

where $f'(x)$ is shorthand for $\frac{df(x)}{dx}$.

Neural networks: Training using gradient descent ^{#Code}

- Gradient Descent Approach

$$w_{n+1} = w_n - \alpha F_w(w_n)$$




Gradient

$$F_w(w) = \begin{bmatrix} \frac{\partial F}{\partial w_1} \\ \frac{\partial F}{\partial w_2} \\ \vdots \\ \frac{\partial F}{\partial w_D} \\ \frac{\partial F}{\partial v_1} \end{bmatrix}$$

Neural networks: Training using gradient descent ^{#Code}

- Gradient Descent Approach

$$w_{n+1} = w_n - \alpha F_w(w_n)$$



We need $\frac{\partial F}{\partial w_1}$

A change in w_1 will affect the output of z^1 : $\frac{\partial z^1}{\partial w_1}$


Gradient

$$F_w(w) = \begin{bmatrix} \frac{\partial F}{\partial w_1} \\ \frac{\partial F}{\partial w_2} \\ \vdots \\ \frac{\partial F}{\partial w_D} \\ \frac{\partial F}{\partial v_1} \end{bmatrix}$$

Neural networks: Training using gradient descent ^{#Code}

- Gradient Descent Approach

$$w_{n+1} = w_n - \alpha F_w(w_n)$$



We need $\frac{\partial F}{\partial w_1}$

A change in w_1 will affect the output of z^1 : $\frac{\partial z^1}{\partial w_1}$

The change in z^1 induced by w_1 will change the predicted output: $\frac{\partial F}{\partial z^1}$

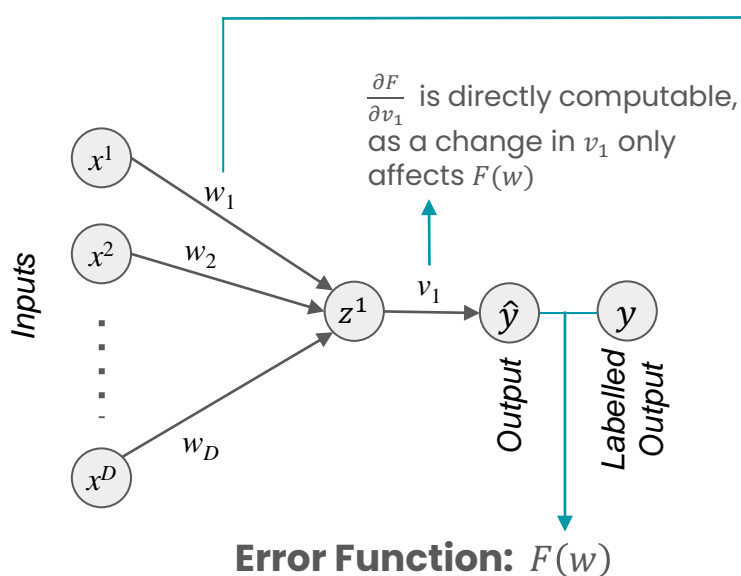
Gradient

$$F_w(w) = \begin{bmatrix} \frac{\partial F}{\partial w_1} \\ \frac{\partial F}{\partial w_2} \\ \vdots \\ \frac{\partial F}{\partial w_D} \\ \frac{\partial F}{\partial v_1} \end{bmatrix}$$

Neural networks: Training using gradient descent

- Gradient Descent Approach

$$w_{n+1} = w_n - \alpha F_w(w_n)$$



We need $\frac{\partial F}{\partial w_1}$

A change in w_1 will affect the output of z^1 : $\frac{\partial z^1}{\partial w_1}$

The change in z^1 induced by w_1 will change the predicted output: $\frac{\partial F}{\partial z^1}$

The total change in the output due to a change in w_1 will then be:

$$\frac{\partial F}{\partial w_1} = \frac{\partial F}{\partial z^1} \frac{\partial z^1}{\partial w_1} \quad \text{(chain rule)}$$

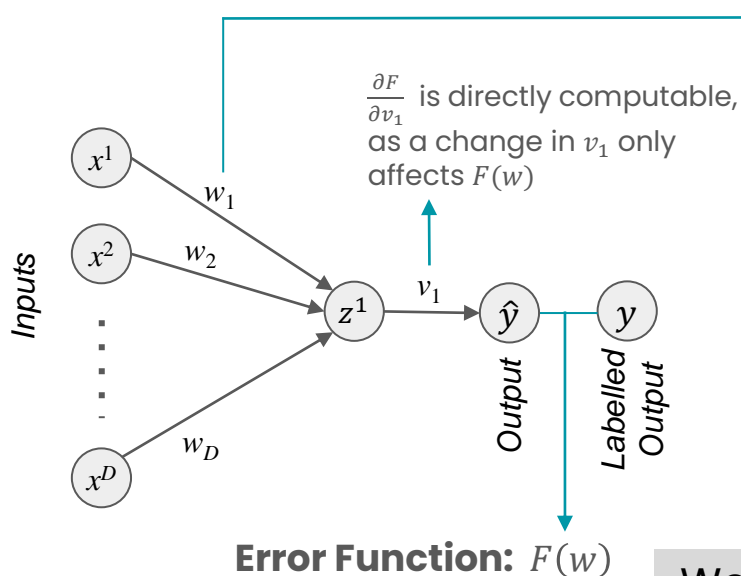
Gradient

$$F_w(w) = \begin{bmatrix} \frac{\partial F}{\partial w_1} \\ \frac{\partial F}{\partial w_2} \\ \vdots \\ \frac{\partial F}{\partial w_D} \end{bmatrix}$$

Neural networks: Training using gradient descent

- Gradient Descent Approach

$$w_{n+1} = w_n - \alpha F_w(w_n)$$



We need $\frac{\partial F}{\partial w_1}$

A change in w_1 will affect the output of z^1 : $\frac{\partial z^1}{\partial w_1}$

The change in z^1 induced by w_1 will change the predicted output: $\frac{\partial F}{\partial z^1}$

The total change in the output due to a change in w_1 will then be:

$$\frac{\partial F}{\partial w_1} = \frac{\partial F}{\partial z^1} \frac{\partial z^1}{\partial w_1} \quad \text{(chain rule)}$$

Gradient

$$F_w(w) = \begin{bmatrix} \frac{\partial F}{\partial w_1} \\ \frac{\partial F}{\partial w_2} \\ \vdots \\ \frac{\partial F}{\partial w_D} \end{bmatrix}$$

We must **propagate gradients** from the output error $F(w)$, all the way through each layer

Training deep networks

TRAINING:

Iterative procedure for minimization of an error function, with adjustments to the weights being made at each step.

STAGE 1:

Evaluate derivatives of the error function wrt the weights

Analytical gradient expressions quickly become intractable

STAGE 2:

Use derivatives to compute adjustments to be made to the weights

Training deep networks

TRAINING:

Iterative procedure for minimization of an error function, with adjustments to the weights being made at each step.

STAGE 1:

Evaluate derivatives of the error function wrt the weights

Analytical gradient expressions quickly become intractable


More sophisticated gradient calculation methods:
backpropagation and **automatic differentiation (AD)**

STAGE 2:

Use derivatives to compute adjustments to be made to the weights
(e.g., gradient descent)

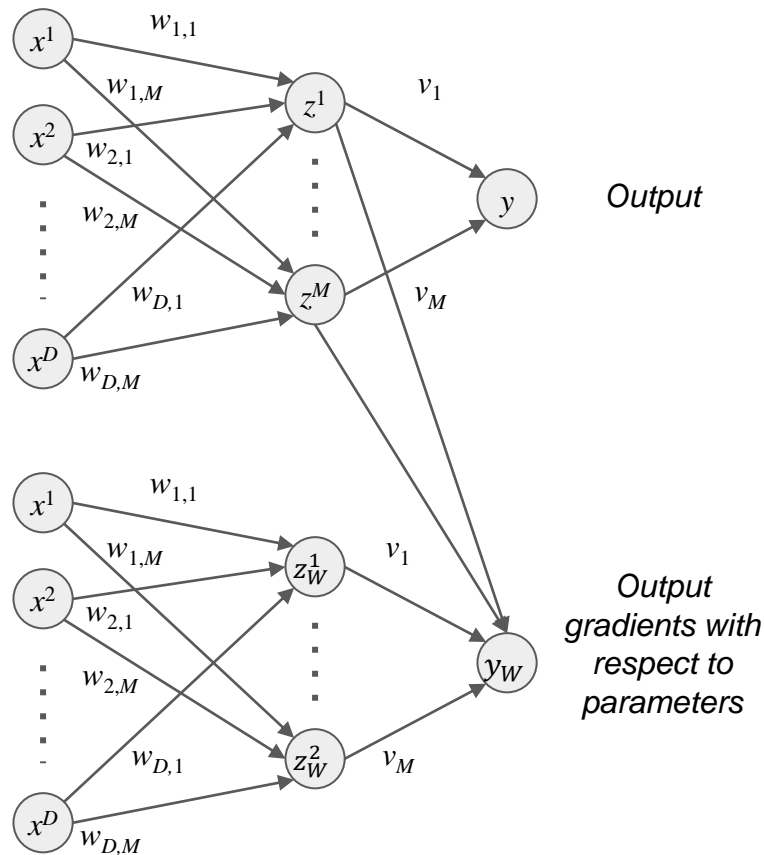
Automatic differentiation

- **"meta-programming"**
approach to gradient calculation
- Obtains the gradients of the output simultaneous with the output of the network



Automatic differentiation

- **"meta-programming"** approach to gradient calculation
- Obtains the gradients of the output simultaneous with the output of the network
- Software packages such as PyTorch, JAX largely avoid the need for any hand-computed gradients in this way



Algebra of dual numbers

- Dual numbers are an easy way to handle chained computations and automate calculations for AD
- Keeps track of current computation's value u and its derivative u' as a pair:
 (u, u')

Algebra of dual numbers

- Dual numbers are an easy way to handle chained computations and automate calculations for AD
- Keeps track of current computation's value u and its derivative u' as a pair:

$$(u, u')$$

- In dual number form, the general chain rule is,

$$f((u, u')) = (f(u), f'(u)u')$$

- Most forms of computational operations used in ML, are special cases of this rule

Algebra of dual numbers

Computations which commonly occur in deep learning, in dual number form

- **Addition:** $f(u, v) = u + v$

Example: Suppose $u = 2x^1$ and $v = x^1 + x^2$.

For $x^1 = 3, x^2 = 4$, we want to compute

1. $z = u + v$

2. $\frac{dz}{dx^1}$

Algebra of dual numbers

Computations which commonly occur in deep learning, in dual number form

- **Addition:** $f(u, v) = u + v$

Example: Suppose $u = 2x^1$ and $v = x^1 + x^2$.

For $x^1 = 3, x^2 = 4$, we want to compute

$$1. z = u + v$$

$$z = 6 + 7$$

$$u = 2 \times 3 = 6$$

$$v = 3 + 4 = 7$$

$$2. \frac{dz}{dx^1} = \frac{d}{dx^1}(u + v) = \frac{du}{dx^1} + \frac{dv}{dx^1} = 2 + 1 = 3$$

Algebra of dual numbers

Computations which commonly occur in deep learning, in dual number form

- **Addition:** $f(u, v) = u + v$

$$(u, u') + (v, v') = (u + v, u' + v')$$

Example: Suppose $u = 2x^1$ and $v = x^1 + x^2$.

For $x^1 = 3, x^2 = 4$, we want to compute

$$1. z = u + v$$

$$z = 6 + 7$$

$$u = 2 \times 3 = 6$$

$$v = 3 + 4 = 7$$

$$2. \frac{dz}{dx^1} = \frac{d}{dx^1}(u + v) = \frac{du}{dx^1} + \frac{dv}{dx^1} = 2 + 1 = 3$$

Algebra of dual numbers

Computations which commonly occur in deep learning, in dual number form

- **Multiplication:** $f(u, v) = uv$

Example: Suppose $u = 2x^1$ and $v = x^1 + x^2$.

For $x^1 = 3, x^2 = 4$, we want to compute

1. $z = uv$

2. $\frac{dz}{dx^1}$

Algebra of dual numbers

Computations which commonly occur in deep learning, in dual number form

- **Multiplication:** $f(u, v) = uv$

Example: Suppose $u = 2x^1$ and $v = x^1 + x^2$.

For $x^1 = 3, x^2 = 4$, we want to compute

1. $z = uv$

$$u = 2 \times 3 = 6$$

$$v = 3 + 4 = 7$$

$$z = 6 \times 7 = 42$$

2. $\frac{dz}{dx^1} = \frac{d}{dx^1}(uv) = \frac{du}{dx^1}v + u\frac{dv}{dx^1} = 2 \times 7 + 6 \times 1 = 20$

Algebra of dual numbers

Computations which commonly occur in deep learning, in dual number form

- **Multiplication:** $f(u, v) = uv$
 $(u, u') \times (v, v') = (uv, u'v + uv')$

Example: Suppose $u = 2x^1$ and $v = x^1 + x^2$.

For $x^1 = 3, x^2 = 4$, we want to compute

1. $z = uv$

$$u = 2 \times 3 = 6$$

$$v = 3 + 4 = 7$$

$$z = 6 \times 7 = 42$$

$$2. \frac{dz}{dx^1} = \frac{d}{dx^1}(uv) = \frac{du}{dx^1}v + u \frac{dv}{dx^1} = 2 \times 7 + 6 \times 1 = 20$$

Algebra of dual numbers

Computations which commonly occur in deep learning, in dual number form

- **Addition:** $(u, u') + (v, v') = (u + v, u' + v')$
- **Multiplication:** $(u, u') \times (v, v') = (uv, u'v + uv')$
- **ReLU activation:** $\text{relu}((u, u')) = (\max(0, u), u' \mathbb{I}[u \geq 0])$
- **Constants** $f(u) = c$

$$f((u, u')) = (c, 0)$$


Here is an example of a chained calculation carried out using dual numbers. Given the constants $y = 3$ and $z = -1$ and variable $x = 2$, compute $u(x, y, z) = \max(yz, y + 2x)$ and its derivative, $u_x(x, y, z)$. Applying the rules above successively (and using additional symbols for intermediate computational results),

$$\begin{aligned}
 \bar{x} &= (2, 1) \\
 \bar{y} &= (3, 0) \\
 \bar{z} &= (-1, 0) \\
 c &= (2, 0) \\
 c\bar{x} &= (2, 0) \times (2, 1) = (4, 2) \\
 r_1 &= \bar{y} \times \bar{z} = (3, 0) \times (-1, 0) = (-3, 0) \\
 r_2 &= \bar{y} + c\bar{x} = (3, 0) + (4, 2) = (7, 2) \\
 \bar{u} &= \max((-3, 0), (7, 2)) = (7, 2),
 \end{aligned} \tag{14.7}$$

therefore $u(x, y, z) = 7$ and $u_x(x, y, z) = 2$. While it is, of course, always possible to find the symbolic derivative of the function $u(x, y, z)$, AD enables entirely ‘mechanical’ calculational steps which lends itself to software implementation.


Automatic differentiation in action

#Code



$$z^1 = \text{relu}(w_{1,1}x^1 + w_{2,1}x^2)$$


Automatic differentiation in action



$$z^1 = \text{relu}(w_{1,1}x^1 + w_{2,1}x^2)$$

$$w_{1,1} \rightarrow (w_{1,1}, 1), w_{2,1} \rightarrow (w_{2,1}, 0), x^1 \rightarrow (x^1, 0), x^2 \rightarrow (x^2, 0)$$

Automatic differentiation in action




$$z^1 = \text{relu}(w_{1,1}x^1 + w_{2,1}x^2)$$

$$= \text{relu}((w_{1,1}, 1) \times (x^1, 0) + (w_{2,1}, 0) \times (x^2, 0))$$

$$w_{1,1} \rightarrow (w_{1,1}, 1), w_{2,1} \rightarrow (w_{2,1}, 0), x^1 \rightarrow (x^1, 0), x^2 \rightarrow (x^2, 0)$$

Automatic differentiation in action



$$z^1 = \text{relu}(w_{1,1}x^1 + w_{2,1}x^2)$$


$$= \text{relu}((w_{1,1}, 1) \times (x^1, 0) + (w_{2,1}, 0) \times (x^2, 0))$$

$$= \text{relu}((w_{1,1}x^1, x^1) + (w_{2,1}x^2, 0))$$

$$w_{1,1} \rightarrow (w_{1,1}, 1), w_{2,1} \rightarrow (w_{2,1}, 0), x^1 \rightarrow (x^1, 0), x^2 \rightarrow (x^2, 0)$$

$$(u, u') \times (v, v') = (uv, u'v + v'u)$$

Automatic differentiation in action



$$z^1 = \text{relu}(w_{1,1}x^1 + w_{2,1}x^2)$$

$$= \text{relu}((w_{1,1}, 1) \times (x^1, 0) + (w_{2,1}, 0) \times (x^2, 0))$$

$$= \text{relu}((w_{1,1}x^1, x^1) + (w_{2,1}x^2, 0))$$


$$= \text{relu}((w_{1,1}x^1 + w_{2,1}x^2, x^1))$$

$$w_{1,1} \rightarrow (w_{1,1}, 1), w_{2,1} \rightarrow (w_{2,1}, 0), x^1 \rightarrow (x^1, 0), x^2 \rightarrow (x^2, 0)$$

$$(u, u') \times (v, v') = (uv, u'v + v'u)$$

$$(u, u') + (v, v') = (u + v, u' + v')$$

Automatic differentiation in action



$$z^1 = \text{relu}(w_{1,1}x^1 + w_{2,1}x^2)$$

$$= \text{relu}((w_{1,1}, 1) \times (x^1, 0) + (w_{2,1}, 0) \times (x^2, 0))$$

$$= \text{relu}((w_{1,1}x^1, x^1) + (w_{2,1}x^2, 0))$$

$$= \text{relu}((w_{1,1}x^1 + w_{2,1}x^2, x^1))$$

$$= (\max(0, w_{1,1}x^1 + w_{2,1}x^2), x^1 \mathbf{1}[w_{1,1}x^1 + w_{2,1}x^2])$$

$$w_{1,1} \rightarrow (w_{1,1}, 1), w_{2,1} \rightarrow (w_{2,1}, 0), x^1 \rightarrow (x^1, 0), x^2 \rightarrow (x^2, 0)$$

$$(u, u') \times (v, v') = (uv, u'v + v'u)$$

$$(u, u') + (v, v') = (u + v, u' + v')$$

$$\text{relu}((u, u')) = (\max(0, u), u' \mathbf{1}[u \geq 0])$$

To recap

- We learned one approach of finding the optimal set of weights for a neural network (i.e., **train the network**)
 - Analytical gradient expressions quickly become intractable
 - Automatic differentiation computes derivatives in a single forward pass
 - We learned the *forward mode* version of AD; there is also *reverse mode*
- **Next:** Probability and probabilistic AI

Further Reading

- **PRML**, Section 5.3
- **R&N**, Section 18.7
- **H&T**, Section 11.4, Section 11.7