

# Week 4

# Databases – Conceptual Design

LM Data Structures, Algorithms, and Databases  
(34141)

Dr Ahmad Ibrahim  
[a.ibrahim@bham.ac.uk](mailto:a.ibrahim@bham.ac.uk)  
Feb 05, 2024



UNIVERSITY OF  
BIRMINGHAM | DUBAI  
دبي

**Content based on lecture notes from Uday Reddy & Achim Jung**

# Topics by Week



Week	Date	Topic
1	15 Jan	Searching algorithms
2	22 Jan	Binary Search Tree
3	29 Jan	Balancing Trees – AVL Tree
4	5 Feb	Databases – Conceptual Design
5	14 Feb	Databases – Logical Design & Relational Algebra
6	19 Feb	Consolidation Week
7	26 Feb	Graph Algorithms
8	4 Mar	Sorting Algorithms
9	11 Mar	Hash tables
10	18 Mar	Databases – Normalization
		Easter break and Eid break
11	22 Apr	Databases – Concurrency
12	29 Apr	Revision Week

# Timetable & Office hours

Day	Time	Event	Location
Monday	6:00-7:00pm	Online support session*	Online*
Tuesday	6:00-7:00pm	Office hour 1 (by appointment)*	Online*
Wednesday	-	-	-
Thursday	6:00-7:00pm	Office hour 2 (by appointment)*	Online*
Friday	6:00-8:00pm	Lecture	Auditorium
	8:00-9:00pm	Tutorial	Auditorium



UNIVERSITY OF  
BIRMINGHAM | DUBAI  
دبي

**\*Zoom link:** <https://bham-ac-uk.zoom.us/j/81310444523?pwd=T01tZlZGdmdUL2lkeHZsVFpjcWxUUT09>

# Assessments

Assessments (Test 1, Test 2, Test 3): **20%**  
Exam: **80%**

## Late Submission Policy:

Submissions between 8:00-8:30pm (Dubai Time) on 08 Feb incur a 10% penalty.

**Zero marks for submissions after 8:30pm.**

Wellbeing approved cases: 1 day extension only.



UNIVERSITY OF  
BIRMINGHAM | DUBAI  
دبي

## ▼ Upcoming assignments



### Test 1

Week 4

Not available until 7 Feb at 20:00 | Due 8 Feb at 20:00 | -/20 pts



### Test 2

Week 7

Not available until 28 Feb at 20:00 | Due 29 Feb at 20:00 | -/20 pts



### Test 3

Week 10

Not available until 20 Mar at 20:00 | Due 21 Mar at 20:00 | -/20 pts

Review of Block-2

- Entities Relations Modelling
- Relational algebra
- SQL
- E-R models into table designs
- Examples

Week	Date	Topic
1	15 Jan	Searching algorithms
2	22 Jan	Binary Search Tree
3	29 Jan	Balancing Trees – AVL Tree
4	5 Feb	Databases – Conceptual Design
5	14 Feb	Databases –Logical Design, Relational Algebra
6	19 Feb	Consolidation Week








Select \* from book  
where ...

# This Week

## Introduction

- Relational databases
- Entity-relationship (ER) diagram /Modeling
- Table design and creation, SQL
- Weak entities
- Hierarchies
- Table design (schemas) Optimisation
- SQL Commands + PostgreSQL





# What is a **database**?

It is a **large** collection of **persistent** data

Typically stored on a **server** somewhere on the net

Accessible from multiple applications on client computers




# What is a **database**?

It is a **large** collection of **persistent** data

Typically stored on a **server** somewhere on the net

Accessible from multiple applications on client computers

➡ **Concurrently** accessible and modifiable



# What is a **database**?

It is a **large** collection of **persistent** data

Typically stored on a **server** somewhere on the net

Accessible from multiple applications on client computers

**Concurrently** accessible and modifiable



Expected to be **secure**.



# What is a **database**?

It is a **large** collection of **persistent** data

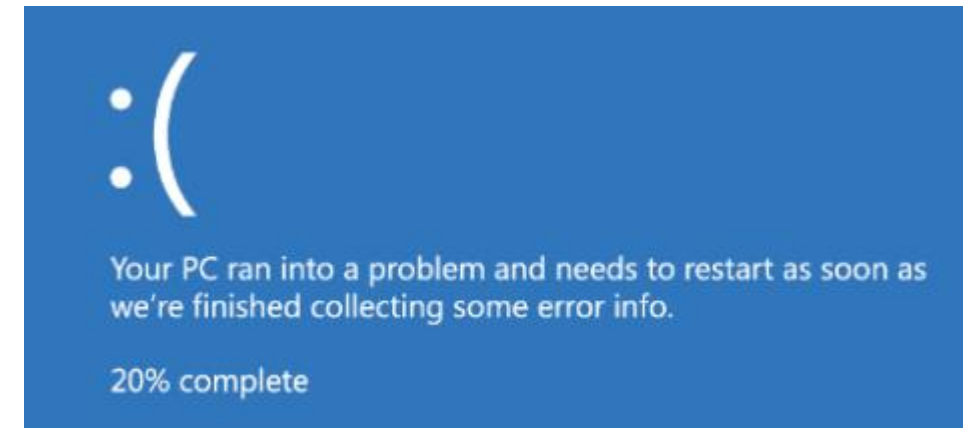

Typically stored on a **server** somewhere on the net

Accessible from multiple applications on client computers

**Concurrently** accessible and modifiable

Expected to be **secure**.

➔ Expected to be **fault-tolerant** (can recover from crashes). No data losses!



# What is a **database**?

It is a **large** collection of **persistent** data

Typically stored on a **server** somewhere on the net

Accessible from multiple applications on client computers


**Concurrently** accessible and modifiable

Expected to be **secure**.

Expected to be **fault-tolerant** (can recover from crashes). No data losses!

➔ Expected to be **efficient**

They were also called **information system**  
(old term, obsolete)



# Why “databases”?

Meant to suggest the “base of data” on which all the applications run.

Typically meant for the internal data of organisations/businesses, but may also be used services provided by the organisation.

## Examples:

University: students, courses, marks, staff, resources, finance

Shop: sales-items, customers, store, sales, staff, finance

Library: books, customers, borrowings, stocks, publishers, staff

Airline: airplanes, parts, airports, flights, customers, travel agents, reservations, staff, finance

Manufacturing company: products, parts, stores, factories, customers, finance



# Brief history

Stage 1: Sequential access files (tape drives)



# Brief history

Stage 1: Sequential access files (tape drives)

Stage 2: Random-access files (disk drives)




# Brief history

Stage 1: Sequential access files (tape drives)

Stage 2: Random-access files (disk drives)

Stage 3: Hierarchically-structured or networked file structures.






# Brief history

Stage 1: Sequential access files (tape drives)

Stage 2: Random-access files (disk drives)

Stage 3: Hierarchically-structured or networked file structures.

Stage 4: Relational databases.



# Brief history

Stage 1: Sequential access files (tape drives)

Stage 2: Random-access files (disk drives)

Stage 3: Hierarchically-structured or networked file structures.

Stage 4: Relational databases.

Stage 5: XML databases/documents

```
<credit>NOAA's National Weather Service</credit>
<credit_URL>https://weather.gov/</credit_URL>
<image>
  <url>https://weather.gov/images/xml_logo.gif</url>
  <title>NOAA's National Weather Service</title>
  <link>https://www.weather.gov</link>
</image>
<suggested_pickup>15 minutes after the hour</suggested_pickup>
<suggested_pickup_period>60</suggested_pickup_period>
<location>Birmingham, Birmingham International Airport, AL</location>
<station_id>KBHM</station_id>
<latitude>33.56556</latitude>
<longitude>-86.745</longitude>
<observation_time>Last Updated on Feb 7 2024, 8:53 am CST</observation_time>
  <observation_time_rfc822>Wed, 07 Feb 2024 08:53:00 -0600</observation_time_rfc822>
<weather>A Few Clouds</weather>
<temperature_string>53.0 F (11.7 C)</temperature_string>
```

# Brief history

Stage 1: Sequential access files (tape drives)

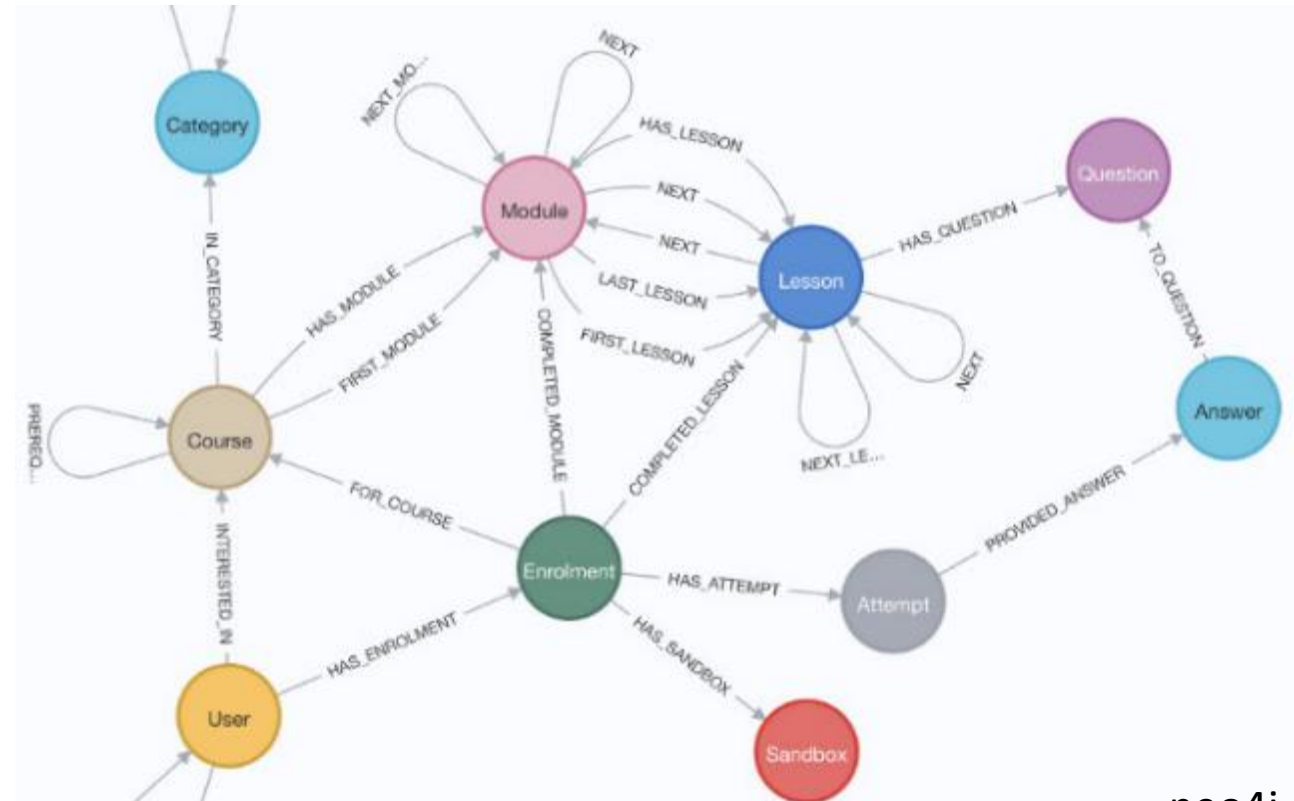
Stage 2: Random-access files (disk drives)

Stage 3: Hierarchically-structured or networked file structures.

Stage 4: Relational databases.

Stage 5: XML databases/documents

Stage 6: No-SQL (Graph Database)



Analytics And Data Science

# **Bad Data Costs the U.S. \$3 Trillion Per Year**

by Thomas C. Redman

September 22, 2016



<https://hbr.org/2016/09/bad-data-costs-the-u-s-3-trillion-per-year>

# This Week

- Introduction
- ➔ • Relational databases
- Entity-relationship (ER) diagram /Modeling
- Table design and creation, SQL
- Weak entities
- Hierarchies
- Table design (schemas) Optimisation
- SQL Commands + PostgreSQL



# Relational databases - history

Invented by Edgar Codd, at IBM Toronto research centre, 1970.

**Key idea:** Most of the data is highly “regular”, i.e., many items of the same kind, e.g., **students** in a University, **products** in a manufacturing.

The regularity allows us to provide a **simplified conceptual view**, making it easy to access/modify the data.


Built as a **relational database management system** (DBMS or RDBMS)

Providers:

IBM, Ingres, Oracle

Microsoft, **Postgres**, MySQL

# Relational Databases – Physical Structure



# Tables (“Relations”)

The textbook uses the term “Relations” a lot.

First name	Last name	Office	Phone number	Email address
Uday	Reddy	210	43740	-
Ahmad	Ibrahim	Dubai Campus	-	-
Mirco	Giacobbe	208	-	-
Jizheng	Wan	-	-	-

Rows  
(Records, tuples)

Columns  
(Attributes, fields)

The entries marked “-” are called “null” entries

Note that all the rows have the same structure.  
But the columns can be quite different from each other.



# This Week

- Introduction
- Relational databases
- ➔ Entity-relationship (ER) diagram /Modeling
- Table design and creation, SQL
- Weak entities
- Hierarchies
- Table design (schemas) Optimisation
- SQL Commands + PostgreSQL



# Entity-relationship (ER) diagram

We start some kind of a requirements description of the problem.



From it, we need to identify:

entities

attributes

relationships between the entities



Entity vs. attribute decision

Entities are “things”

Entities should have their own attributes!



Entity vs. relationship decision

Sometimes, nouns are used for relationships as well! e.g., “parent”, “supervisor”, “rental”, “payment”, etc.

# Requirements description of the problem:

## Library management system

*We would like to build a library management system that can keep a record of books, copies of the books, and the currently borrowed books.*

*Each library member can borrow only one book at a time for a certain duration. A late penalty/fine will be applied in case of a late return.*



# Entities - examples

## People (roles)

Student, Lecturer, Staff member, engineer, doctor, patient, customer, passenger, ...

## Objects

Cars, airplanes, products, parts, sales items, offices, buildings, ...

## Organisations

Companies, suppliers, departments, clubs, committees, ...

## Conceptual

Course, degree programme, project, design, exhibition, ...

## Events

Course deliveries, lectures, exams, concerts, sales, ...

# Library management system: **entities**

*We would like to build a **library management system** that can keep a record of **books**, copies of the books, and the currently borrowed books.*

*Each library **member** can borrow only one book at a time for a certain duration. A late penalty/fine will be applied in case of a late return.*



# Example – Library management system



Book



Member

We depict entities in our design by rectangles.

# Library management system: **Attributes** for entities



- **Book**

- Author(s)
- Title
- Publisher
- Year of publication
- Edition

- **Member**

- Membership number
- Name (first name, last name)
- Address
- Date of joining

# Example – Library management system





# Relationships between entities

It is these relationships that tie the tables together and are the key to doing database design *correctly*.

## For example:

- Students *register for* courses.
- Students *sit for* exams (and receive marks).
- These exams *assess* particular courses.
- Courses *constitute* degree programmes.
- Lecturers *teach* courses.
- Lecturers *set* exams (and mark them).
- Human relationships
  - Spouse, offspring, parent, manager, client, ...
- Role relationships
  - Teaching, studying, supervising (a department), managing (a project), selling, buying, borrowing, ...
- Organic relationships
  - Belongs to, Part of, Located at, ...
- Event occurrences
  - Lecturing at, concert at, ...

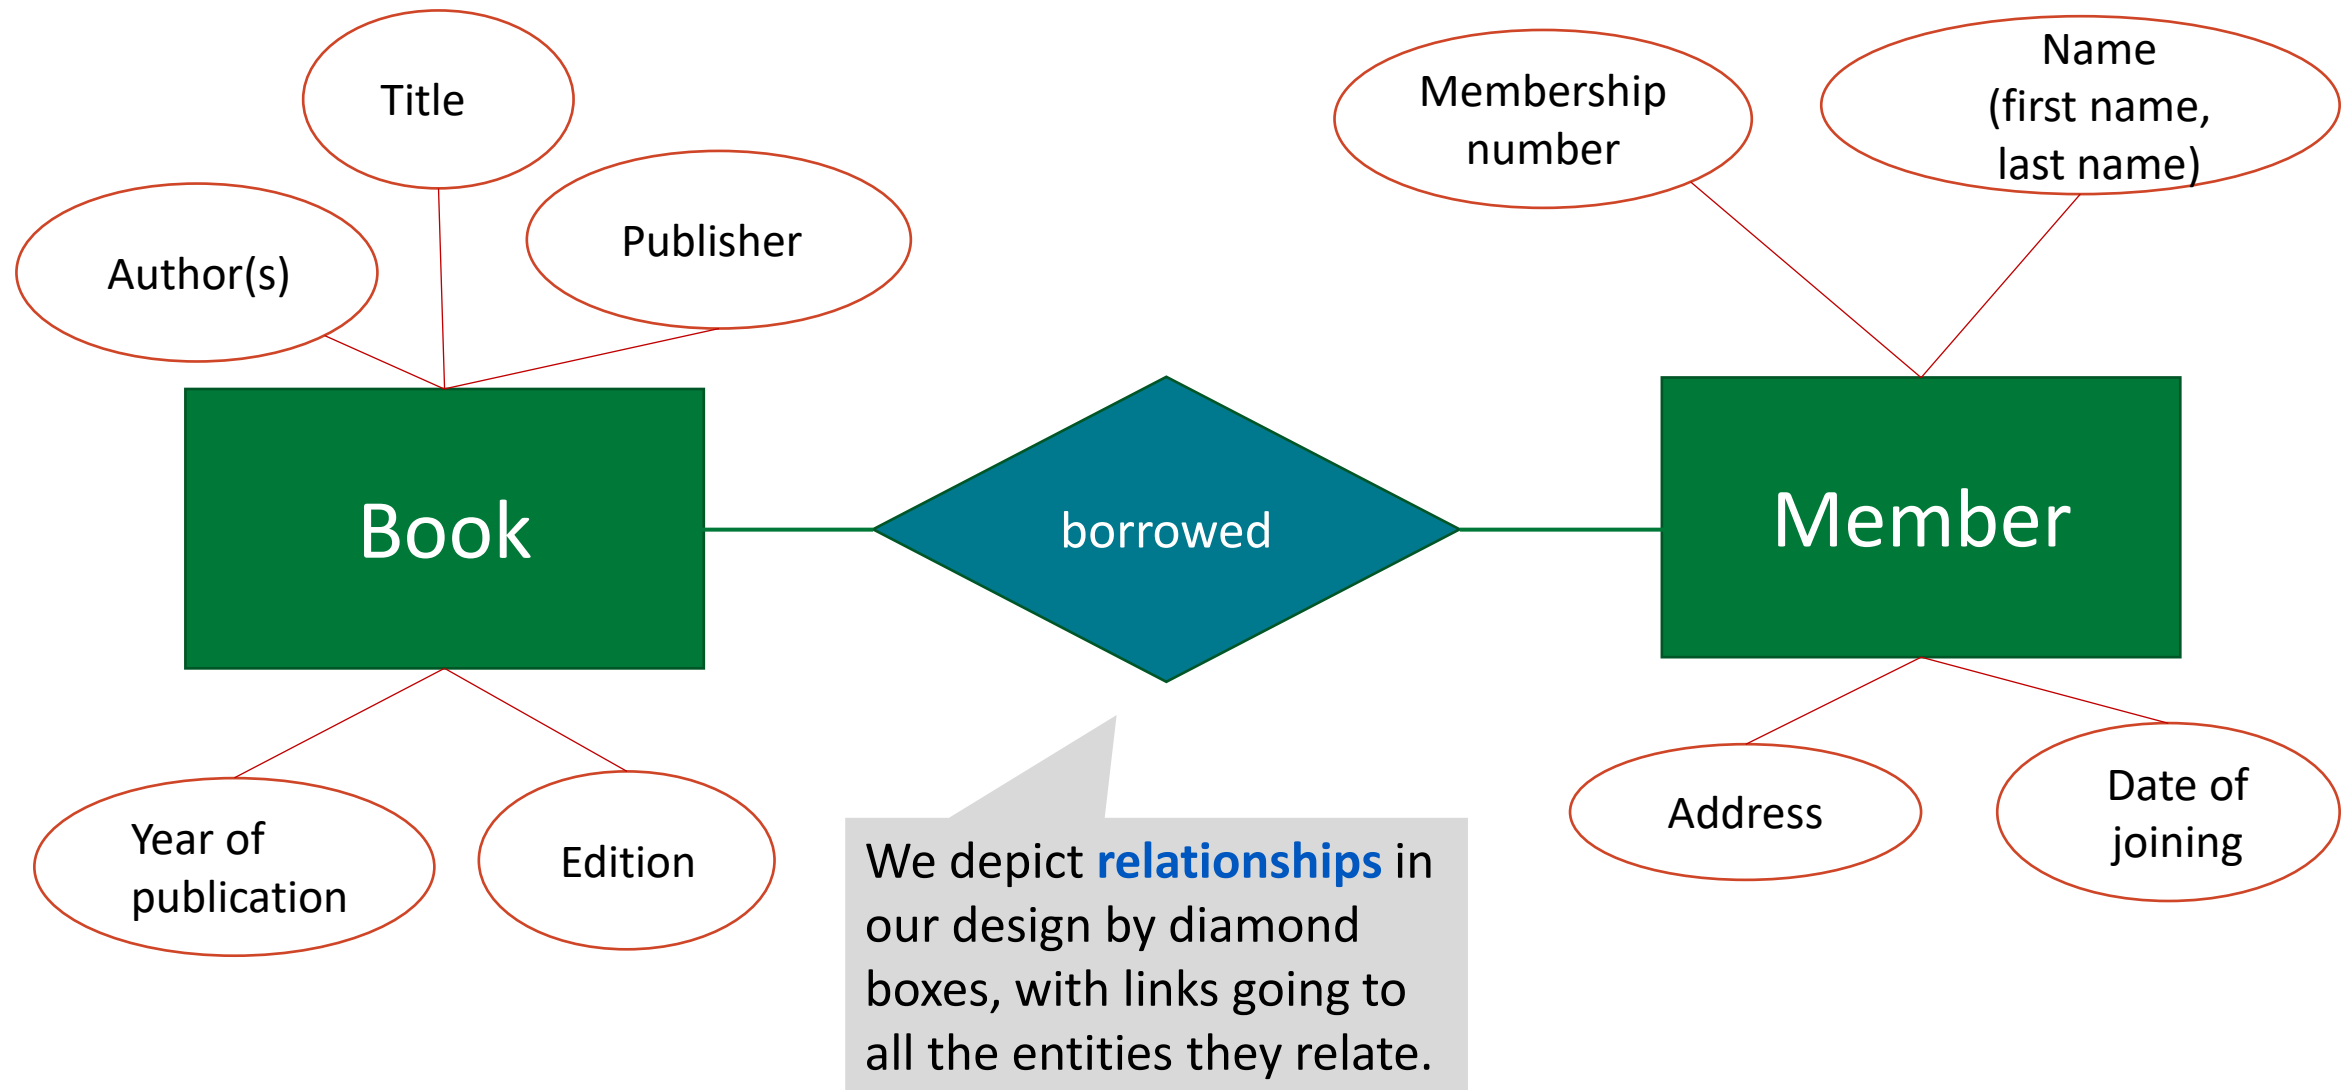
# Library management system: Relationships

*We would like to build a **library management system** that can keep a record of books, copies of the books, and the currently borrowed books.*

*Each library member can **borrow** only one book at a time for a certain duration. A late penalty/fine will be applied in case of a late return.*



# Example – Library management system



# Attributes (for both entities and relationships)

- Book

- Author(s)
- Title
- Publisher
- Year of publication
- Edition

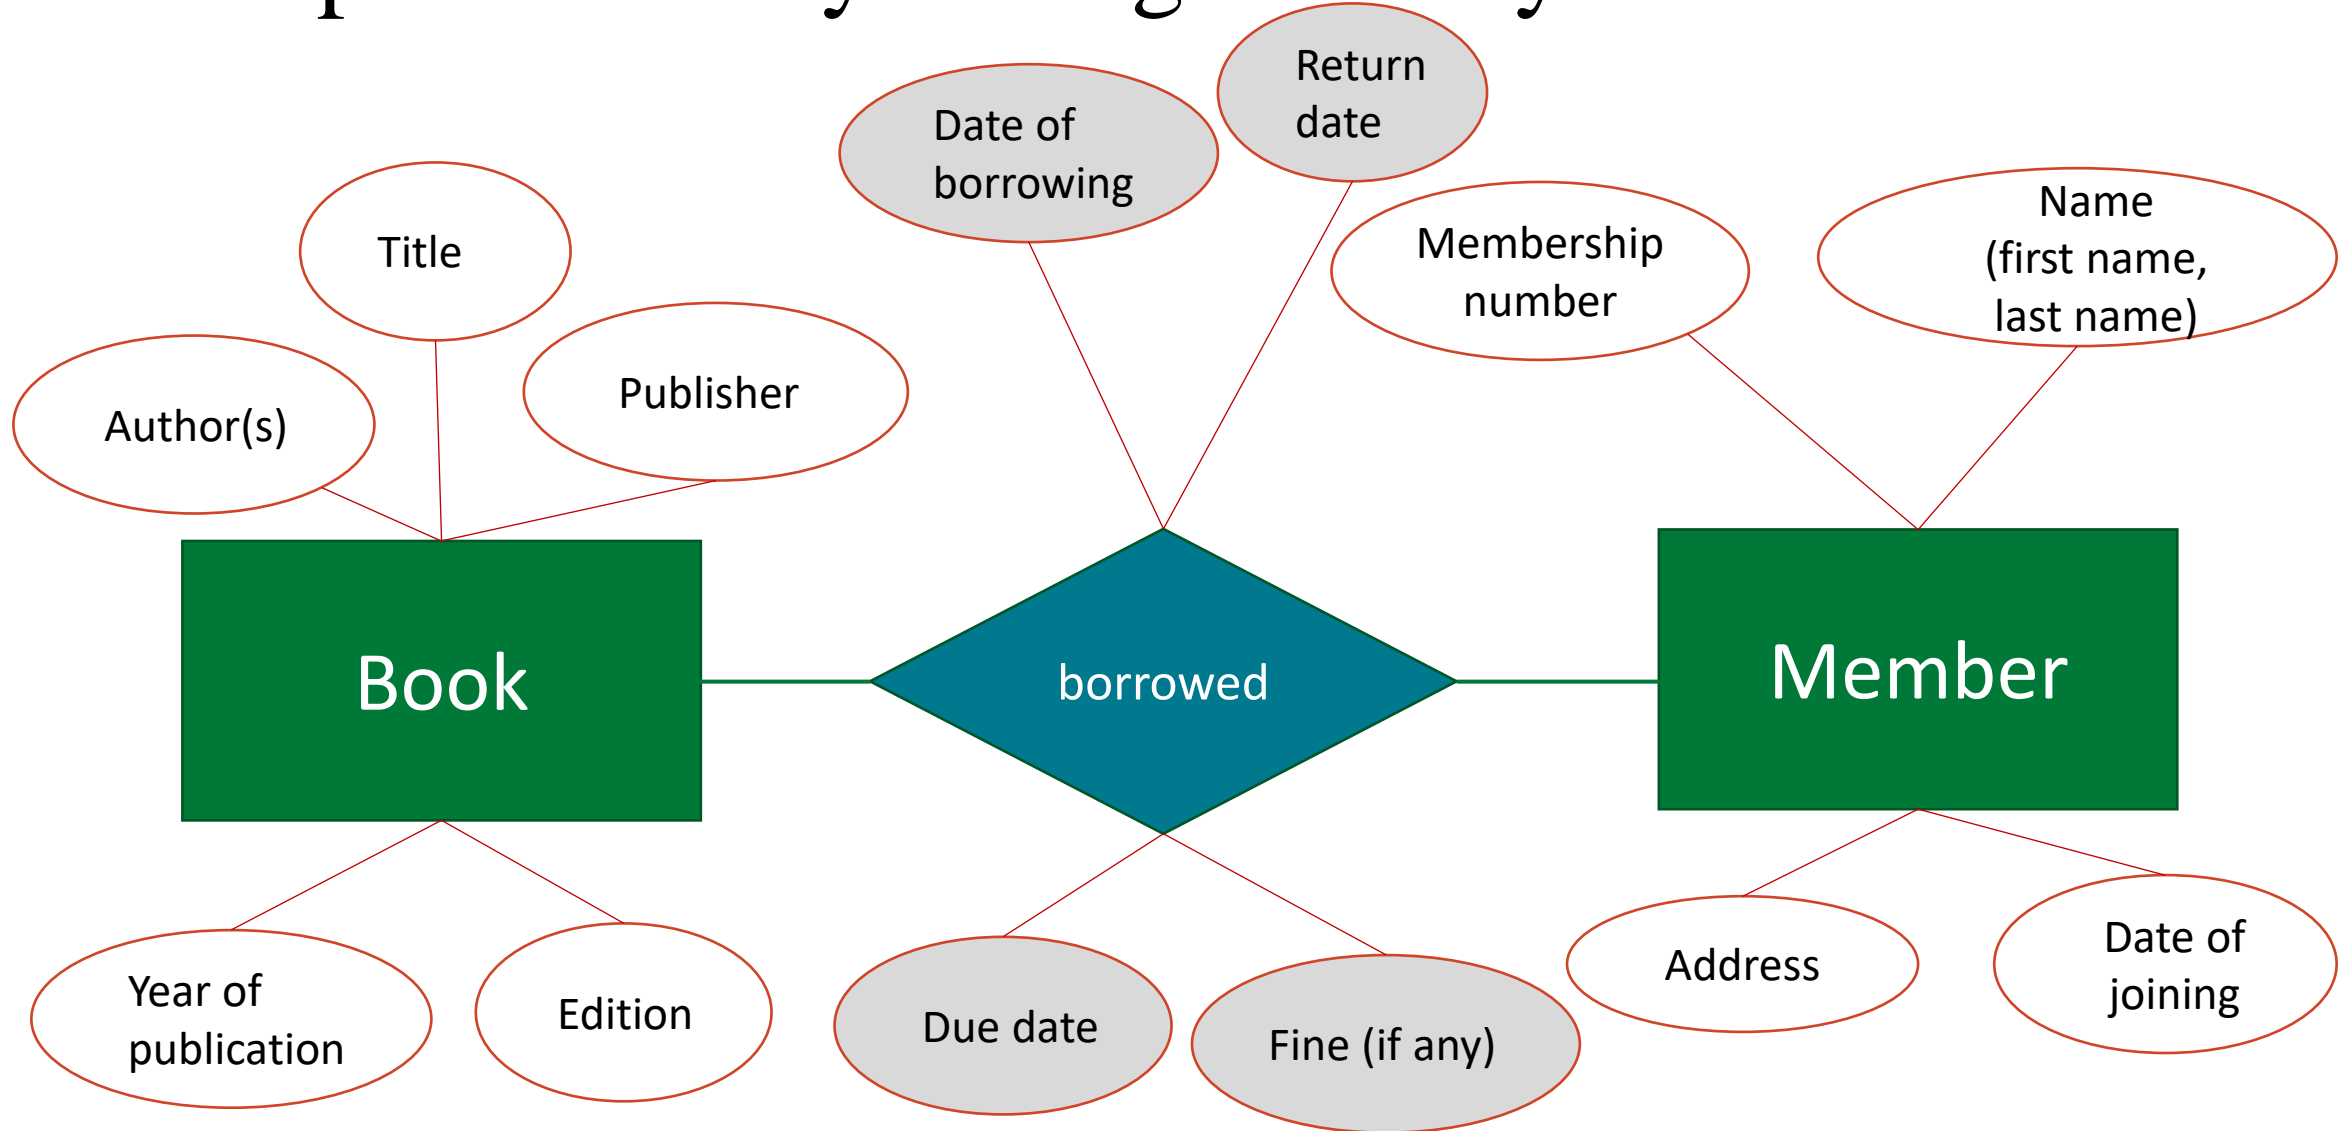
- Member

- Membership number
- Name (first name, last name)
- Address
- Date of joining

borrowed

- Date of borrowing
- Due date
- Return date
- Fine (if any)

# Example – Library management system



# Example – Library management system

We prefer not to put attributes in the ER diagrams.  
They make the diagram look too crowded and noisy.



# Example – Library management system



We assume:

- There is only one copy of each book in the library.
- Member can borrow only one book at a time.
- We only record the current borrowings (After return the information is deleted.)

**How can we  
represent this  
information?**

# Multiplicities (Cardinalities)

For every **relationship** and every side (**entity**) of the relationship, we think about: **minimum**, and **maximum** number of times the **entity** can participate in the **relationship**.



A book may or may not be borrowed.  
If it is borrowed, there is only one borrowing.  
So, minimum = 0, maximum = 1.

A member may or may not borrow a book.  
But if they borrowed, they borrowed only one book.  
So, minimum = 0, maximum = 1.



# Example – Library management system (with multiplicities)

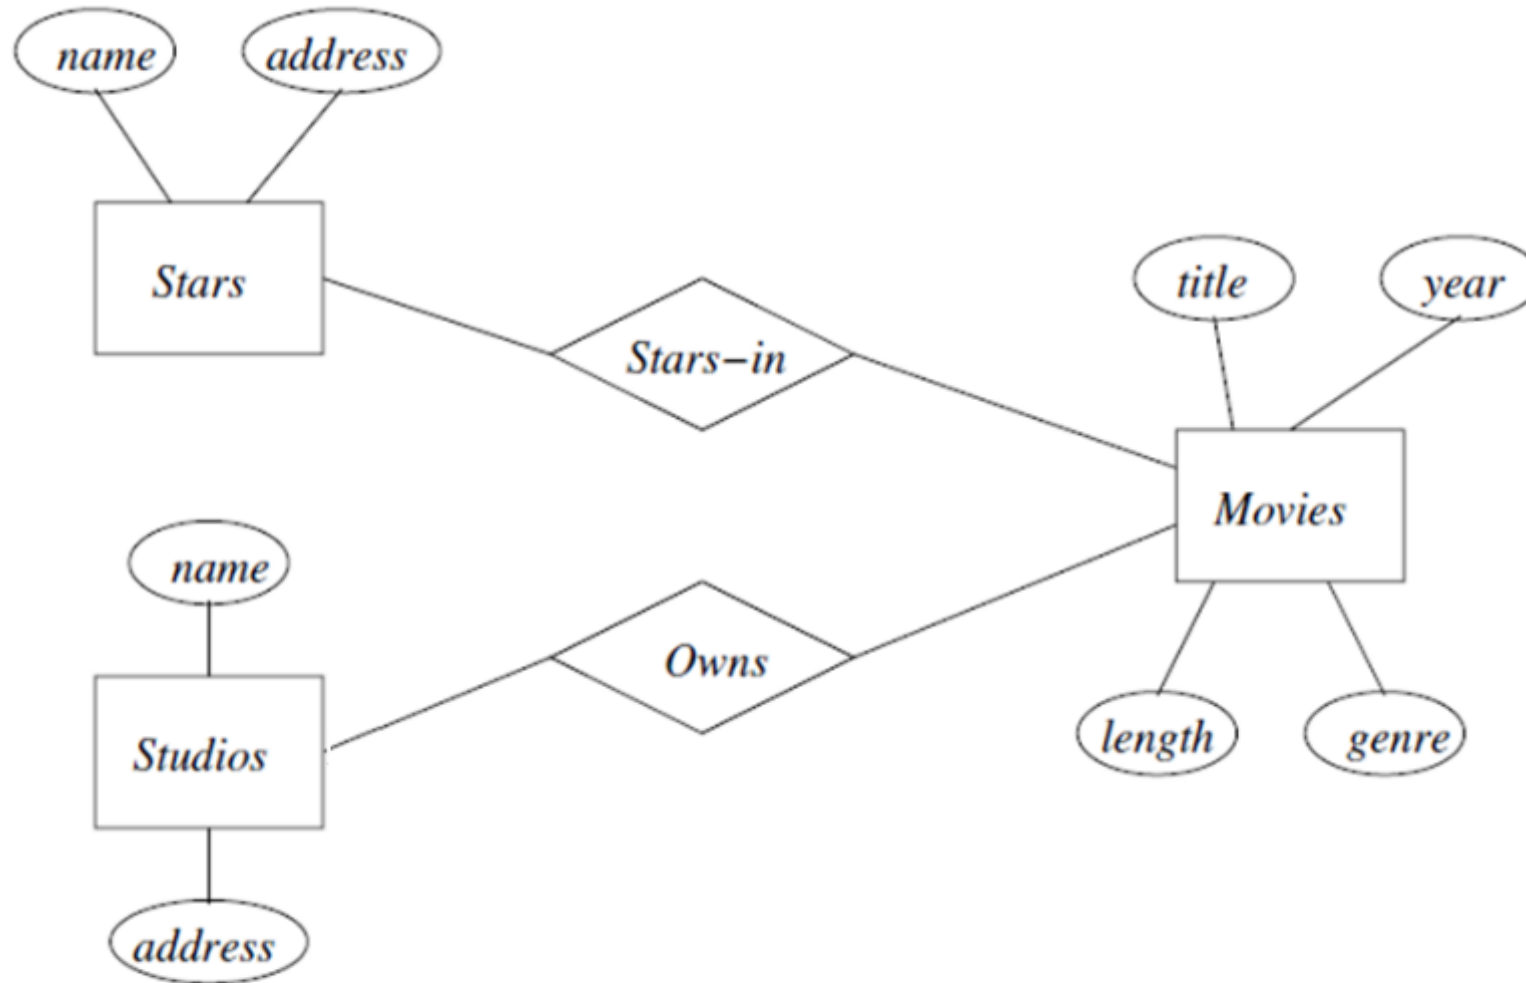
Each **side** of a relationship (link connecting to an entity) is given a **multiplicity**, specifying the minimum and maximum number of times the entity can participate in the relationship.



A book may or may not be borrowed.  
If it is borrowed, there is only one borrowing.  
So, minimum = 0, maximum = 1.

A member may or may not borrow a book.  
But if they borrowed, they borrowed only one book. So, minimum = 0, maximum = 1.

# Additional Example: Database about Movies



- “Stars”, “Movies” and “Studios” are **entities**.
- “Stars-in” and “Owns” are **relationships**.
- **attributes** are:
  - name, address
  - title, year, length, genre
  - name, address

Figure 2: An entity-relationship diagram for the movie database

# Additional Example: Database about Birthplaces

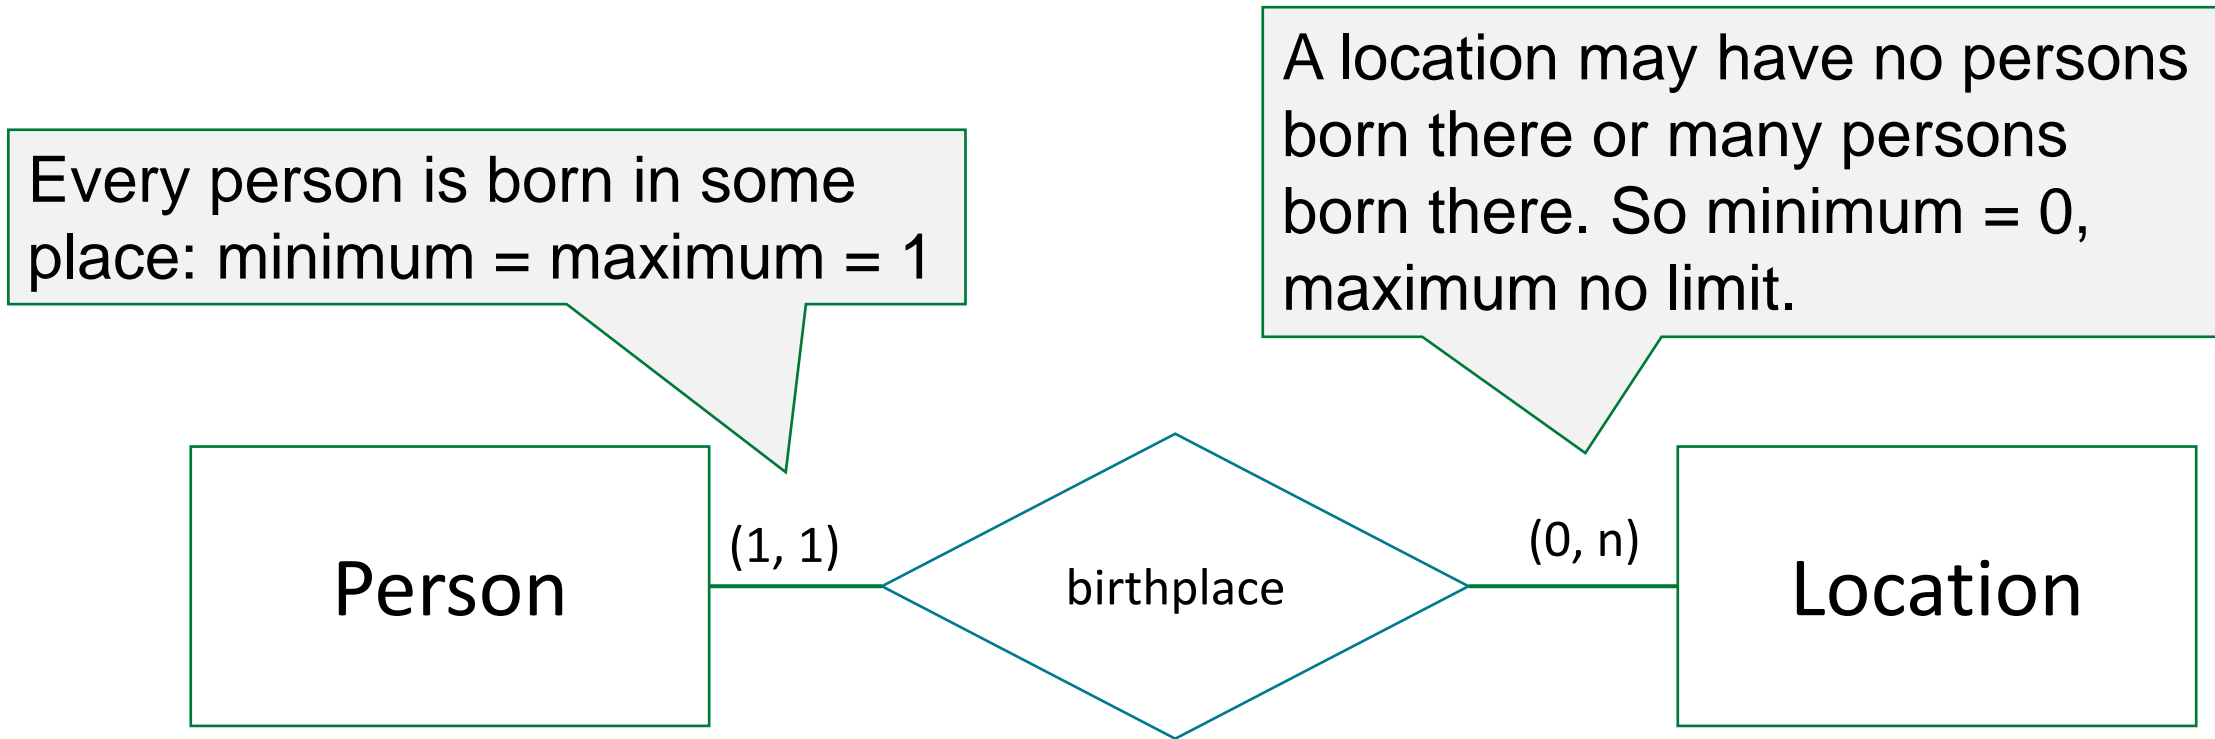
Every person is born in some place.

A location may have no persons born there or many persons born

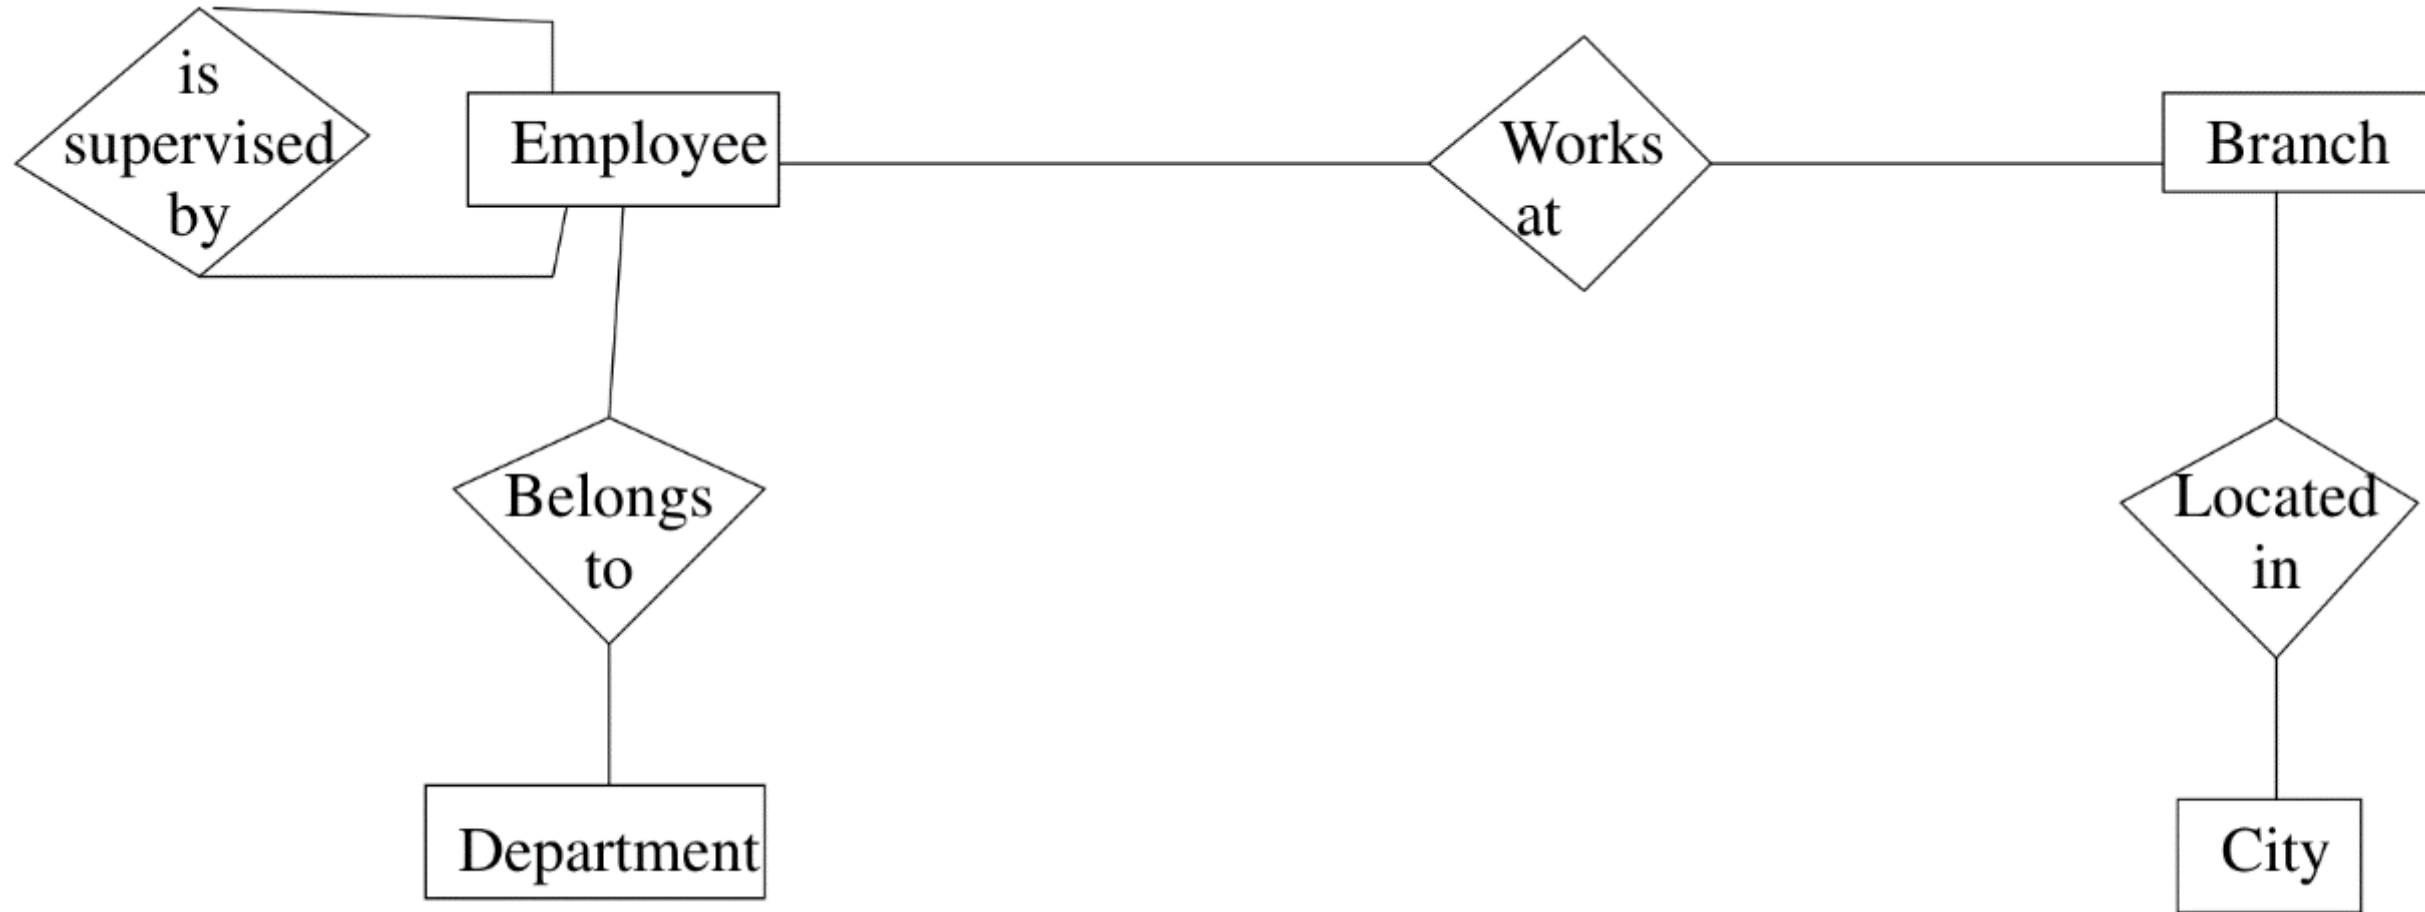
**What about multiplicity?**



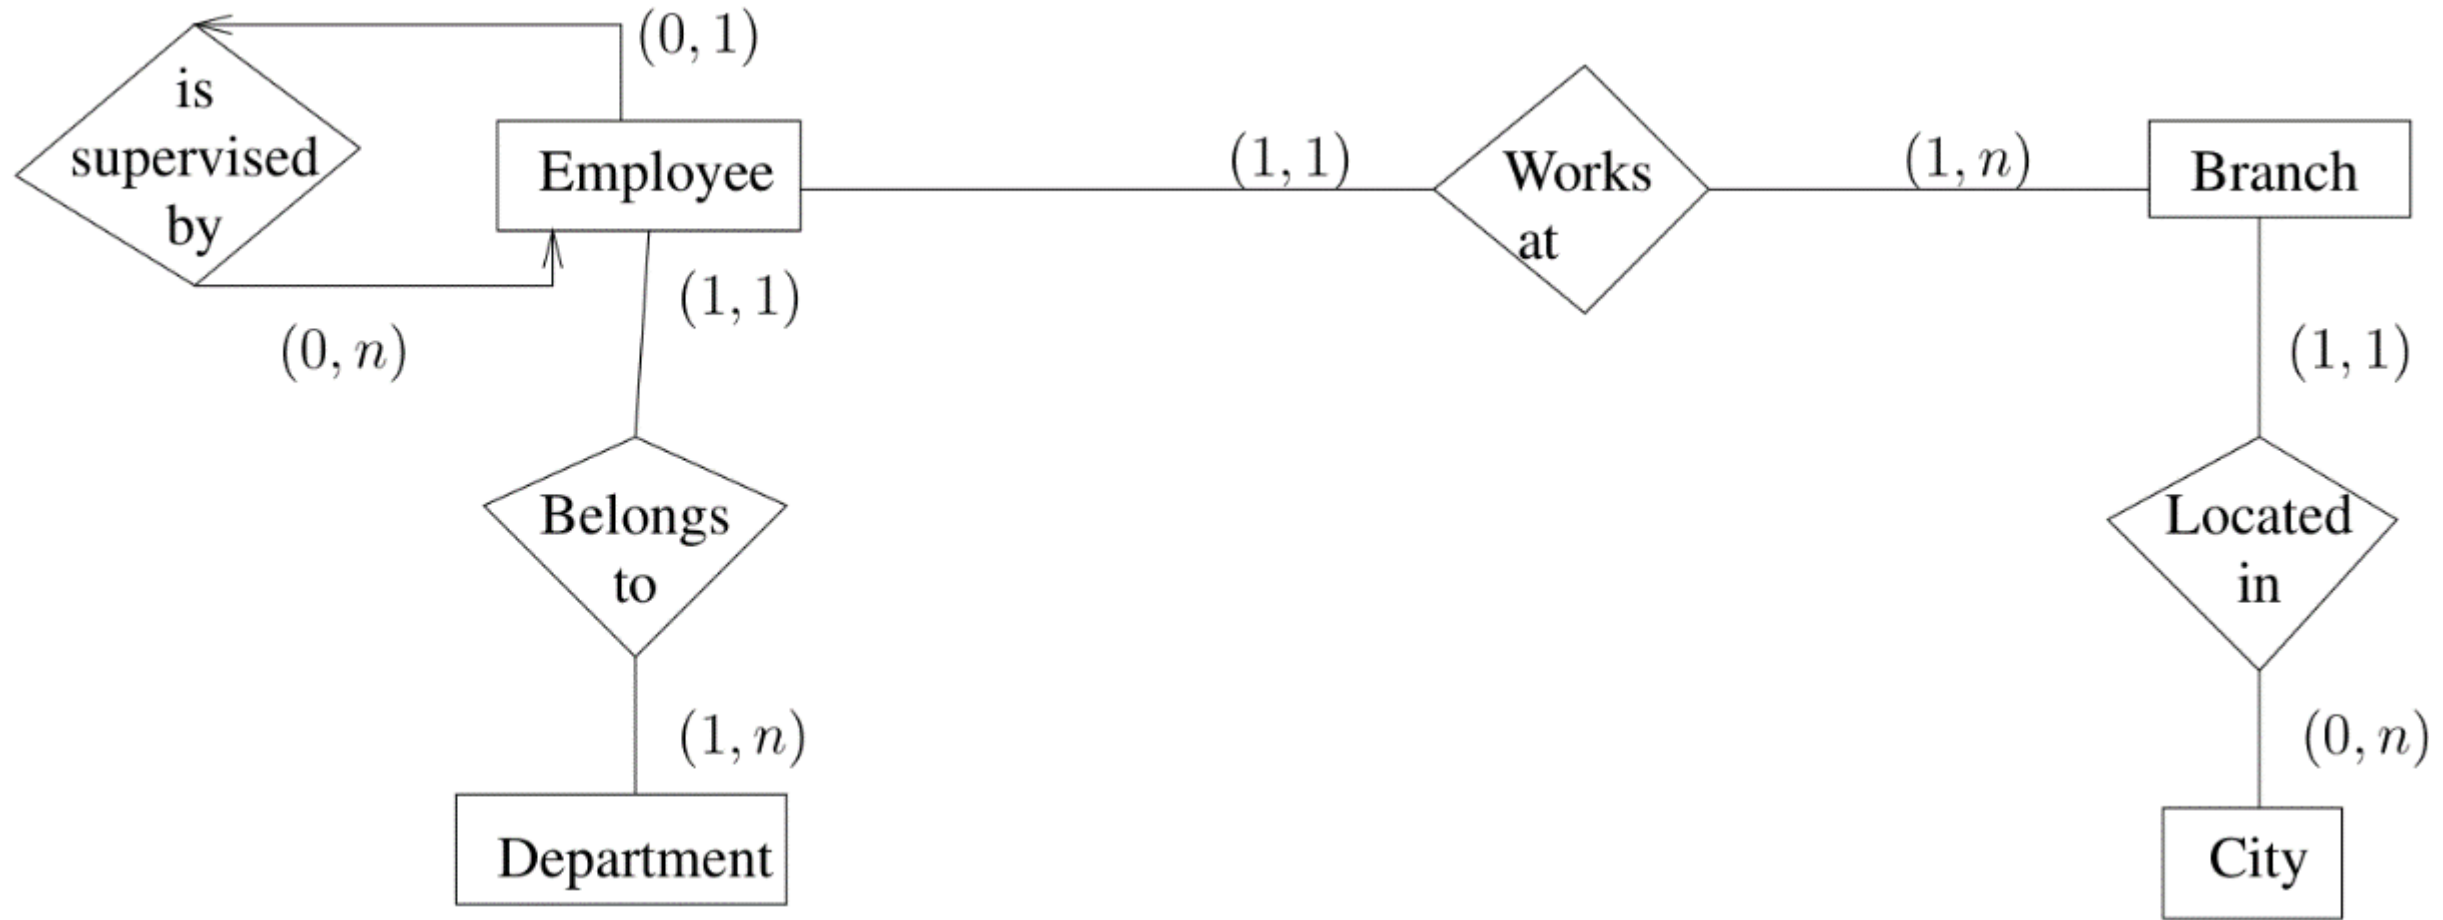
# Additional Example: Database about Birthplaces



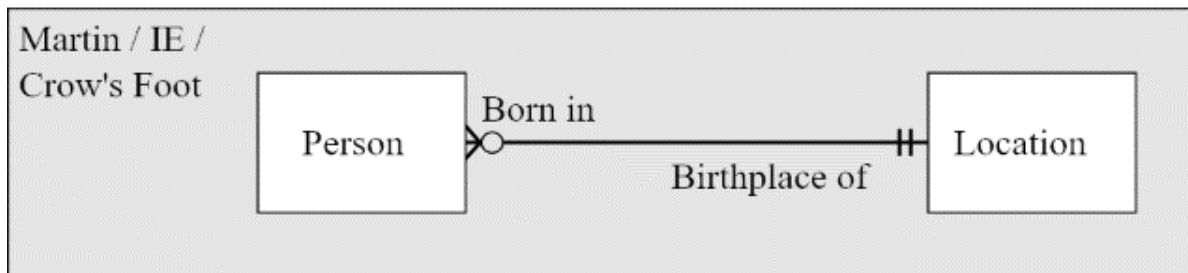
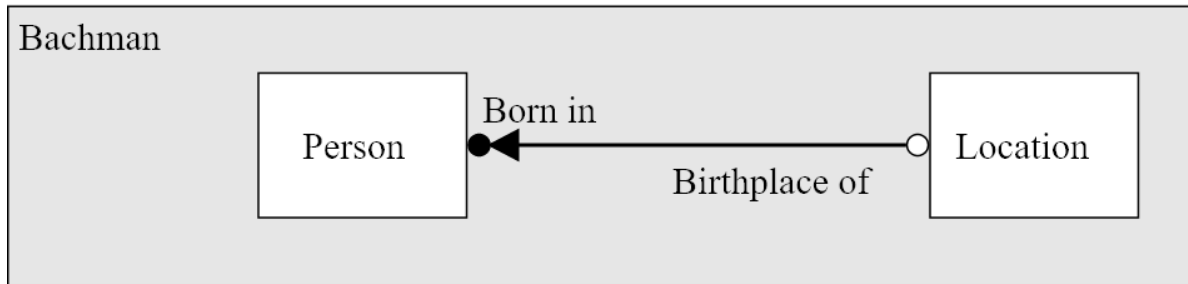

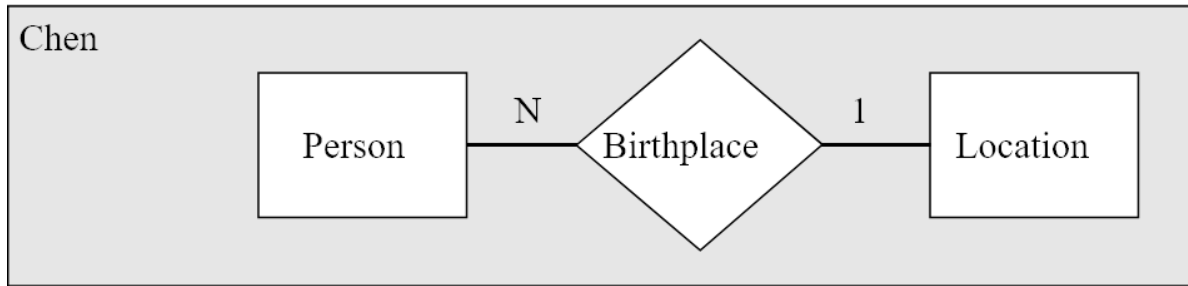
# Additional Example: human resources database



# Additional Example: human resources database (multiplicities)




# Multiplicity: A plethora of other notations



DO NOT USE  
IN THIS MODULE!



- The more modern notations **omit** the diamond box for relations.
- They draw a **line** for the relation, and sometimes label the line and sometimes not.
- Sometimes two labels are placed at the two ends.

# What exactly do you mean by “entity”, “entity set”?





# What exactly do you mean by “entity”, “entity set”?



# This Week

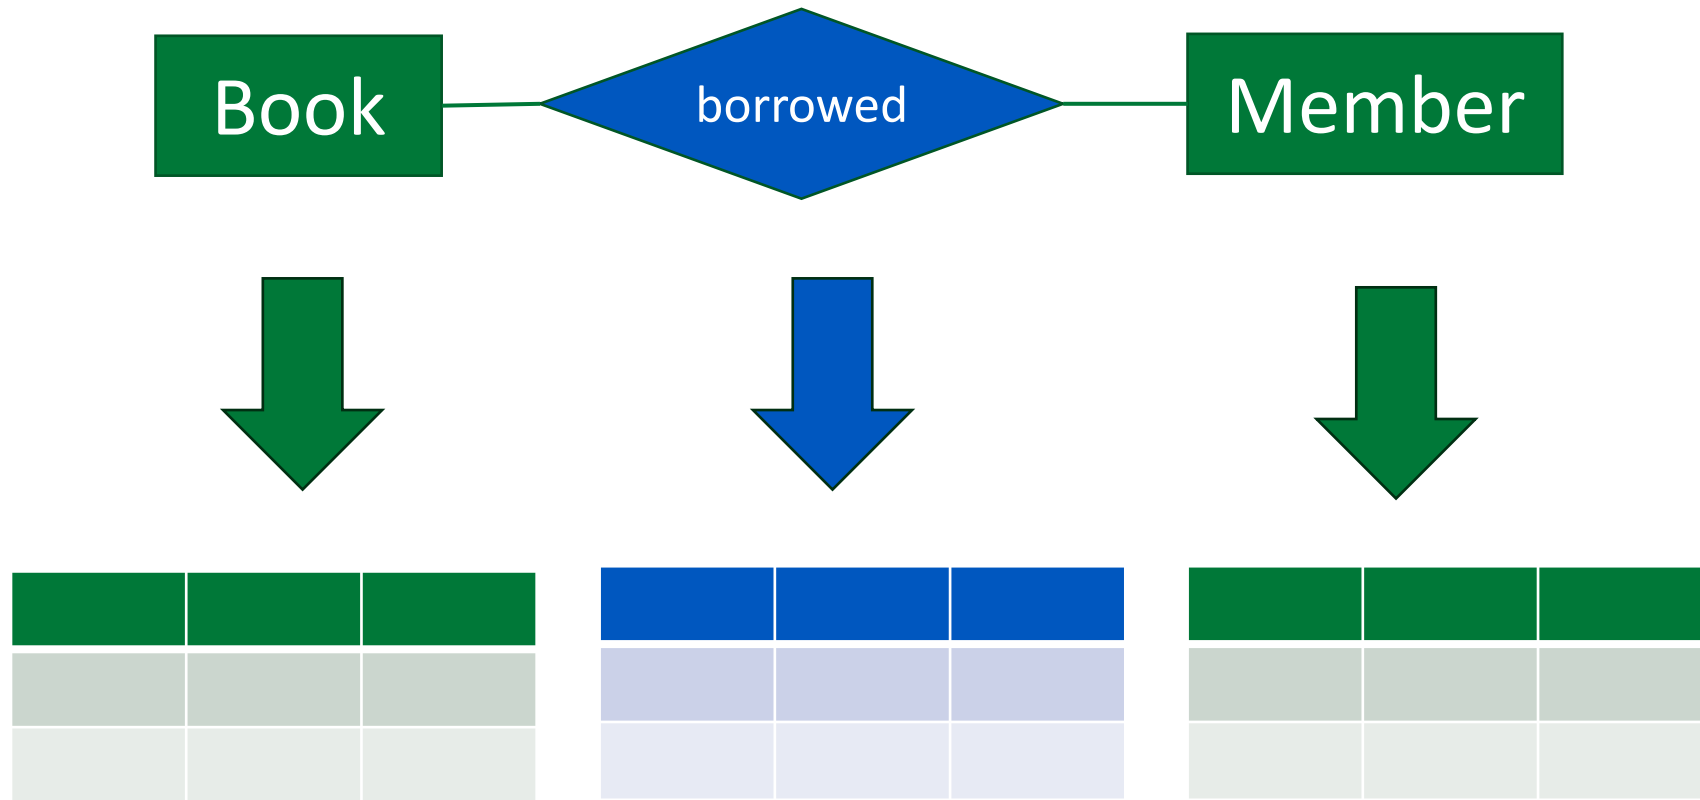
- Introduction
- Relational databases
- Entity-relationship (ER) diagram /Modeling
- ➔ Table design and creation, SQL
- Weak entities
- Hierarchies
- Table design (schemas) Optimisation
- SQL Commands + PostgreSQL



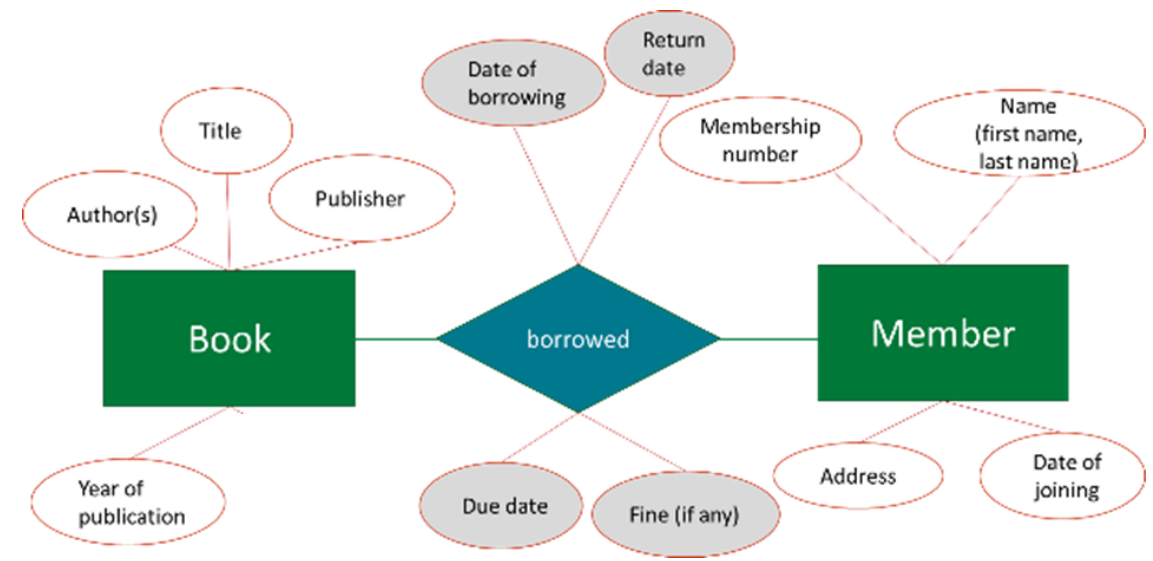
# Table design and creation

By default, every entity and relationship in our ER model becomes a table in the database.

(This can be optimised. We will get to that later.)



# Table design (schemas)



Unique  
(no duplicate)


The table designs are written in the **schema** notation.

Book(book id, authors, title, publisher, year)

Member(member num, last name, first name, address, date of joining)

Borrow(book id, member num, date, due date, return date, fine)

# Table design (schemas)



Unique  
identifiers  
(no duplicate)

The table designs are written in the **schema** notation.

Book(book id, authors, title, publisher, year)


Identifier in ER  
language

Member(member num, last name, first name, address, date of joining)

Primary key in  
Table

Borrow(book id, member num, date, due date, return date, fine)

# Table design (schemas)



The table designs are written in the **schema** notation.

Book(book id, authors, title, publisher, year)

Member(member num, last name, first name, address, date of joining)

Borrow(book id, member num, date, due date, return date, fine)

2 fields of  
primary key

must necessarily occur in the Book and Member tables (constraint)

# Relational DBMS (SQL)

A **database management system** (DBMS) is a software system (made by providers like Oracle, Postgres, MySQL).

- Creating tables.
- Querying tables for finding information.
- Adding, deleting or modifying records in tables.
- Ensuring that **constraints** continue to be satisfied during modifications.

All this is done using a “little” programming language called **SQL** (“Structured Query Language”).

**Constraints example:**  
Two books should not have same ID number

Member can borrow only 1 book. Database should not have 2 records for the same member at the same time. System should give error, when trying to insert 2nd record for the same member



# SQL command to create the book table

Book(book id, authors, title, publisher, year)

```
create table book (bookid    integer not null unique,  
                  authors    character varying(100),  
                  title       character varying(40),  
                  publisher    character varying(20),  
                  year         integer,  
  
                  primary key (bookid));
```



- The **unique** keyword ensures that a particular bookid occurs in only one record in the book table.
- “**not null**” says this field cannot be null. All fields can be null by default (a bad feature of SQL!)
- The **primary key** declaration says that the bookid field can be used for uniquely identify records.



# SQL command to create the book table

Book(book id, authors, title, publisher, year)

```
create table book (bookid integer not null unique,
                  authors character varying(100),
                  title character varying(40),
                  publisher character varying(20),
                  year integer,
                  primary key (bookid));
```




## Values that can be stored in this table:

bookid	authors	title	publisher	year
1	Dennis M. Ritchie	C Programming	Prentice Hall	1978
2	Herbert	Java: A Beginner's Guide	McGraw-Hill Education	2018
3	Eric Matthes	Python Crash Course		2015
...				

# SQL command to create the member table

Member(member num, last name, first name, address, date of joining)


```
create table member(member_num integer not null unique,  
                    last_name  character varying(20),  
                    first_name character varying(20),  
                    address    character varying(100),  
                    join_date  date,  
  
                    primary key (member_num));
```



# SQL command to create the member table

Member(member\_num, last name, first name, address, date of joining)

```
create table member(member_num integer not null unique,  
                    last_name  character varying(20),  
                    first_name character varying(20),  
                    address     character varying(100),  
                    join_date  date,  
  
                    primary key (member_num));
```



Values that can be stored in this table:


member_num	last_name	first_name	address	join_date
1	Smith	John	123 Street, London	15/01/2023
2	Brown	James	789 Road, Birmingham	10/03/2023
3				

# SQL command to create the borrow table

Borrow(book id, member num, date, due date, return date, fine)

```
create table borrow(bookid      integer not null references book(bookid),
                    member_num integer not null references member(member_num),
                    date         date not null,
                    due_date     date not null,
                    return_date  date,
                    fine          integer,

                    primary key (bookid, member_num));
```



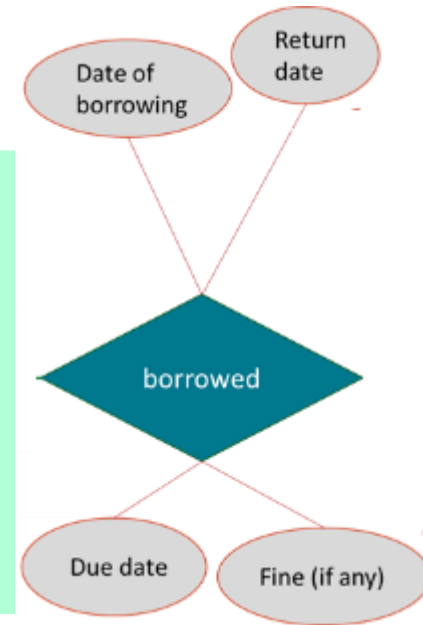
- We declare that the bookid and member\_num fields should **reference** the corresponding fields in the book and member tables. So we can't issue non-existent books to non-existent members. These kinds of declarations are **constraints**.
- The primary key declaration says that the bookid and member\_num fields **together** uniquely identify a record in the table.

# SQL command to create the borrow table

Borrow(book id, member num, date, due date, return date, fine)

```
create table borrow(bookid      integer not null references book(bookid),
                    member_num integer not null references member(member_num),
                    date         date not null,
                    due_date     date not null,
                    return_date  date,
                    fine         integer,

                    primary key (bookid, member_num));
```



## Values that can be stored in this table:

bookid	member_num	borrow_date	due_date	return_date	fine
1	1	16/01/2023	15/02/2023	NULL	NULL
2	2	05/04/2023	05/05/2023	NULL	NULL

# SQL query examples

book

bookid	authors	title	publisher	year
1	Dennis M. Ritchie	C Programming	Prentice Hall	1978
2	Herbert	Java: A Beginner's Guide	McGraw-Hill Education	2018
3	Eric Matthes	Python Crash Course		2015
...				

member

member_num	last_name	first_name	address	join_date
1	Smith	John	123 Street, London	15/01/2023
2	Brown	James	789 Road, Birmingham	10/03/2023
3				

borrow

bookid	member_num	borrow_date	due_date	return_date	fine
1	1	16/01/2023	15/02/2023	NULL	NULL
2	2	05/04/2023	05/05/2023	NULL	NULL

# SQL query examples

book

bookid	authors	title	publisher	year
1	Dennis M. Ritchie	C Programming	Prentice Hall	1978
2	Herbert	Java: A Beginner's Guide	McGraw-Hill Education	2018
3	Eric Matthes	Python Crash Course		2015
...				

member

member_num	last_name	first_name	address	join_date
1				15/01/2023
2	Brown	James	789 Road, Birmingham	10/03/2023
3				

Finds the id number of 'Python Crash Course' by 'Eric Matthes'?

borrow

bookid	member_num	borrow_date	due_date	return_date	fine
1	1	16/01/2023	15/02/2023	NULL	NULL
2	2	05/04/2023	05/05/2023	NULL	NULL

# SQL query examples

book

bookid	authors	title	publisher	year
1	Dennis M. Ritchie	C Programming	Prentice Hall	1978
2	Herbert	Java: A Beginner's Guide	McGraw-Hill Education	2018
3	Eric Matthes	Python Crash Course		2015
...				

member

member_num	join_date
1	15/01/2023
2	10/03/2023
3	

borrow

bookid	member_num	borrow_date	return_date	fine
1			NULL	NULL
2	2	05/04/2023	05/05/2023	NULL

Finds the id number of 'Python Crash Course' by 'Eric Matthes'?

```
select bookid
from book
where authors = 'Eric Matthes'
and title = 'Python Crash Course';
```



# SQL query examples

book

bookid	authors	title	publisher	year
1	Dennis M. Ritchie	C Programming	Prentice Hall	1978
2	Herbert	Java: A Beginner's Guide	McGraw-Hill Education	2018
3	Eric Matthes	Python Crash Course		2015
...				

member

member_num	last_name	join_date
1	Smith	15/01/2023
2	Brown	10/03/2023
3		

Check if the book 'C Programming' by 'Dennis M. Ritchie' is available for checkout.

borrow

bookid	member_num	borrow_date	due_date	return_date	fine
1	1	16/01/2023	15/02/2023	NULL	NULL
2	2	05/04/2023	05/05/2023	NULL	NULL

# SQL query examples

bookid	authors	title	publisher	year
1	Dennis M. Ritchie	C Programming	Prentice Hall	1978
2	Herbert	Java: A Beginner's Guide	McGraw-Hill Education	2018
3	Eric Matthes	Python Crash Course		2015
...				

book

member_num	last_name	first_name
1	Smith	
2	Brown	
3		

member

date
2023
2023

bookid	member_num	borrowed
1	1	16/
2	2	05/

borrow

fine
NULL
NULL

Check if the book 'C Programming' by 'Dennis M. Ritchie' is available for checkout.

```
select 'not available'
from book, borrow
where book.bookid = borrow.bookid
and book.authors = 'dennis m. ritchie'
and book.title = 'c programming';
```

# This Week

- Introduction
- Relational databases
- Entity-relationship (ER) diagram /Modeling
- Table design and creation, SQL
- ➔ Weak entities
- Hierarchies
- Table design (schemas) Optimisation
- SQL Commands + PostgreSQL




# Weak entities

An entity, which depends on some other entity for its identity. Entities that are not weak are called “strong entities” by some tutors. A weak entity’s primary key consists of the primary key of its **main entity** (or **owner entity**) and some additional attributes. Example: **Copy** dependent on **Book**.



How to represent multiple copies of same book in library in ER Diagram?

# Weak entities - library management system



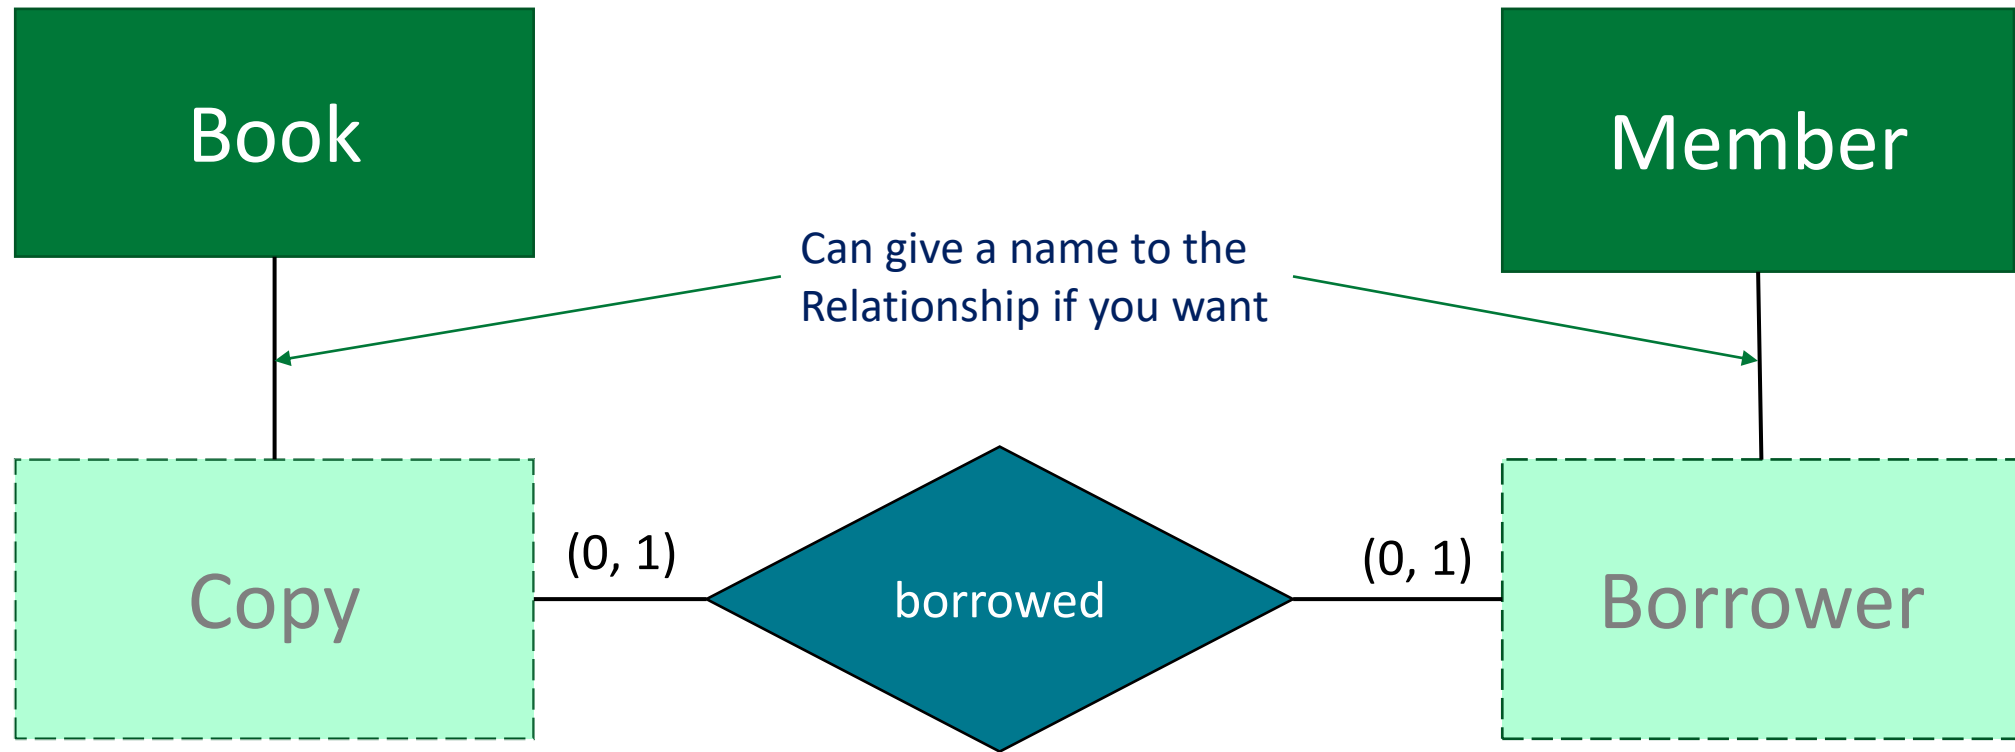
Suppose our library has  
**multiple copies** of some of  
the books.

How do we handle them?

Suppose our library allows  
**family memberships**.

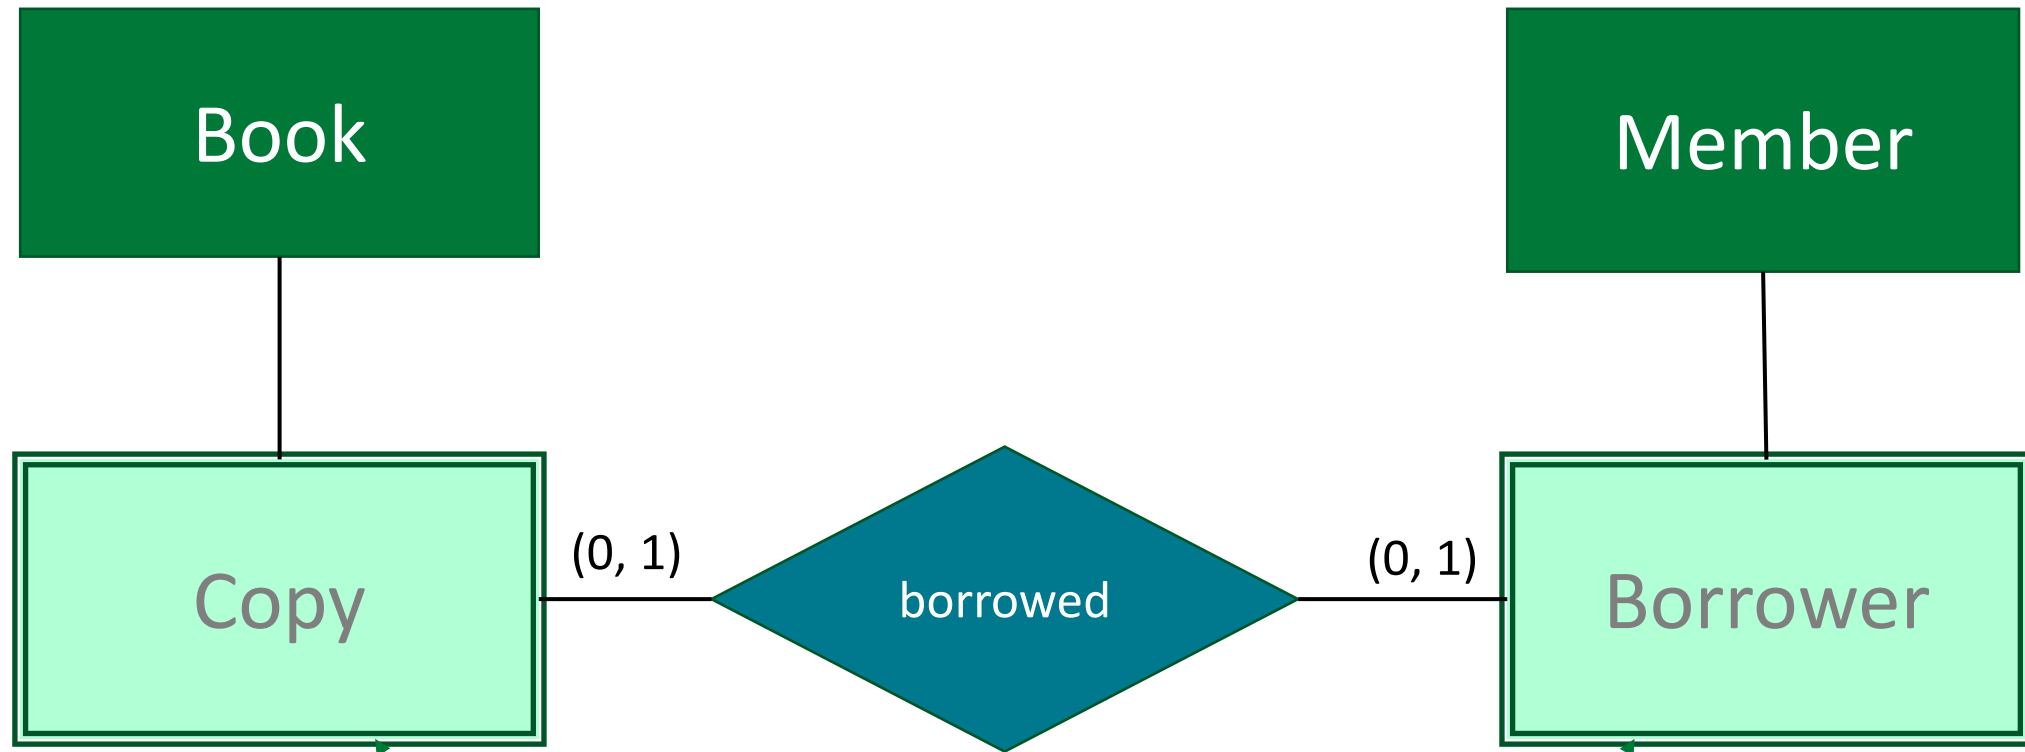
In addition to the primary  
member, a number of **relatives**  
can also borrow books.

# Revised ER-diagram with Weak entities



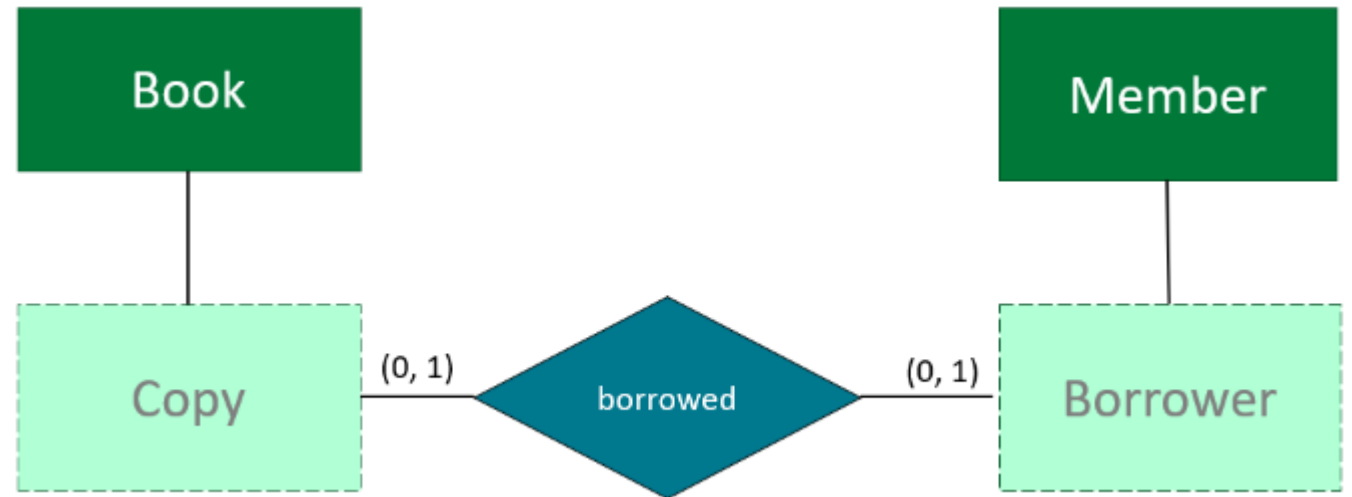
**Copies** of books are borrowed by **borrowers** associated with primary members.  
Think about what needs to be done if a book has a single copy.  
And when the primary member wants to borrow a book on their own.

# Another notation for weak entities



Some people use double lines to depict weak entities.

# Table for **weak** entities



The table designs in the **schema** notation.

Book(book id, authors, title, publisher, year)

Copy(book id, copy num)

Member(member num, last name, first name, address, date of joining)

Borrower(member num, first name)

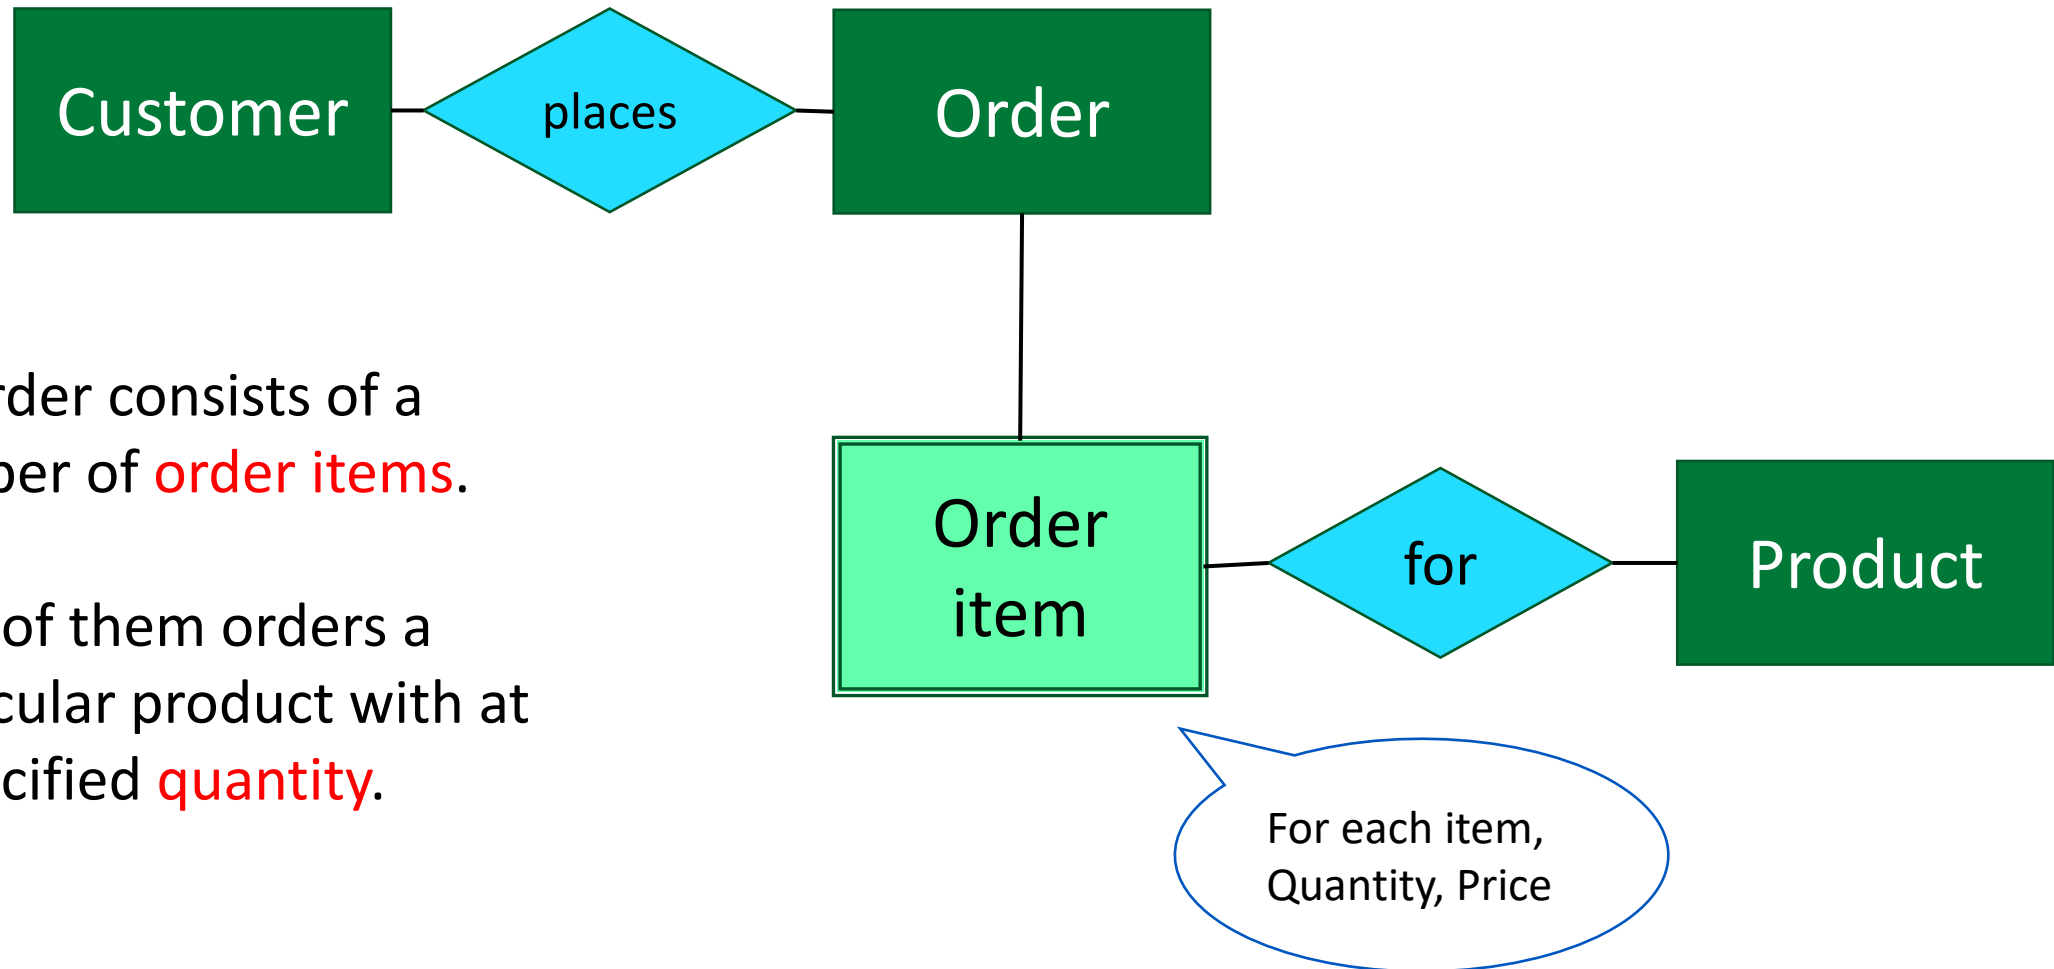
Assuming last name is common

Borrow(book id, copy num, member num, first name, date, due date, return date, fine)

Composite primary key



# Another common example for **weak** entities



An order consists of a number of **order items**.

Each of them orders a particular product with at a specified **quantity**.


# This Week

- Introduction
- Relational databases
- Entity-relationship (ER) diagram /Modeling
- Table design and creation, SQL
- Weak entities
- ➔ Hierarchies
  - Table design (schemas) Optimisation
  - SQL Commands + PostgreSQL




# Hierarchies (generalisation hierarchies)

Hierarchies represent “is-a” relationships just like in Object Oriented Programming.



# Hierarchies

Hierarchies represent “is-a” relationships just like in Object Oriented Programming.



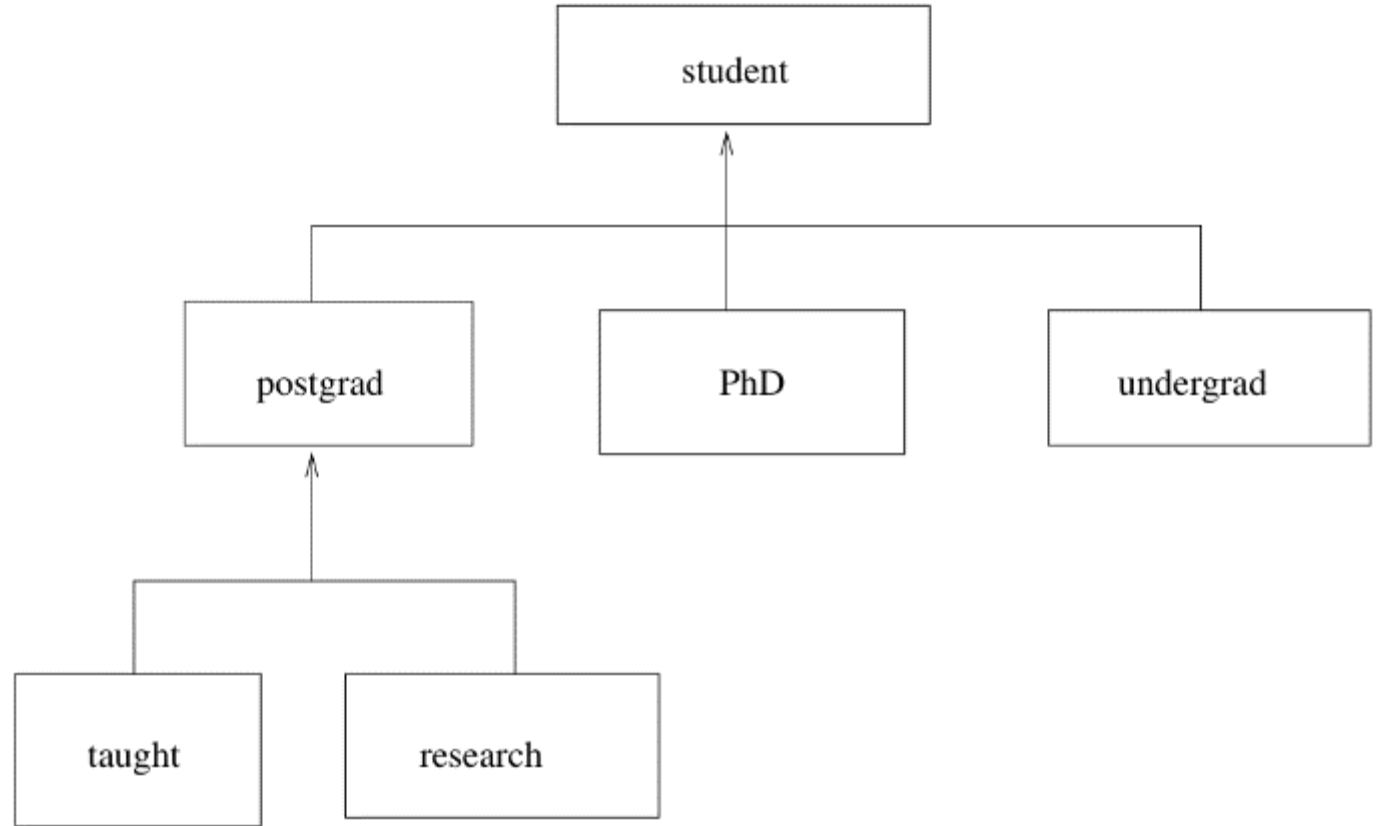
**Example-1** Cars, trucks and motor cycles are special cases of a more general **vehicle** entity.

Vehicle hierarchy

# Hierarchies

Hierarchies represent “**is-a**” relationships just like in Object Oriented Programming.

**Example-2** undergrad and postgrad student are special cases of a more general **student** entity (“super-entity”).

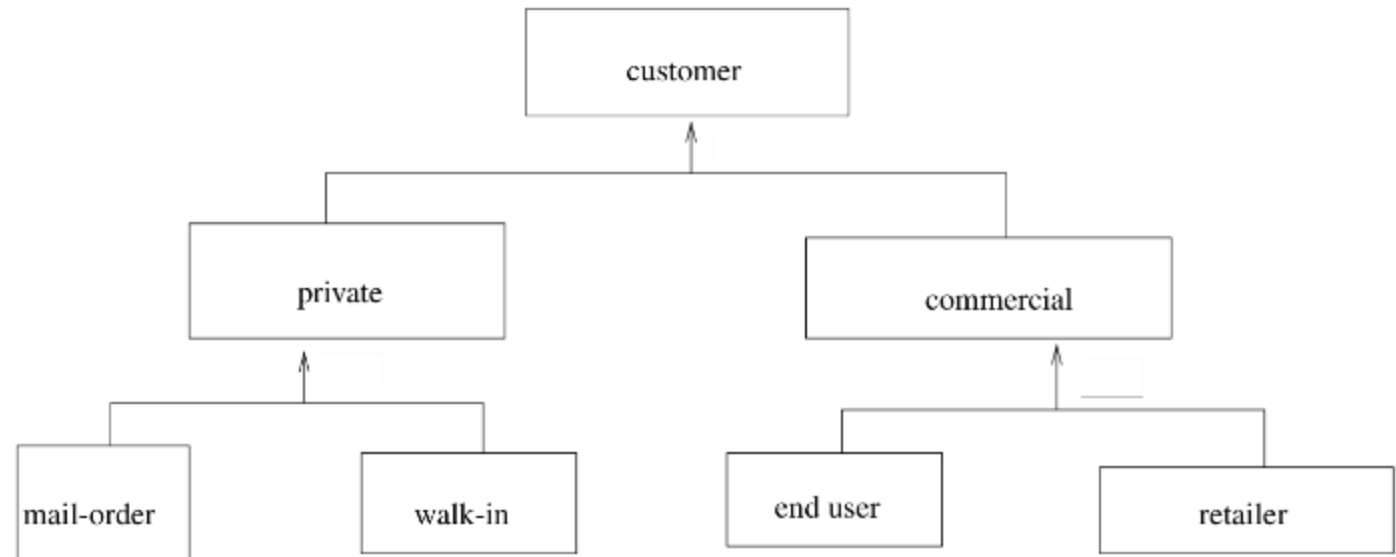


Student  
hierarchy

# Hierarchies

Hierarchies represent “is-a” relationships just like in Object Oriented Programming.

## Example-3 Customer hierarchy



Customer hierarchy

# Annotations for hierarchy relationships


Two features need to be specified.

Coverage **c**:

Do the subclasses **totally** (t) cover all of the superclass entity, or **partially** (p)?

Overlap **x**:

Are the subclasses **overlapping** (o), or are they **mutually exclusive** (e)?

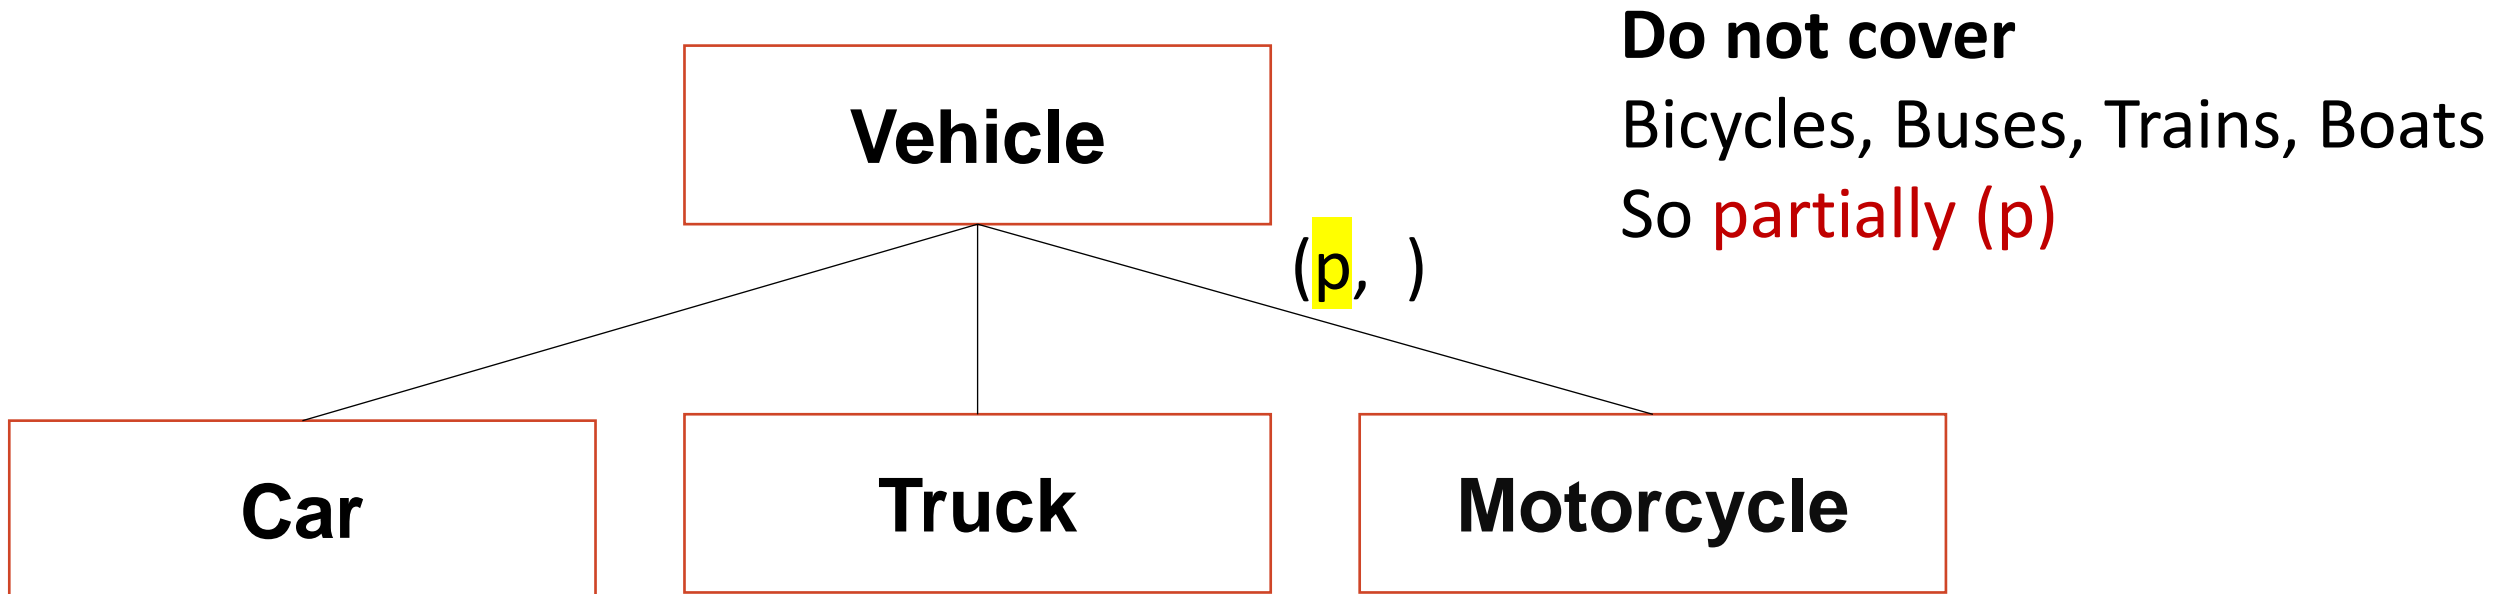


# Annotations for hierarchy relationships

Two features need to be specified.

Coverage **c**:

Do the subclasses **totally** (t) cover all of the superclass entity, or **partially** (p)?



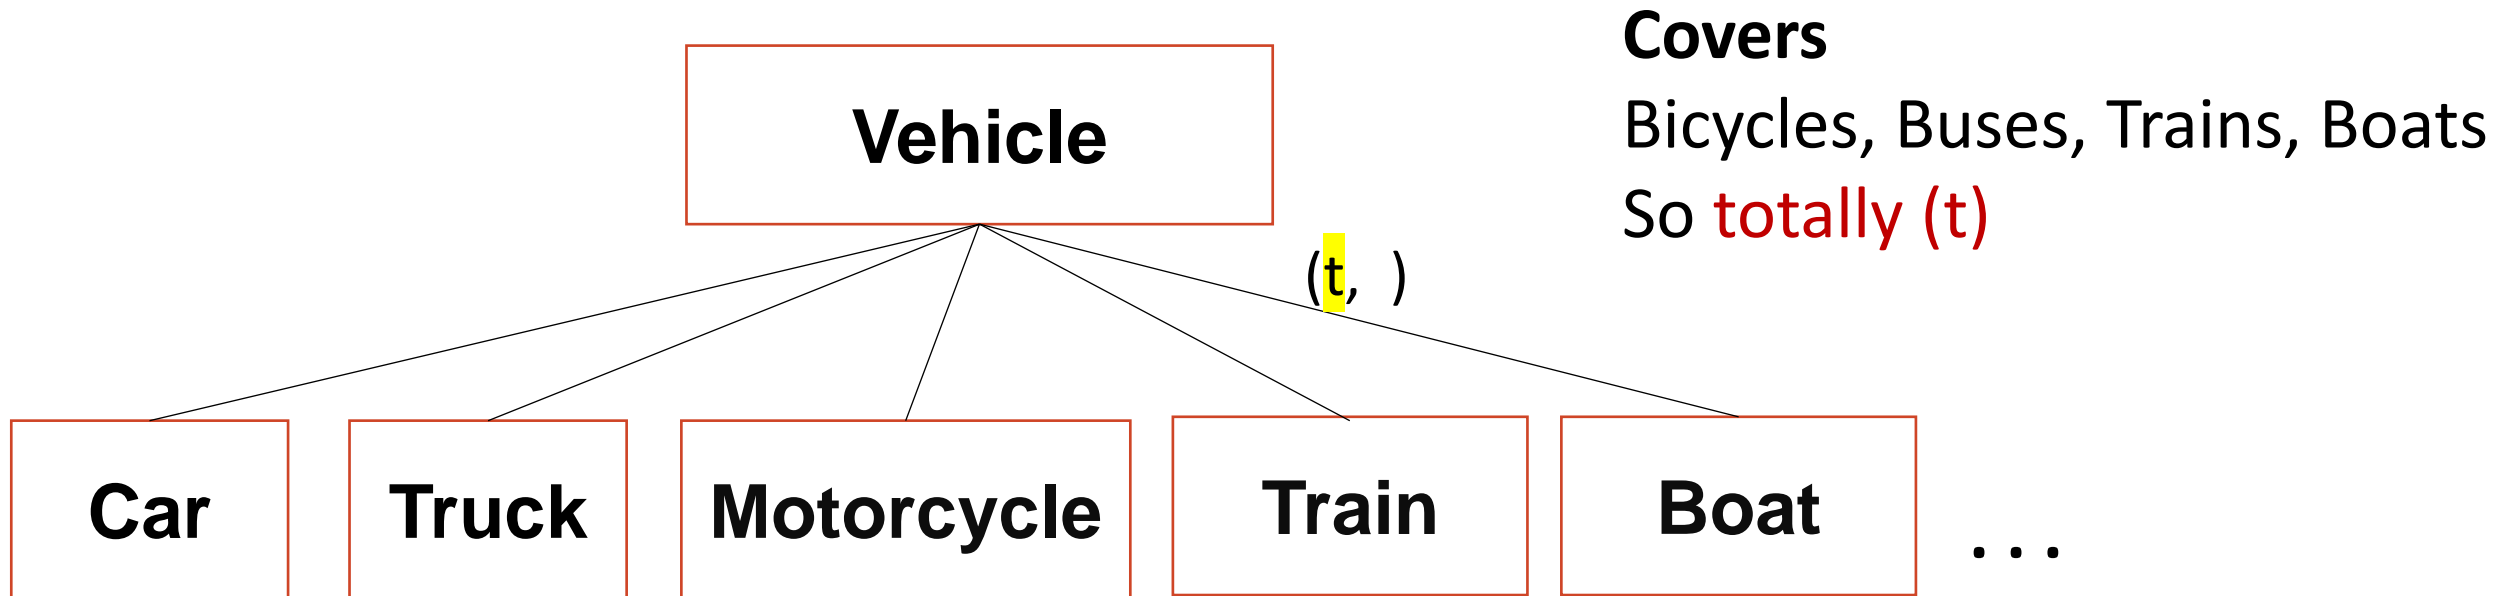


# Annotations for hierarchy relationships

Two features need to be specified.

Coverage **c**:

Do the subclasses **totally** (t) cover all of the superclass entity, or **partially** (p)?

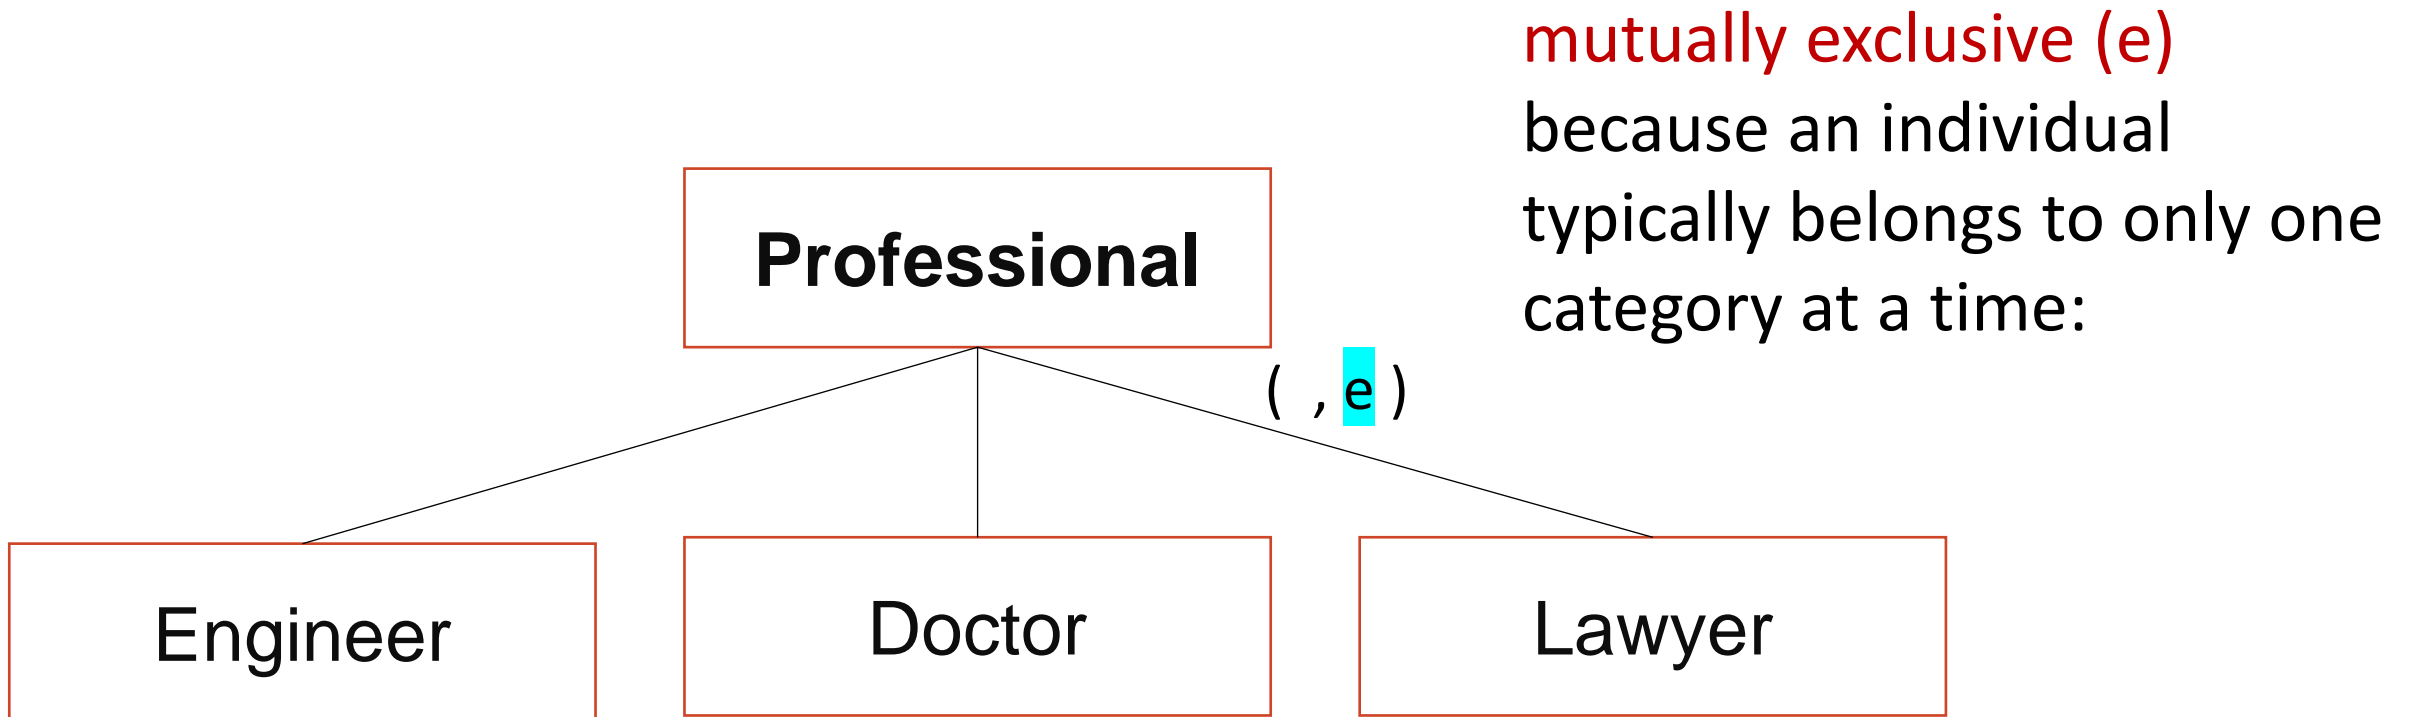


# Annotations for hierarchy relationships

Two features need to be specified.

Overlap **x**:

Are the subclasses **overlapping** (o), or are they **mutually exclusive** (e)?

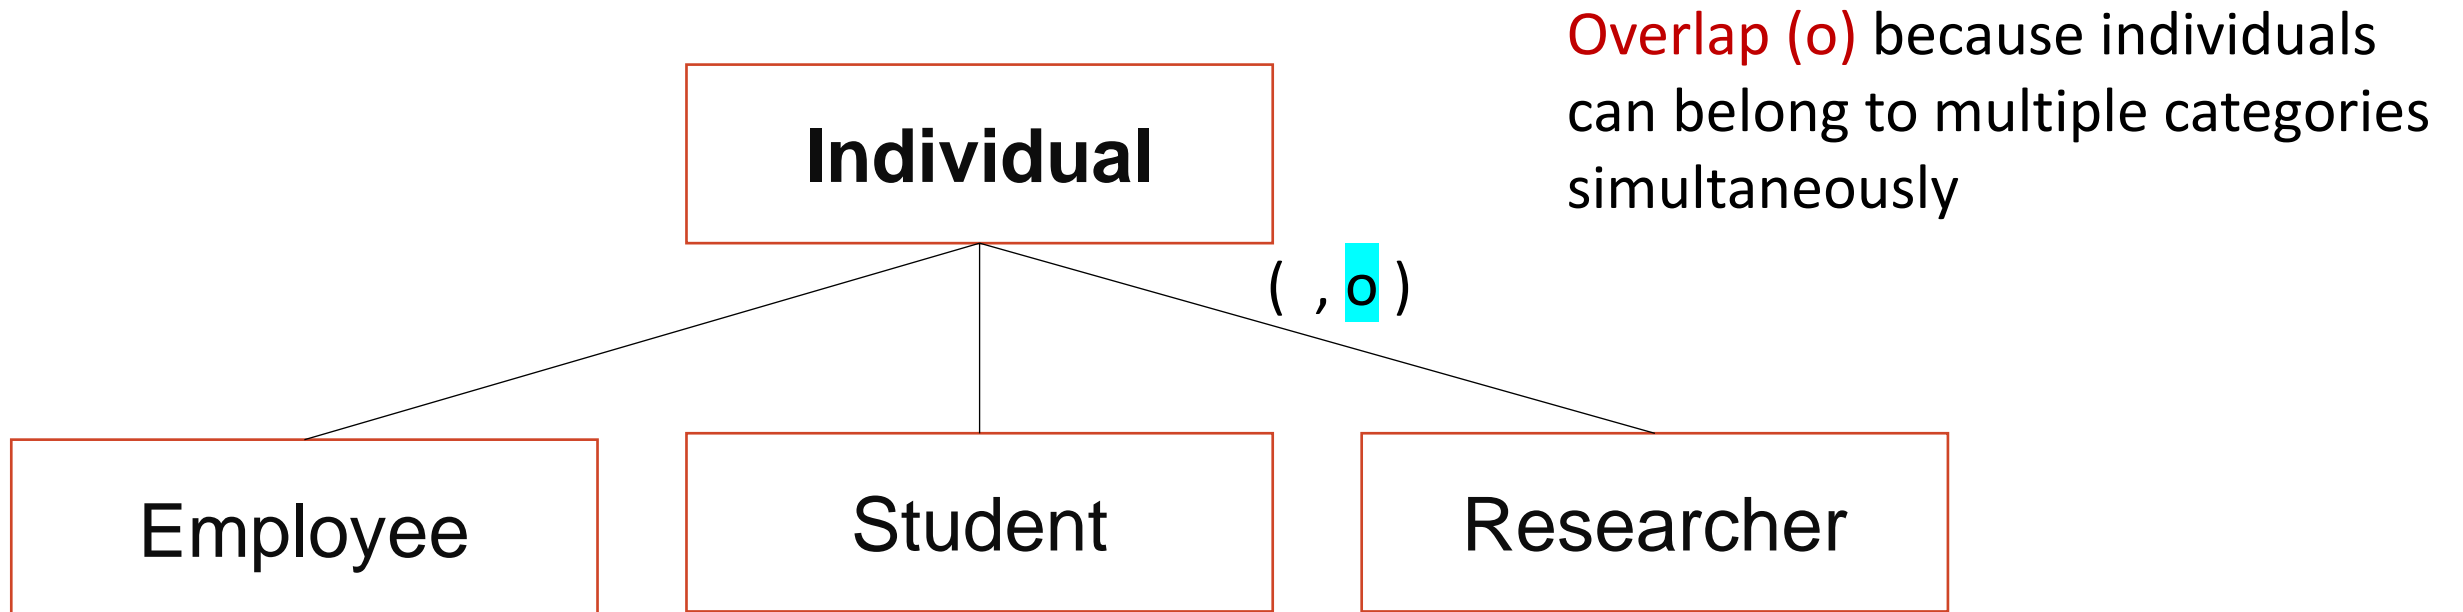


# Annotations for hierarchy relationships

Two features need to be specified.

Overlap **x**:

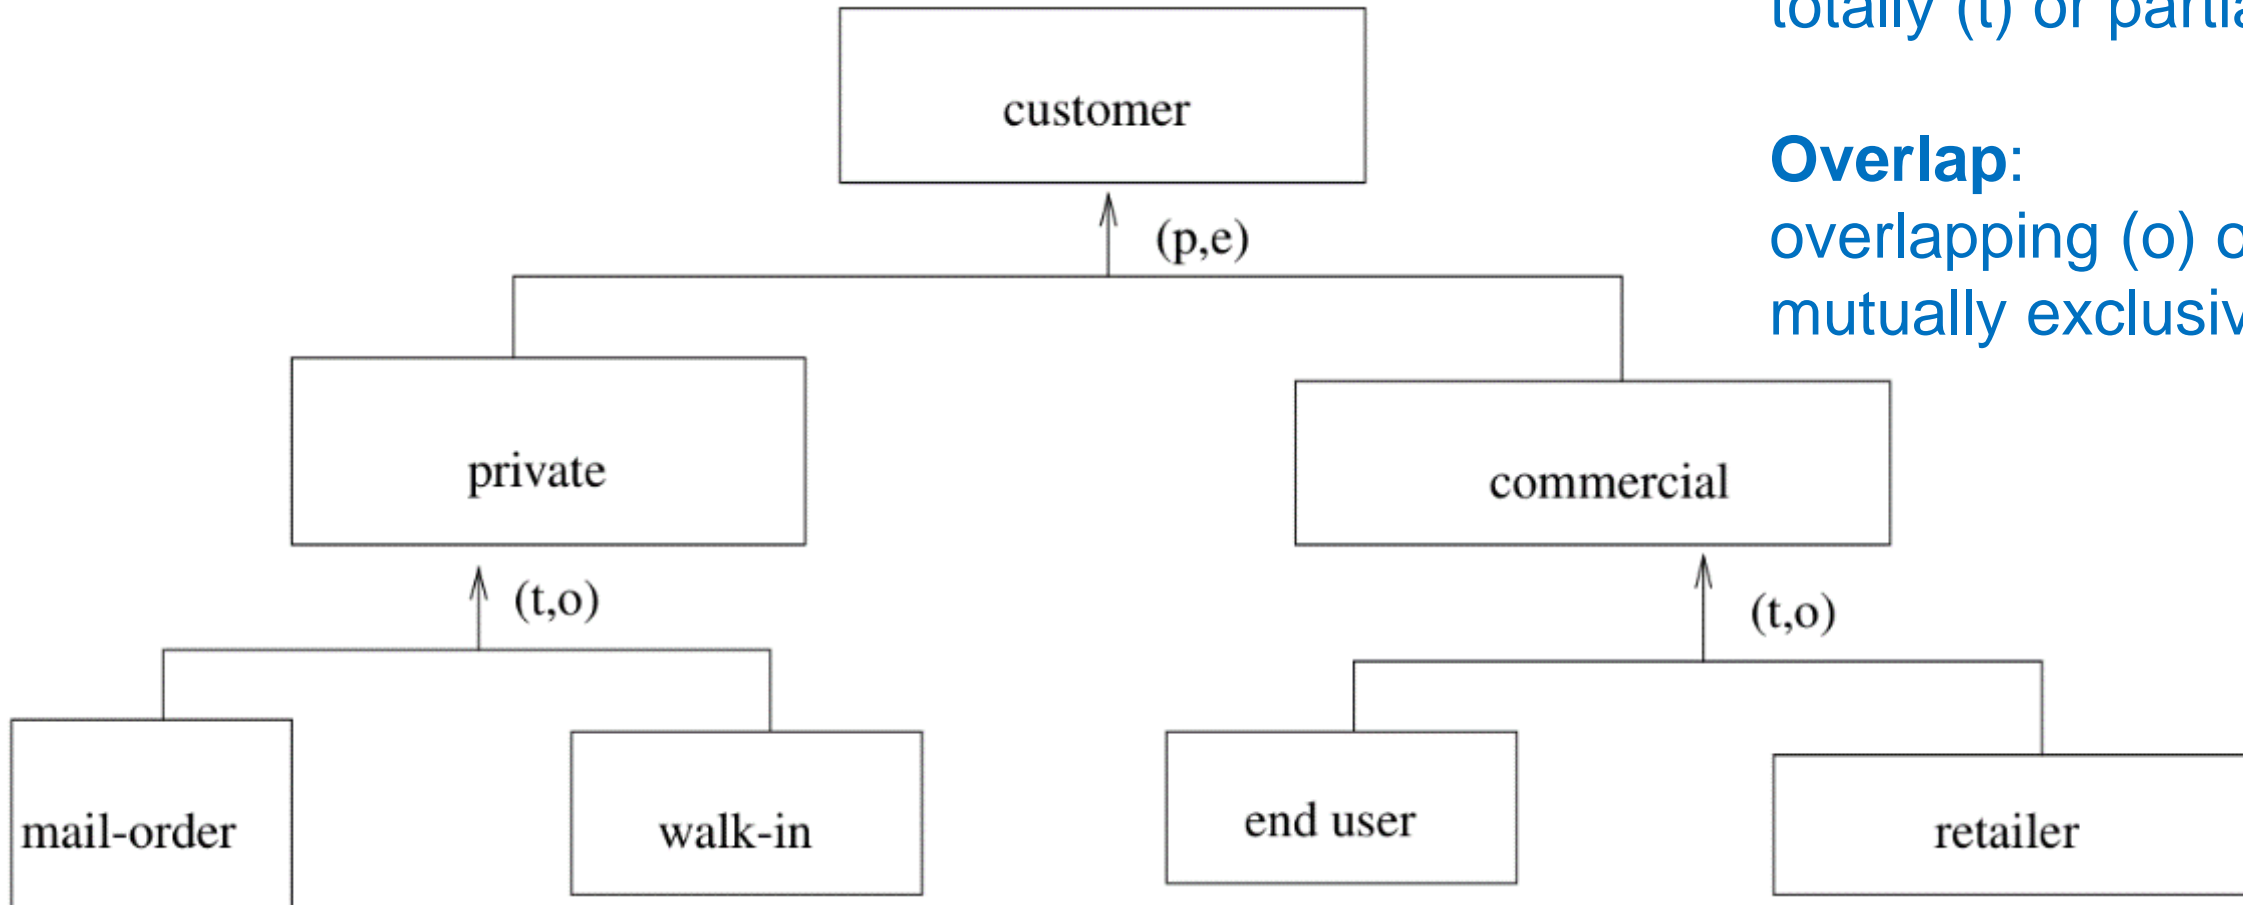
Are the subclasses **overlapping** (o), or are they **mutually exclusive** (e)?



# Example of Coverage and Overlap

**Coverage:**  
totally (t) or partially (p)

**Overlap:**  
overlapping (o) or  
mutually exclusive (e)



# This Week

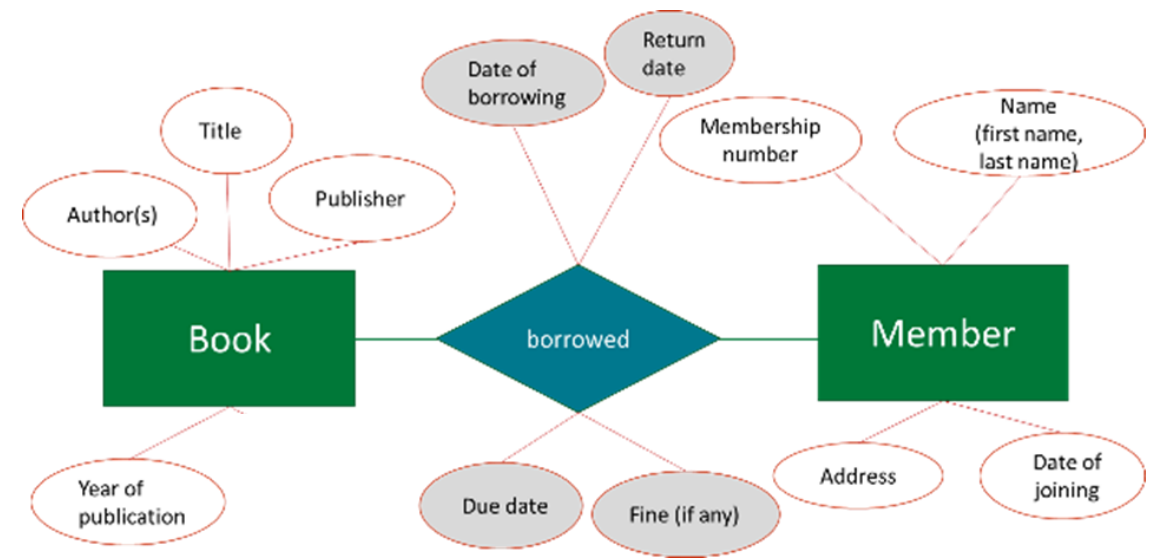
- Introduction
- Relational databases
- Entity-relationship (ER) diagram /Modeling
- Table design and creation, SQL
- Weak entities
- Hierarchies
- ➔ Table design (schemas) Optimisation
- SQL Commands + PostgreSQL



# Table design (schemas)

Do we really need 3 tables for this problem.

Perhaps 2 tables would be enough?



The table designs are written in the **schema** notation.

Book(book id, authors, title, publisher, year)


Member(member num, last name, first name, address, date of joining)

Borrow(book id, member num, date, due date, return date, fine)

# Table design (schemas)

Optimisation 1: We can add the **Borrow** fields to the **Member** table.

Optimisation 2: If we ignore fines, we can also add the **Borrow** fields to the **Book** table.




Book(book id, authors, title, publisher, year)

Member(member num, last name, first name, address, date of joining)

Borrow(book id, member num, date, due date, return date, fine)

# Optimisation 1



Original schema for 3 tables (2 entities + 1 relationship).

**Book** (book id, authors, title, publisher, year)

**Member** (member num, last name, first name, address, date of joining)

**Borrow** (book id, member num, borrow date, due date, ~~return date~~, ~~fine~~)


**Optimisation 1:** We can add the **Borrow** fields to the **Member** table.

**Book** (book id, authors, title, publisher, year)

**Member** (member num, last name, first name, address, date of joining,  
book id, borrow date, due date, ~~return date~~, ~~fine~~)



# Optimisation 2



Original schema for 3 tables (2 entities + 1 relationship).

**Book** (book id, authors, title, publisher, year)

**Member** (member num, last name, first name, address, date of joining)

**Borrow** (book id, member num, borrow date, due date, ~~return date~~, fine)

**Optimisation 2:** We can add the **Borrow** fields to the **Book** table.

**Book** (book id, authors, title, publisher, year,  
member num, borrow date, due date)

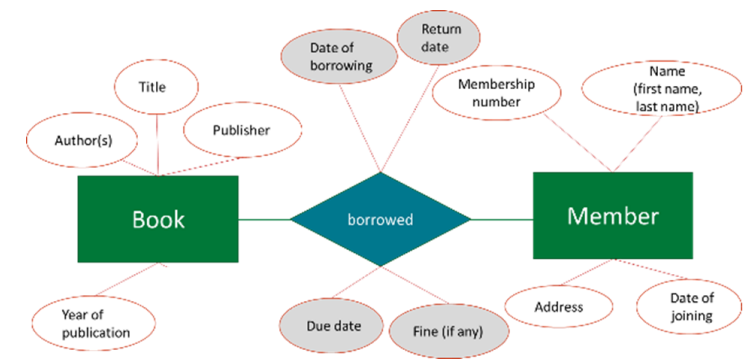
**Member** (member num, last name, first name, address, date of joining)

# Issues with optimisation - 1

Adding **Borrow** fields to the **Member** table.

**Book** (book id, authors, title, publisher, year)

**Member** (member num, last name, first name, address, date of joining,  
book id<sup>o</sup>, borrow date<sup>o</sup>, due date<sup>o</sup>, ~~return-date<sup>o</sup>~~, fine<sup>o</sup>)



## Issues

What happens when the member is not borrowing any book?

The **book id** field and other related fields would be **null**.

Represented by the <sup>o</sup> notations in our schemas.


What happens if the library decides to allow multiple books to be borrowed?

This design wouldn't work any more.

We can only merge relations into tables if the multiplicity is **maximum 1**.

**It is not a good idea to hard-wire business policies into the database design!**  
**They can change in future.**

# Issues with optimisation - 2



Adding **Borrow** fields to the **Book** table.

**Book** (book id, authors, title, publisher, year,  
member num°, borrow date°, due date°)

**Member** (member num, last name, first name, address, date of joining)

## Issues

What happens if the book is not borrowed at present?

The **book id** field and other related fields would be **null**.

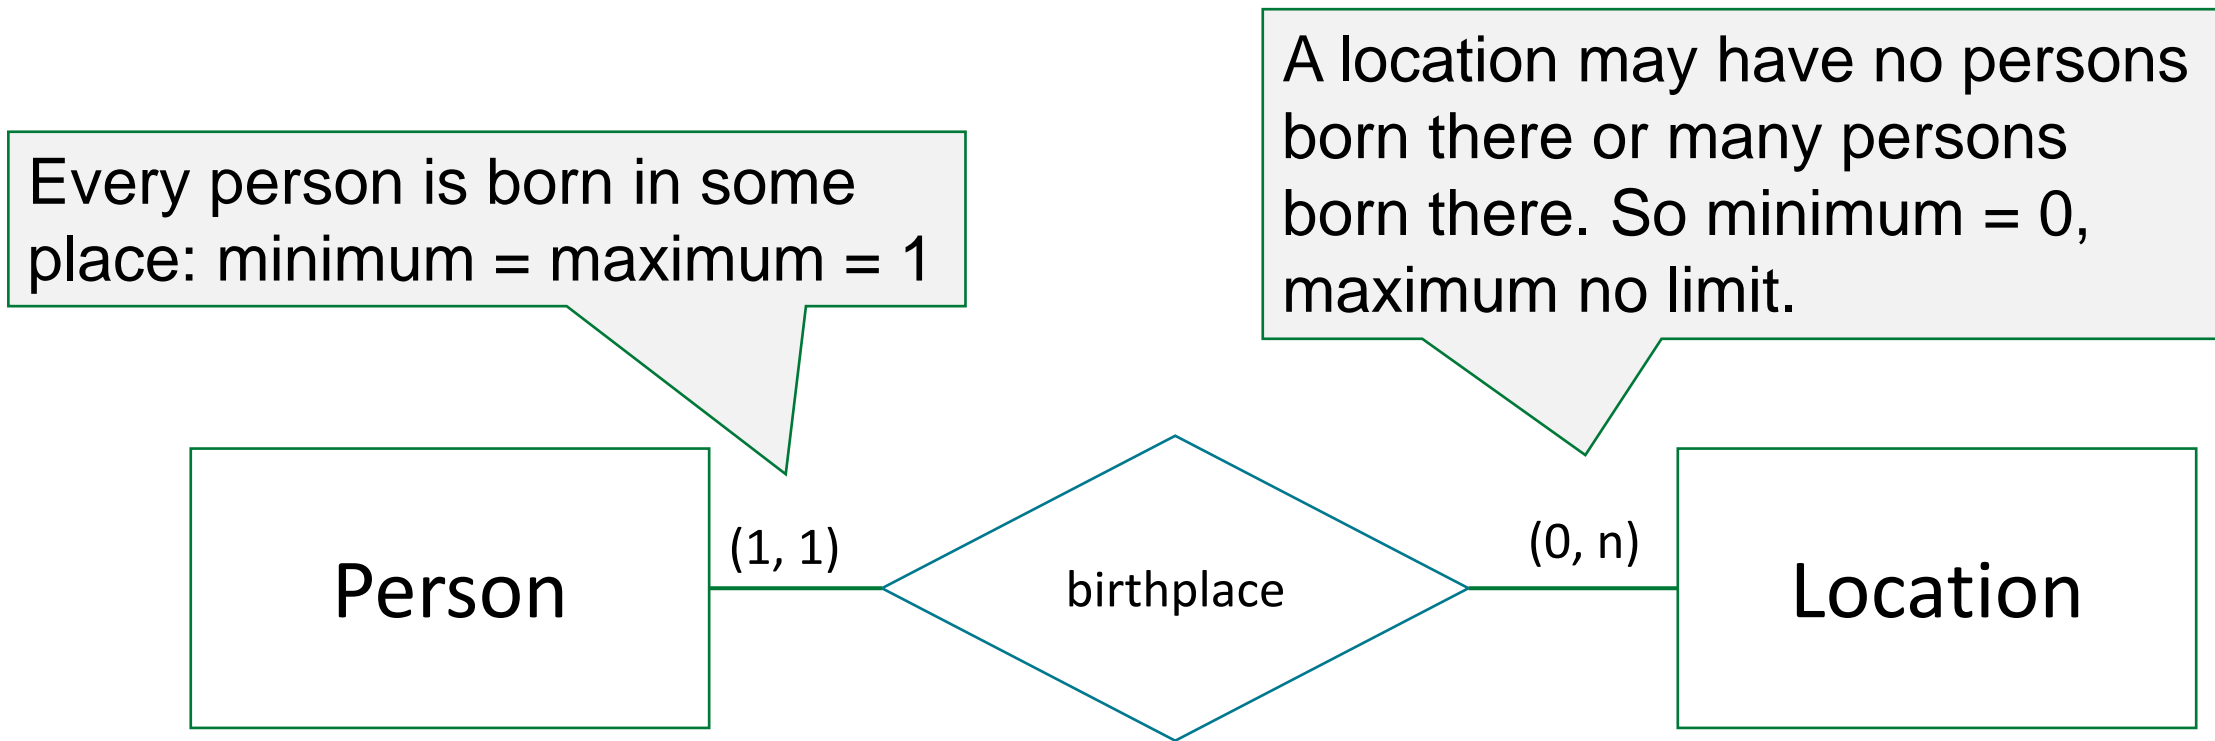
But we cannot represent the fines satisfactorily in this design. It is not possible to represent a book having one member with an outstanding fine, and another member that is currently borrowing.

Perhaps the fine attribute can be moved to the Member table? Yes, but we wouldn't know which book the fine was charged for.

Perhaps we can add the Borrow fields to both the Book table and the Member table? It would introduce **redundancy** in the design, which can be problematic.

**All said and done, perhaps the 3-table design is the best design for this problem!**

# Additional Example: Database about Birthplaces



# Table design



**Straight design:** Three tables (2 entities + 1 relationship)

**Person**(pid, last name, first name)

**Location**(locid, name)

**Birthplace**(pid, locid)

**Optimised design:**

Since the multiplicity of **Birthplace** at the **Person** end is (1,1), the **Birthplace** table can be merged into the **Person** table.

**Person**(pid, lastname, first name, birthplace)

Here, “birthplace” is the locid of the Location

Note: The **Birthplace** table cannot be merged into the **Location** table (because the maximum multiplicity is not 1).

# Keypoints of representing relationships


In a straight design, each entity and each relationship is made into a separate table.

If a relationship has (0, 1) or (1, 1) multiplicity with an entity, its table can be merged with that of the entity.

(But it is *not necessary* to do so.)

If the multiplicity is (0, 1) then the relationship attributes should allow for the NULL possibility.

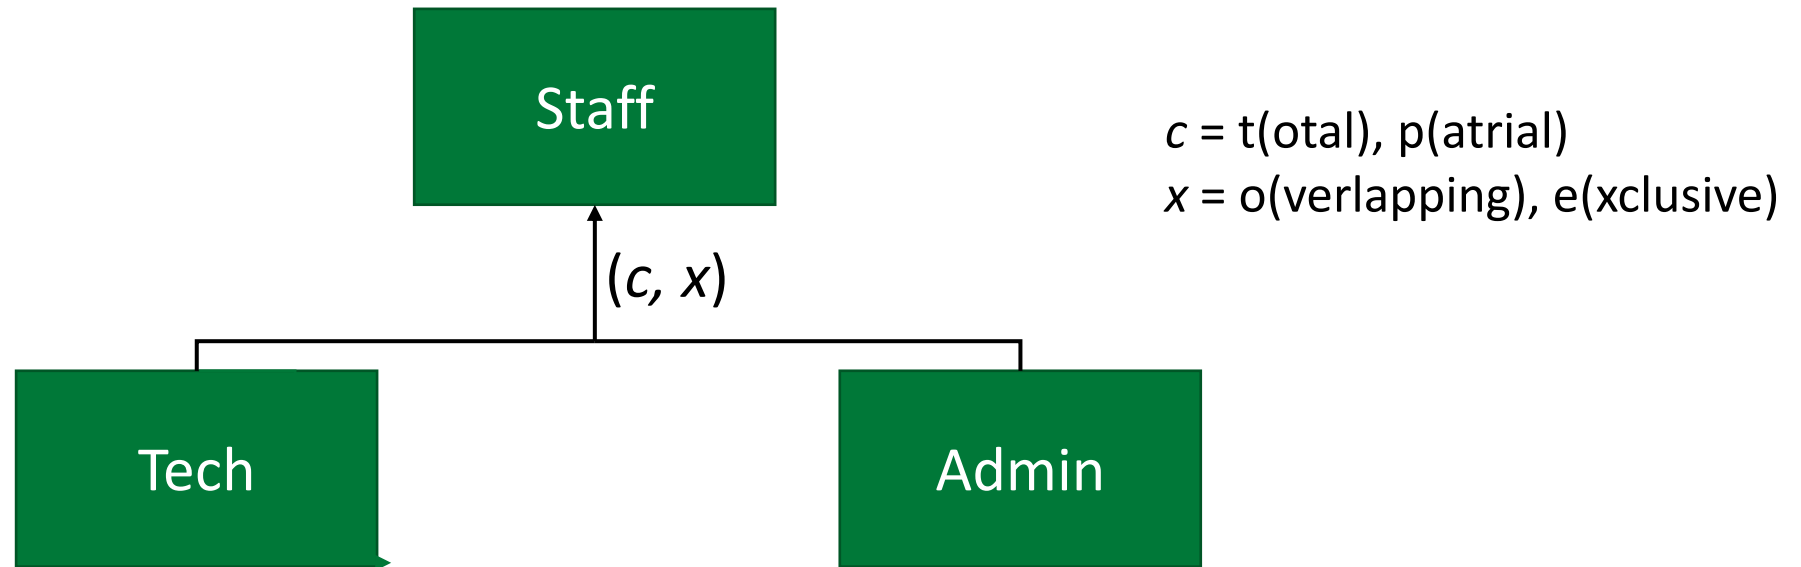
# Hierarchy



Three possible choices for table design:

1. Keep all three
2. Keep B and C, and omit A
3. Keep A, and omit B and C

# Hierarchy example




Three possible choices for table design:

- Keep all three
- Keep Tech and Admin, and omit Staff
- Keep Staff, and omit Tech and Admin



# 1. Keeping all three (superclass + subclasses)

This is the simplest method and always applicable.



Remember that an entity of type B or C (subclasses) is **also** of type A.

So, the **primary identity** of the entity is established in type A, and **inherited** by B and C.

For example, **staffid** is defined in the Staff table, and then inherited in the Tech and Admin tables.

The superclass table contains all the *common* attributes (name, address, salary, date of joining etc.)


The subclass tables contain the *additional* attributes, e.g., **speciality** for tech staff, and **role** for admin staff.

# Sample schemas

Staff(staffid, lastname, firstname, joindate, grade, salary)

Techstaff(staffid, speciality, department)

Adminstaff(staffid, role, unit)




## 2. Keep only the subclasses

This method is only possible if the **coverage** is **total**, and **overlap** is **exclusive**.

The **primary identity** is established with respect to the superclass, even though the superclass table is not stored (a bit confusing!).

For example, **staffid** should be unique across both the Tech and Admin tables, i.e., unique within an imaginary Staff table.

Whatever attributes might have been stored in the superclass are duplicated in all the subclasses (e.g., staffid, name, salary etc.)



### 3. Keep only the superclass


The subclass tables are omitted.

So all their attributes need to be stored in the superclass.

In addition, we need to add a **variant** attribute, which specifies the subclass type of an entity. (e.g., **staff\_type** = tech or admin)


If a particular entity belongs to one variant, the attributes corresponding to the other variant would be NULL.

For example, the records for technical staff would have NULL for **role**. The records for administrative staff would have NULL for **speciality**.



# Key points of representing hierarchies

1. The simplest method is to make tables for the superclass as well as subclasses.
2. Making tables for only subclasses is possible if the annotation for the hierarchy is (t, e). Identity has to be defined with respect to an imaginary superclass table.
3. Keeping only the superclass and merging the subclass attributes is always possible, but
  - Involves adding a variant attribute for the subclass type.
  - Perhaps wastes space for the NULL attribute values.



# This Week

- Introduction
- Relational databases
- Entity-relationship (ER) diagram /Modeling
- Table design and creation, SQL
- Weak entities
- Hierarchies
- Table design (schemas) Optimisation

SQL Commands + PostgreSQL



UNIVERSITY OF  
BIRMINGHAM | DUBAI  
دبي

# SQL command example

```
create table borrow(bookid      integer not null references book(bookid),
                    member_num integer not null references member(member_num),
                    date         date not null,
                    due_date     date not null,
                    return_date  date,
                    fine         integer,

                    primary key (bookid, member_num));
```

- The general format is: *field-name, type, field-constraints*.
- We use “not null” for all fields that are not designated as possibly null in the schema.
- We use the “references” constraint for all references into other tables.
- This happens in:
  - Relationships (for referencing the entities)
  - Weak entities (for referencing the **owner** entity)
  - Hierarchies (for referencing superclass entities)

# Types in SQL

INT (or INTEGER)

DECIMAL(n, m) – decimal numbers of length n, with m digits after the point

BOOLEAN

CHAR – single character

CHAR(n) – fixed length character string

VARCHAR(n) – variable length character string with a limit

TEXT – arbitrary length string

DATE and TIME



# Generating id numbers

See Paragraph 34.

In standard SQL, we can declare

```
create sequence staffid_seq;
```

This creates a sequence **generator** in the database, so that evaluating `nextval(staff_seq)` generates a new integer value.

We can use this as a “default value” in the create table statements:

```
create table Tech(staffid integer DEFAULT nextval(staffid_seq);  
                ... )
```

```
create table Admin(staffid integer DEFAULT nextval(staffid_seq);  
                  ... )
```

# Foreign key issue: ON DELETE

Fields with “references” constraints are called **foreign keys**.

They arise in relationships, weak entities and hierarchies.

The DBMS ensures that, whenever we insert a record with a foreign key value, that value is *actually present* in the referenced table.

What happens when the referenced value is **deleted** from the foreign table?  
Then the foreign key constraint would be violated!

The DBMS blocks such deletion and gives an error (by default).

This is called “ON DELETE NO ACTION” constraint.

“NO ACTION” may be a bit confusing. It really means that the deletion would get blocked!

# ON DELETE actions

**ON DELETE NO ACTION** – default, no need to declare

**ON DELETE SET NULL**

This sets the foreign key to null, if possible.

If null is not allowed, it would give rise to an error.

**ON DELETE CASCADE**

The record with the foreign key gets deleted automatically.

Think of what the effect of these declarations would be when a library member quits and we try to delete their record.

What should happen to a borrow record that might still be referencing that member?

# Constraints in relational databases

In general, constraints are good! They allow us to detect errors early and make sure that the data in the system is **consistent** and **valid**.

**Field constraints** – constraint on the value of a field (e.g., NOT NULL)

**Record constraints** – constraints on an entire record (e.g., you might check that the due\_date is greater than the borrow\_date)

**Table constraints** – constraints on the entire table (e.g., PRIMARY KEY or UNIQUE).

**Database constraints** – constraints that span multiple tables (e.g., REFERENCES)

# Postgres SQL

Next Week Support Session



<https://www.postgresql.org/download/>



UNIVERSITY OF  
BIRMINGHAM | DUBAI  
دبي

# Summary

- Introduction
- Relational databases
- Entity-relationship (ER) diagram /Modeling
- Table design and creation, SQL
- Weak entities
- Hierarchies
- Table design (schemas) Optimisation
- SQL Commands + PostgreSQL



# Some questions you might think about

- Our table has a single field for all the authors of a book. Can we list each author individually, somehow?
- If we have multiple copies of books in the library, what can we do?
- If we want to allow a member to borrow to multiple books (say 4 books max), what changes do we need in the design?
- If membership is actually “family membership” so that anybody in the family can borrow, do we need to change the design in anyway?
- The “primary key” of the borrow table consists of two fields. Do we need two? Can we make do with one field only?

# Weekly Reading



## Reading material on Canvas

- Handout 1, 2, 3

## Textbook

**Chapter 4** – “High-level database models”, Ullman and Widom, A First Course in Database Systems, Third edition ([eBook link](#))

**Chapter 6** Database Design Using the E-R Model by Silberschatz, Database System Concepts, 7th Ed ([eBook link](#))





# Exercise Questions

**ER Modeling** (Handout 2, Page 3)

Exercise 16.....

**Generalisation hierarchy** (Handout 2, Page 7)

Exercise 24 ...

# Mid Semester Feedback to the Module Team (Dubai)





UNIVERSITY OF  
BIRMINGHAM

DUBAI  
دبي

Next  
Week



Week	Date	Topic
1	15 Jan	Searching algorithms
2	22 Jan	Binary Search Tree
3	29 Jan	Balancing Trees – AVL Tree
4	5 Feb	Databases – Conceptual Design
5	14 Feb	Databases – Logical Design & Relational Algebra
6	19 Feb	Consolidation Week
7	26 Feb	Graph Algorithms
8	4 Mar	Sorting Algorithms
9	11 Mar	Hash tables
10	18 Mar	Databases – Normalization
		Easter break and Eid break
11	22 Apr	Databases – Concurrency
12	29 Apr	Revision Week

Thank you.  
Questions?



UNIVERSITY OF  
BIRMINGHAM

DUBAI  
دبي

# Attendance